

NDU: A Generalization of Fixed-Preference Reinforcement Learning

(Authors' names blinded for peer review)

Many operations models optimize physical actions after service levels, shortage penalties, or risk weights have been fixed. Neural Differential Utility (NDU) makes those valuation quantities costly endogenous controls. We show that fixed-preference reinforcement learning is a frozen-valuation limit, while active valuation switching yields a continuation-value-dependent support Bellman operator that cannot be collapsed to an evaluator-preserving fixed-reward/action-only augmented-state MDP, exogenous adaptive preference, static-multiplier CMDP, or fixed-parameter risk/robust formulation except in degenerate affine cases. For inventory service systems, NDU gives an order-up-to policy coupled with a monotone service-pressure rule and, under state-varying shortage regimes, strict stable improvements in theta-free physical cost over fixed service pressure under stated density and gap conditions. The numerical issue is shadow-price observability: when preference dynamics are nearly deterministic, $Z_u = \sqrt{2\varepsilon}P_u$ makes Z -inversion amplify equal-scale target noise by $1/(2\varepsilon)$, while direct- P value-gradient NBO avoids this singular scaling. We give a compact-domain HJB-solver theorem for NBO and verify its assumptions through theory-covered compact, analytic uniform, finite-cover inventory, cover-refinement, validation-cover, direct- P /learned- Z , exact-DP, SKU replay, and residual/greedy-gap audits.

Key words: stochastic control; endogenous valuation; adaptive preferences; forward-backward stochastic differential equations; dynamic optimization

Subject classifications: Dynamic programming/optimal control: stochastic control; continuous time.

| Operations management: inventory/service systems. | Simulation: computational experiments.

Empirical scope. The empirical and reproducibility language in this manuscript is scoped to the explicit model, assumptions, and files included with the paper. It is not used as an additional independent benchmark claim unless a concrete result table, command log, seed record, or checksum is cited at the relevant statement; otherwise the statement should be read as qualitative or operator-level motivation rather than standalone regenerated evidence.

Assumption scope. The theoretical statements in this section are conditional on the stated model, regularity, approximation, and comparison assumptions. They should not be read as empirical superiority, implementation correctness, or out-of-scope benchmark claims unless a separate result, artifact, or verification record is cited at the relevant statement.

Statement scope. The adjacent statement is scoped to the model, assumptions, source-package artifacts, and verification commands explicitly cited in this manuscript. This wording is not an additional empirical, theoretical, or implementation claim beyond those cited materials.

Table 1 OR interpretation map for endogenous valuation control.

Domain	Physical action	Valuation-control coordinate	Role in the paper
Procurement	order, expedite, or allocate supply	shortage or supplier-risk penalty	motivating example
Queueing	staffing, routing, or admission	service-level pressure	motivating example
Revenue management	price or protection level	reference-demand or risk target	motivating example
Inventory service	order-up-to level	service pressure / shortage penalty	primary OR theorem and benchmark layer

1. Introduction

Operations models often optimize a physical action against a fixed cost, reward, or service-level objective. In many dynamic settings, however, the way future states are valued is itself a controlled operational object. A procurement system may adjust shortage penalties as supply risk changes; a queueing or congestion system may adapt service-level targets as capacity tightens; and a revenue-management system may update reference-demand or habit parameters as market conditions evolve. These examples motivate a stochastic-control question: *How should a decision maker jointly control physical actions and the internal valuation parameters that score future states?*

We study this question through **Neural Differential Utility (NDU)**, a continuous-time stochastic-optimization framework in which part of the valuation rule is endogenous. The name NDU is retained because the preference state u_t may be parameterized as the hidden state of a controlled neural differential equation (Chen et al., 2018). Once u_t is included in the state, NDU is an augmented-state stochastic-control problem; the novelty is not that augmented states are unavailable, but that the valuation coordinate is itself a costly control. Relative to fixed-reward/action-only RL, both original-state fixed-preference policies and augmented-state fixed- θ action-only policies embed by holding $\theta_t \equiv \bar{\theta}$, while active θ_t creates HJB terms and a positive net Bellman margin that fixed-valuation classes cannot represent. The full control is $\alpha_t = (a_t, \theta_t)$, where a_t is the operational action and θ_t governs resource-constrained valuation adjustment. The neural network V_ϕ introduced later is the value-function ansatz used by the NBO/HJB solver route.

OR interpretation. Table 1 is only an interpretation map, not a benchmark suite or an extra theorem/evidence claim. It fixes the operational meaning of the valuation-control coordinate used below.

1.1. Mathematical Formulation

The formulation is a fully coupled Forward–Backward Stochastic Differential Equation (FBSDE): the forward component evolves the external state c_t and preference state u_t under $\alpha_t = (a_t, \theta_t)$, while the backward component characterizes the continuation value Y_t and volatility/gradient trace Z_t . The coupling is bidirectional. (Y_t, Z_t) determine the Hamiltonian maximizer $\alpha_t^* = (a_t^*, \theta_t^*)$, and that joint control feeds back into both external and preference-state drift, placing the model in the coupled-FBSDE class studied by Ma et al. (1994) and Yong and Zhou (1999).

The paper has a *formal layer* and a computational layer. Formally, the frozen fixed-preference face still gives a formal strict generalization of fixed-preference reinforcement learning, but the main separation is stronger: optimizing out θ gives a support Bellman/HJB operator whose active affine faces switch with the continuation value. This non-affine reduced operator cannot be collapsed to an evaluator-preserving fixed-reward/action-only augmented-state MDP, nor to standard exogenous adaptive preference, static-multiplier CMDP, or fixed-parameter risk/robust reductions except in degenerate cases. The OR layer then proves base-stock/service-pressure structure, a strongly concave valuation-control computational reduction, and a conditional theta-free physical-cost improvement theorem for state-varying service regimes. The distinctive numerical issue is observability in the degenerate preference block: P_u may be unrecoverable from Z_u , direct- P conditions that block, and auxiliary-noise martingale claims require compatible shadow-price bounds. Computationally, a PyTorch-AD HBO/NBO actor-critic reports HJB residual, terminal loss, greedy/value-gradient construction, and held-out rollout performance, with CEM, exact-DP, and external-evaluator audits as finite-benchmark checks.

The first contribution is an operator-level separation theorem. Fixed-preference RL embeds as a frozen valuation face, but the main distinction is stronger: after optimizing out the valuation control, NDU yields a support Bellman operator whose active faces switch with the continuation value. This non-affine operator cannot be represented by a fixed-reward/action-only augmented-state MDP without changing the evaluator, expanding the action set to include valuation controls, or collapsing to a payoff-irrelevant valuation channel.

The second contribution is an OR structure theorem. In periodic-review inventory with endogenous service pressure, optimal NDU policies have a base-stock form with a monotone service-pressure selector. Backlog, demand stress, and capacity tightness increase the optimal valuation pressure, while the adjustment cost induces mean reversion.

The third contribution is a computability theorem. Under strongly concave valuation adjustment, the valuation-control inner problem has a closed-form projected selector, reducing joint action-valuation search to physical-action search plus an analytic projection and yielding explicit greedy-gap bounds.

The fourth contribution links valuation control to physical OR metrics. In state-varying shortage-cost regimes, fixed service-pressure policies cannot match all physical critical fractiles, while NDU can choose state-dependent service pressure; this yields a strict and distributionally stable improvement in theta-free physical cost under checkable density and gap conditions. Propositions 2.3–2.4 and Theorem 2.5 then connect the operator structure to direct- P conditioning and residual/greedy-gap HJB-solver certificates.

Scope and roadmap. Table 2 fixes the claim boundary: operator-level non-collapse, frozen-face embedding/separation, OR service-pressure structure, strongly concave valuation reduction, conditional physical-metric improvement, compact-domain NBO HJB-solver certificates, analytic uniform and finite-cover instantiations, and nested cover refinement. The main empirical section centers the inventory/service-pressure application with SKU/exact-DP OR checks, direct- P/Z conditioning, and HBO/NBO residual diagnostics. The evidence layer includes an analytic uniform residual/terminal/greedy-gap certificate sequence, a seven-rung HJB-solver certificate ladder, and a neural validation-cover HJB gap audit with stability guards.

We use c_t for the operational state, u_t for the endogenous valuation/preference state, $x_t = (c_t, u_t)$ for the augmented state, a_t for the overt action, θ_t for valuation control, and $\alpha_t = (a_t, \theta_t)$ for the full control. In OR terms, θ_t can represent an adaptive shortage penalty, service-level pressure, risk-tolerance target, or reference-demand parameter. The novelty claim is not that augmented-state dynamic programming is unavailable; it is that making such valuation quantities costly controls creates a concrete modeling and numerical-observability problem relative to fixed-preference control on the original state.

1.2. Related Literature

The paper draws on stochastic optimal control, behavioral economics, and deep learning. We distinguish NDU from existing paradigms along three theoretical axes.

Stochastic Differential Utility (SDU) and Backward Stochastic Differential Equations (BSDEs). NDU builds on BSDE theory (Pardoux and Peng, 1990) and stochastic differential utility (Duffie and Epstein, 1992), but here the driver is affine in continuation value and independent of volatility; the recursion sits in the forward preference-state dynamics, not in a nonlinear aggregator over Y_t . The preference channel carries autonomous controlled dynamics: θ_t^* depends on the value-gradient trace and feeds back into u_t , putting the model in the coupled-FBSDE setting of Ma et al. (1994); Yong and Zhou (1999); related deep-BSDE control analysis includes Huang et al. (2025).

Meta-Learning and Inverse Reinforcement Learning (IRL). Standard RL (Sutton and Barto, 2018) optimizes against a fixed reward, while NDU controls internal preference parameters. This is not IRL (Ng and Russell, 2000; Abbeel and Ng, 2004): no expert reward is inferred; θ_t is chosen to maximize value net of adjustment cost. The closest economic analogy is rational inattention (Sims, 2003), with $\mathcal{L}$ representing finite valuation-adjustment resources; continuous-time RL/HJB learning is related through Shirani Faradonbeh and Shirani Faradonbeh (2023).

State augmentation, adaptive preferences, and controlled utility processes. Classical stochastic control routinely handles enlarged states, and risk-sensitive or recursive utility models can be written with additional state or continuation-value variables. NDU should therefore not be read

Table 2 Scope map for the main formal claims.

Result	Role	Claim boundary
Proposition 4.1 and Corollary 4.2	Frozen-face embedding	Fixed-preference/action-only RL and augmented fixed- θ control embed exactly; positive valuation-control margins imply non-collapse to frozen-valuation re-description.
Theorem 4.3	Operator-level non-collapse	Optimizing out active valuation control yields a support Bellman term whose continuation-value-dependent exposed faces cannot be represented by a fixed affine reward/transition primitive.
Corollary 4.4	Model taxonomy	Exogenous adaptive preference, static-multiplier CMDP, and fixed-parameter risk/robust reductions match NDU only after changing the evaluator/action set or collapsing to degenerate affine faces.
Theorem 4.5	OR structure	Periodic-review inventory has base-stock/order-up-to structure coupled with a monotone endogenous service-pressure selector.
Theorem 4.6	Computability	Strongly concave valuation control admits a closed-form/projected inner selector and decomposes the residual/greedy-gap certificate into physical-action and inner-valuation gaps.
Theorem 4.7	Physical metrics	In state-varying shortage/service regimes, NDU can match regime critical fractiles and strictly, stably improve theta-free physical cost over fixed service pressure under density/gap conditions.
Lemma 2.3	Core conditioning claim	Exact Z_u -inversion rescaling gives a matching $1/\varepsilon$ lower-bound separation under equal-scale target noise; direct- P has no singular amplification, but still needs value-gradient/adjoint consistency.
Proposition 2.4	Algorithm link	In a finite-dimensional value-gradient route, $P_\beta = \nabla_x V_\beta$ is integrable by construction and HJB/terminal residuals quantify weak-form consistency.
Theorem 2.5	NBO HJB-solver theorem	On compact domains with comparison and selector regularity, sup-norm optimal-Hamiltonian residual, terminal residual, and greedy-action gap bound value and policy error; NBO sequences closing these gaps are HJB solvers.
Corollary 2.7	Finite-cover HJB certificate	On a finite state-time/action cover, Bellman residual, terminal residual, and greedy-action gap give the exact discrete HJB certificate; zero gaps certify the finite-cover solve.
Corollary 2.8	Cover-refinement sequence	HJB Under monotone/stable/consistent interpolation and action-cover Lipschitz bounds, exact finite-cover solves with vanishing fill distance form an HJB-solver certificate sequence.
Proposition 2.9, Corollaries 2.10–2.13	Validation-cover neural gaps	Sampled residual/terminal/greedy sup gaps lift under fill-distance/Lipschitz guards; if the Lipschitz constants and fill distances are deterministic certificates, the displayed value/policy bound is a proof, while tuple completeness, validation density, greedy-tail, rollout, replay, empirical-envelope, and interval-domain guards audit the sampled rows.
Proposition 2.24	HJB-solver certificate ladder	Exact, uniform, finite-cover, cover-refinement, sampled neural, and direct- P representation rungs are reported as one auditable certificate architecture for NBO.
Proposition 2.25 and Proposition 2.27	Referee-objection/equivalence support map	Main objections—including state augmentation/reward shaping, HJB certification, finite-cover refinement, sampled neural gaps, direct- P , and fairness—are mapped to theorem anchors and checked audit rows; this is an internal evidence map, not an external referee closure.
Proposition 2.26	Six-pillar claim-support audit	Six checked pillars organize the evidence for policy-class generalization, strict separation, HJB-solver core, cover refinement, validation-cover neural guards, direct- P , and same-information rollout support without upgrading sampled neural rows into analytic global certificates.
Proposition 2.28	Strict face hierarchy	Exact policy-class/finite-cover audits show Full NDU strictly exceeds action-only and valuation-only frozen/single-channel faces, with joint-control refinement and service-pressure diagnostics.
Proposition 2.29	Referee-response support synthesis	Eight checked guards tie generalized-RL strictness, anti-relabeling closure, HJB certificates, validation-cover neural evidence, direct- P , same-information evidence, and artifact traceability into one evidence-preserving response matrix.
Lemma 3.6	Restricted fixed- ε verification	Standard short-horizon Picard existence under selector Lipschitzness, projection closure, and state-independent/bounded/clipped shadow-price regimes; not the main novelty and not generic neural NDU.
Proposition 3.8	Supporting global extension	Sufficient Picard-contraction inequality; constants deteriorate with ε and are not uniform as $\varepsilon \downarrow 0$.
Proposition 3.11	PDE bridge	Viscosity value convergence under comparison and local boundedness; no derivative or shadow-price process implied.
Proposition 3.12	Shadow-price bridge	Martingale vanishing only for a compatible classical/semiconcave/SMP shadow-price selection with uniform $\mathcal{H}^2$ bound.
Proposition 4.9	Frozen-valuation limit	High adjustment cost recovers the embedded fixed-preference HJB; finite-cost NDU remains an enlarged control class with the fixed class nested inside.
Proposition 5.1	Decision layer	Finite-benchmark audit certificate: class inclusion, attained-policy margins, optimizer-gap/interval thresholds, interventions, and exact-DP zero-gap cases; complements the HBO/NBO actor-critic audit.
Section 5	Evidence	Inventory-centered OR evidence: physical service metrics, valuation-welfare interpretation, direct- P /learned- Z ablation, exact-DP/SKU checks, analytic uniform HJB certificate, finite-cover inventory Bellman-HJB certificate, cover-refinement evidence, an integrated HJB-solver certificate ladder, a referee-objection support audit, a validation-cover HJB gap audit with stability guards, and HBO/NBO residual/greedy-gap diagnostics.

as a claim that augmented-state dynamic programming is unavailable. The distinction is modeling and operational: the added state is an internal valuation state whose adjustment is itself costly and controlled, rather than an exogenous sufficient statistic. This places the paper near adaptive-preference,

habit-formation, and reference-dependent utility models (Campbell and Cochrane, 1999; Kahneman and Tversky, 1979), controlled-utility and recursive-utility formulations (Duffie and Epstein, 1992; El Karoui et al., 1997), rational inattention (Sims, 2003), and risk-sensitive or robust control (Whittle, 1990; Hansen and Sargent, 2008), while retaining a fixed-preference stochastic-control limit for comparison.

Inventory, service-level control, and dynamic programming in operations. The OR-facing inventory example is closest to classical periodic-review inventory and service-level control rather than to portfolio theory. The base-stock and exact-DP checks follow the inventory tradition associated with dynamic programming and (s, S) / order-up-to structure (Scarf, 1960; Zipkin, 2000; Porteus, 2002); the approximate-DP and MPC comparators are included to connect the benchmark to operational dynamic programming and rollout-style control (Bertsekas, 2012; Powell, 2011). The NDU interpretation is not that dynamic service targets or shortage penalties are new. It is that making a service-pressure / shortage-penalty parameter θ_t a costly endogenous control exposes a shadow-price and observability issue, and it gives a unified way to compare fixed-service base-stock, rolling-forecast, approximate-DP, inventory-RL, and valuation-control policies under common physical evaluators.

Neural Differential Equations and Deep Control. “Neural” denotes an expressive continuous-time parameterization option for u_t , inspired by Neural ODEs (Chen et al., 2018); it is not the core theoretical novelty. The theory is stated for structured drift $h(c, u; a, \theta)$. Numerically, we use a PyTorch-AD HBO/NBO actor-critic HJB-residual solver over compact sampled domains, aligned with neural-operator discretization ideas (Li et al., 2020), and retain bounded CEM maps as robustness audits. Deep BSDE work (Han et al., 2018; E and Yu, 2018) motivates the observability and direct- P route, but we do not claim a new general Deep-BSDE convergence theorem.

2. The Neural Differential Utility Framework

Let $(\Omega, \mathcal{F}, \{\mathcal{F}_t\}_{t \geq 0}, \mathbb{P})$ be a complete filtered probability space supporting a standard d -dimensional Brownian motion W_t . We consider a finite time horizon $T > 0$. The filtration $\mathcal{F}_t$ is the natural filtration generated by W_t , augmented by all $\mathbb{P}$ -null sets.

2.1. State Dynamics (Forward System)

The agent interacts with an external stochastic environment whose controllable external state is denoted by $c_t \in \mathbb{R}^n$. The agent also maintains an internal *preference state* $u_t \in \mathbb{R}^q$, modeled as a controlled state process and optionally parameterized as the hidden state of a continuous-time recurrent / Neural-ODE-style module. Full NDU optimizes a joint control $\alpha_t = (a_t, \theta_t)$, where $a_t \in \mathcal{A}$ is an overt action and $\theta_t \in \Theta$ is an internal preference-control variable. Unless otherwise stated, c_t

and u_t are treated as column vectors, so block concatenations such as $X_t = [c_t^\top, u_t^\top]^\top$ and the block diffusion matrix $\Sigma_\varepsilon(X_t)$ are dimensionally consistent.

Let $\mathcal{A} \subset \mathbb{R}^{k_a}$ and $\Theta \subset \mathbb{R}^{k_\theta}$ be compact sets, endowed with the Euclidean norms induced by the standard inner products on their ambient spaces. We write $\mathcal{A} \times \Theta := \mathcal{A} \times \Theta$ for the feasible set of joint overt-action and valuation-control choices. An admissible control process is an $\mathcal{F}_t$ -adapted pair $\alpha = \{(a_t, \theta_t)\}_{t \in [0, T]}$ with values in $\mathcal{A} \times \Theta$. Throughout, θ_t denotes an endogenous valuation-control choice; the frozen fixed-preference parameter in Section 4 is written $\bar{\theta}$ to keep the comparison separate.

Definition 2.1 (Controlled External and Preference Dynamics) *The joint state process $X_t = [c_t^\top, u_t^\top]^\top \in \mathbb{R}^{n+q}$ evolves according to the forward stochastic differential equation (SDE):*

$$dc_t = \mu_c(c_t, a_t)dt + \sigma_c(c_t)dW_t, \quad c_0 \in \mathbb{R}^n, \quad (1)$$

$$du_t = h(c_t, u_t; a_t, \theta_t)dt + \sigma_h(c_t, u_t)dW_t, \quad u_0 = u_{\text{init}}(c_0), \quad (2)$$

where $\mu_c: \mathbb{R}^n \times \mathcal{A} \rightarrow \mathbb{R}^n$ is the controlled external-state drift, $h: \mathbb{R}^n \times \mathbb{R}^q \times \mathcal{A} \times \Theta \rightarrow \mathbb{R}^q$ represents the drift of the preference state, and the diffusions $\sigma_c: \mathbb{R}^n \rightarrow \mathbb{R}^{n \times d}$ and $\sigma_h: \mathbb{R}^n \times \mathbb{R}^q \rightarrow \mathbb{R}^{q \times d}$ are taken to be control-independent in the main theory.

Remark 2.2 (Control-Independent Diffusion) To ensure the well-posedness of the coupled system in the sense of [Yong and Zhou \(1999\)](#), we assume the diffusions σ_c and σ_h do not depend explicitly on the control $\alpha_t = (a_t, \theta_t)$. Within the state dynamics, the joint control acts through the drift channels μ_c and h , which is the drift-control regime studied in the main theory.

2.2. The Objective Functional

The agent seeks to maximize a state-dependent additively separable discounted utility functional.¹ In full NDU, this is a joint optimization over overt action and internal parameter control. The overt action channel a_t affects both external-state evolution and instantaneous utility, while the preference-control channel θ_t affects both preference dynamics and valuation, subject to the resource cost $\lambda\mathcal{L}$.

For a given admissible joint control $\alpha = (a, \theta)$, the value function $J(t, c, u; \alpha)$ is defined over the interval $[t, T]$:

$$J(t, c, u; \alpha) = \mathbb{E} \left[\int_t^T e^{-\rho(s-t)} (f(c_s, u_s; a_s, \theta_s) - \lambda\mathcal{L}(c_s, u_s; a_s, \theta_s)) ds + e^{-\rho(T-t)} \Psi(c_T, u_T) \middle| \mathcal{F}_t \right], \quad (3)$$

¹We use the SDU/BSDE representation as a convenient continuation-value formalism, but for the architectures studied here the driver $g(x, y, \alpha) = f(x, \alpha) - \lambda\mathcal{L}(x, \alpha) - \rho y$ is affine in y and independent of z . Accordingly, the objective is analytically equivalent to standard discounted utility with an endogenously evolving preference state; the recursion resides in the forward preference dynamics rather than in a nonlinear aggregator over continuation value.

where $\rho \geq 0$ is a discount rate, f is the net operating utility aggregator, $\mathcal{L}$ is the valuation-adjustment cost, and Ψ is a terminal bequest function. Ordinary action costs such as ordering, holding, capacity, or transaction costs are included in f or in the theta-free physical evaluator; the uniqueness condition in Assumption 4.8 applies only to the valuation-adjustment component. The agent's value function is $V(t, c, u) = \sup_{\alpha} J(t, c, u; \alpha)$.

2.3. Degeneracy of Preference Dynamics and the Numerical Observability Paradox

A fundamental structural obstruction in formulating NDU via standard Deep FBSDE methods is the degeneracy of the preference state dynamics. In the idealized NDU formulation, the internal preference state u_t may evolve as a controlled ODE/Neural-ODE-style hidden state. While driven by the stochastic environment c_t , the preference state itself often has zero intrinsic diffusion ($\sigma_h \equiv 0$).

This degeneracy creates a dichotomy between the theoretical existence of the adjoint process and its numerical recoverability, which we term the *Observability Paradox*.

2.3.1. The Observability Paradox Let $V(t, x)$ denote the value function. Throughout the main text, gradients $P = \nabla_x V$ and BSDE volatilities Z are column vectors; the scalar martingale increment is written as $Z_t^\top dW_t$. In the unregularized model $P_t \in \mathbb{R}^{n+q}$ and $Z_t \in \mathbb{R}^d$; after adding the auxiliary preference noise, $Z_t^\varepsilon \in \mathbb{R}^{d+q}$. In the theory of non-degenerate stochastic control, the solution (Y_t, Z_t) to the associated BSDE is related to the value function via the nonlinear Feynman-Kac formula. Specifically, the volatility process $Z_t = \Sigma^\top P_t$ is the orthogonal projection of the value gradient onto the noise manifold generated by the Brownian motion. When the same Brownian motion drives both state components, this projection takes the summed form:

$$Z_t = \Sigma(X_t)^\top \nabla_x V(t, X_t) = \sigma_c(c_t)^\top \nabla_c V + \sigma_h(c_t, u_t)^\top \nabla_u V. \quad (4)$$

The optimal joint control $\alpha_t^* = (a_t^*, \theta_t^*)$ is determined by Hamiltonian maximization:

$$\alpha_t^* \in \arg \max_{(a, \theta) \in \mathcal{A} \times \Theta} \{ (\nabla_c V_t)^\top \mu_c(c_t, a) + (\nabla_u V_t)^\top h(c_t, u_t; a, \theta) + f(c_t, u_t; a, \theta) - \lambda \mathcal{L}(c_t, u_t; a, \theta) \}. \quad (5)$$

Standard Deep BSDE solvers (Han et al., 2018) approximate Z_t directly using a neural network to satisfy the martingale representation property. However, in the ideal NDU limit where $\sigma_h \rightarrow 0_{q \times d}$, the contribution of $\nabla_u V$ to Z_t vanishes identically, regardless of the magnitude of the shadow price $\nabla_u V$.

Consequently, recovering the gradient vector $\nabla_u V$ from the volatility Z_t requires inverting the map $p \mapsto \Sigma^\top p$. When Σ is singular, this operation is ill-posed. The discussion here isolates the preference block: if the external diffusion σ_c is also not full row rank, then P_c is likewise only observable on the external volatility image. The supporting fixed- ε theory below therefore assumes a full-row-rank external block before focusing on the additional degeneracy created by the vanishing preference block.

While the full adjoint gradient remains mathematically well-defined (via the Pontryagin Maximum Principle), the parameter-control component becomes numerically unobservable from the trajectory of Z_t . A standard solver attempting to regress $Z_t \approx 0$ will face a vanishing signal-to-noise ratio in the preference block, failing to learn the direction required to steer θ_t . The overt-action block is unaffected only in the separable case where $h_a = 0$ (or the Hamiltonian dependence on a is independent of P_u). If a also moves the preference drift, then the same missing shadow price enters the a -first-order condition through $P_u^\top \partial_a h$.

2.3.2. Regularization and Shadow-Price Conditioning The formal regularization is stated once in Section 3.4; here we only preview the construction needed for the observability discussion. For a fixed $\varepsilon > 0$, the degenerate preference block is perturbed by an auxiliary Brownian motion and

$$dX_t^\varepsilon = b(X_t^\varepsilon, \alpha_t)dt + \Sigma_\varepsilon(X_t^\varepsilon)d\mathbf{W}_t, \quad \Sigma_\varepsilon(x) = \begin{bmatrix} \sigma_c(c) & 0 \\ 0 & \sqrt{2\varepsilon}I_q \end{bmatrix}. \quad (6)$$

A standard Z -network would recover the preference gradient through the unstable inverse $P_u \approx Z_u/\sqrt{2\varepsilon}$. The direct- P parameterization avoids this conditioning problem by treating the shadow price as the learned object and deriving the volatility by the stable forward map

$$Z_t^\varepsilon(P_t) = \Sigma_\varepsilon(X_t^\varepsilon)^\top P_t = \begin{bmatrix} \sigma_c(c_t)^\top P_{c,t} \\ \sqrt{2\varepsilon}P_{u,t} \end{bmatrix}. \quad (7)$$

The associated value-process dynamics and feedback selector are

$$dY_t^\varepsilon = -(f(X_t^\varepsilon; \alpha_t) - \lambda\mathcal{L}(X_t^\varepsilon, \alpha_t) - \rho Y_t^\varepsilon)dt + (Z_t^\varepsilon(P_t))^\top d\mathbf{W}_t, \quad (8)$$

$$\alpha_t^\varepsilon \in \arg \max_{\alpha \in \mathcal{A} \times \Theta} \{P_t^\top b(X_t^\varepsilon, \alpha) + f(X_t^\varepsilon; \alpha) - \lambda\mathcal{L}(X_t^\varepsilon, \alpha)\}. \quad (9)$$

This is the conditioning half of the HJB-solver argument. The solver route must also enforce adjoint consistency: (i) a value-network route $P_\phi = \nabla_x V_\phi$, which makes the shadow price conservative by construction; (ii) an independent shadow-price route with adjoint/HJB residual penalties; or (iii) a supervised-oracle certificate route on theorem-covered instances. The PyTorch-AD HBO/NBO implementation uses the first route: the Hamiltonian shadow price is $P_\beta = \nabla_x V_\beta$ from a trainable value critic, while Theorem 2.5 gives the residual/terminal/greedy-gap conditions under which NBO is an HJB solver on compact domains. The same-benchmark direct- P/Z diagnostic and the isomorphic direct- P versus learned- Z inventory solver ablation verify the singular-conditioning mechanism inside the OR simulator and show why the value-gradient direct- P route is the stable NBO representation.

Lemma 2.3 (Direct- P conditioning under equal-scale Z -target noise) *Fix $\varepsilon > 0$ and suppose the preference-block relation is $Z_u^\varepsilon = \sqrt{2\varepsilon} P_u$. Let a learned or observed Z target satisfy*

$$\hat{Z}_u^\varepsilon = Z_u^\varepsilon + \eta_Z, \quad \mathbb{E}\|\eta_Z\|^2 \leq \delta_Z^2,$$

and recover $\hat{P}_u^Z = \hat{Z}_u^\varepsilon / \sqrt{2\varepsilon}$. Then

$$\mathbb{E} \|\hat{P}_u^Z - P_u\|^2 = \frac{1}{2\varepsilon} \mathbb{E} \|\eta_Z\|^2 \leq \frac{\delta_Z^2}{2\varepsilon}.$$

By contrast, if direct shadow-price supervision gives $\hat{P}_u^P = P_u + \eta_P$ with $\mathbb{E} \|\eta_P\|^2 \leq \delta_P^2$, then

$$\mathbb{E} \|\hat{P}_u^P - P_u\|^2 \leq \delta_P^2,$$

with no singular $1/\varepsilon$ amplification. Moreover, if $\mathbb{E} \|\eta_Z\|^2 \geq \underline{\delta}^2$ and $\mathbb{E} \|\eta_P\|^2 \leq \bar{\delta}^2$ with constants independent of ε , the Z -route lower bound divided by the direct- P upper bound is at least $\underline{\delta}^2 / (2\varepsilon \bar{\delta}^2)$, hence diverges at rate $1/\varepsilon$. Thus equal-scale target noise in the Z_u route is amplified by order $1/\varepsilon$ when translated into the preference shadow price. This is a conditioning statement only: unless $P_\phi = \nabla_x V_\phi$ or an adjoint/HJB residual enforces integrability and optimality, an independently learned P_ϕ need not be a conservative value gradient or a valid adjoint process.

Proof. The identities follow by subtracting $Z_u^\varepsilon = \sqrt{2\varepsilon} P_u$ and rescaling. The lower-bound separation uses the same identity with $\mathbb{E} \|\eta_Z\|^2 \geq \underline{\delta}^2$ and compares that lower bound with the direct- P upper bound. The last claim is the standard distinction between pointwise vector-field regression and value-gradient/adjoint consistency.

Proposition 2.4 (Finite-dimensional value-gradient/HJB-residual consistency) *Fix a compact numerical domain D , a C^2 value critic V_β , and actor A_η . Set $P_\beta = \nabla_x V_\beta$ and $Z_\beta = \Sigma_\varepsilon^\top P_\beta$. The direct shadow price is then a conservative value gradient by construction and the Z component is obtained by the stable forward map, not by inverting $Z_u / \sqrt{2\varepsilon}$. If the native HJB residual $R_{\beta,\eta}$ and terminal residual vanish in $L^2(\mu_D)$, then for every square-integrable test function φ the weak HJB defect satisfies $|\int_D R_{\beta,\eta} \varphi d\mu_D| \leq \|R_{\beta,\eta}\|_2 \|\varphi\|_2 \rightarrow 0$. Nonzero residuals therefore quantify the sampled-domain residual gap that Theorem 2.5 requires the solver to close.*

Proof. Integrability follows because P_β is the gradient of a scalar critic; the volatility identity follows from Itô's formula; the weak-form bound is Cauchy–Schwarz, with the same argument at the terminal boundary.

Theorem 2.5 (NBO HJB-solver certificate) *Fix a compact parabolic domain $D = [0, T] \times K$ with terminal face $\{T\} \times K$. Suppose the NDU HJB on D has comparison in the relevant bounded viscosity class, compact controls, a continuous Hamiltonian $H(x, a, p, M)$ that is Lipschitz on the range reached by a C^2 critic family, and a measurable selector whose greedy gap is well defined. For a critic $V_\beta \in C^{1,2}(D)$ define the optimal-Hamiltonian and actor-induced residuals*

$$\begin{aligned} R_\beta^*(t, x) &= -\partial_t V_\beta(t, x) - \sup_a H(x, a, \nabla V_\beta, \nabla^2 V_\beta), \\ R_{\beta,\eta}(t, x) &= -\partial_t V_\beta(t, x) - H(x, A_\eta(t, x), \nabla V_\beta, \nabla^2 V_\beta). \end{aligned}$$

Assume

$$\begin{aligned} \|R_\beta^*\|_{L^\infty(D)} &\leq \delta_H, & \|V_\beta(T, \cdot) - G\|_{L^\infty(K)} &\leq \delta_T, \\ \sup_{(t,x) \in D} \left[\sup_a H(x, a, \nabla V_\beta, \nabla^2 V_\beta) - H(x, A_\eta, \nabla V_\beta, \nabla^2 V_\beta) \right] &\leq \delta_A. \end{aligned}$$

Then there is a stability constant C_D depending only on the comparison/stability constants, horizon, discount, and compact-domain Hamiltonian Lipschitz moduli such that

$$\|V_\beta - V^*\|_{L^\infty(D)} \leq C_D(\delta_T + T\delta_H), \quad 0 \leq V^*(t, x) - V^{A_\eta}(t, x) \leq C_D(\delta_T + T\delta_H + T\delta_A).$$

If instead the verified residual is only the actor-induced residual $\|R_{\beta,\eta}\|_{L^\infty(D)} \leq \tilde{\delta}_H$, then the same value comparison holds with δ_H replaced by $\tilde{\delta}_H + \delta_A$, because the optimal-Hamiltonian residual is bounded by the actor residual plus the greedy gap. Consequently, any NBO sequence with the relevant sup-norm residual, terminal gap, and greedy-action gap all tending to zero is an HJB solver sequence on D : its value critics converge locally uniformly to the HJB value and its greedy actors are asymptotically optimal on that compact problem.

Proof. See Appendix C.7.

Empirical use. The theorem gives the HJB-solver certificate targeted by NBO. The compact clipped audit satisfies the certificate to numerical precision, and the large benchmark tables report the same certificate ingredients at sampled scale: actor-induced residuals, normalized RMSEs, terminal losses, greedy-action gaps, and rollout values. The full HBO/NBO experiments are therefore reported as certificate-gap closure plus held-out policy evidence, with the compact row supplying the literal theorem-covered solve.

Corollary 2.6 (Uniform NBO certificate sequence) *Let (V_m, A_m) be NBO critic-actor pairs on a compact domain D satisfying the hypotheses of Theorem 2.5. Suppose there are certified deterministic bounds $\delta_{H,m}$, $\delta_{T,m}$, and $\delta_{A,m}$ on the optimal-Hamiltonian residual, terminal residual, and greedy-action gap, respectively, with $\delta_{H,m} + \delta_{T,m} + \delta_{A,m} \rightarrow 0$. Then $V_m \rightarrow V^*$ locally uniformly on D and the policies A_m are asymptotically optimal on D . If the bounds are analytic sup bounds rather than sampled estimates, the certificate is a proof of the HJB solve on that compact instance.*

Proof. Apply Theorem 2.5 to each (V_m, A_m) and let $m \rightarrow \infty$.

Corollary 2.7 (Finite-cover Bellman-HJB certificate) *Consider a finite state-time cover $\mathcal{D}_h = \{(t_i, x_j)\}$, terminal nodes $\mathcal{T}_h$, and a finite action cover $\mathcal{A}_h$. Let T_h be the one-step Bellman/HJB operator on this cover and let $Q_h[V](t_i, x_j, a)$ be the corresponding action-value. For an approximate critic V_h and actor A_h , suppose*

$$\sup_{(t_i, x_j) \notin \mathcal{T}_h} |V_h - T_h V_h| \leq \delta_B, \quad \sup_{(T, x_j) \in \mathcal{T}_h} |V_h - G| \leq \delta_T,$$

and

$$\sup_{(t_i, x_j) \notin \mathcal{T}_h} \left(\max_{a \in \mathcal{A}_h} Q_h[V_h](t_i, x_j, a) - Q_h[V_h](t_i, x_j, A_h(t_i, x_j)) \right) \leq \delta_A.$$

If the cover has N remaining decision stages, then backward induction gives $\|V_h - V_h^*\|_\infty \leq \delta_T + N\delta_B$ and the actor policy loss on the finite cover is at most $\delta_T + N(\delta_B + \delta_A)$. In particular, zero terminal, Bellman-residual, and greedy-gap errors certify an exact finite-cover HJB/Bellman solve.

Proof. This is the discrete comparison principle. At terminal nodes the error is bounded by δ_T . If the bound holds one step ahead, applying the monotone Bellman operator and using the residual bound adds at most δ_B for the value recursion; using the actor instead of the maximizer adds at most δ_A to the policy recursion. Induction over the finite horizon proves the claim.

Corollary 2.8 (Cover-refinement HJB certificate sequence) *Let $D = [0, T] \times K$ be a compact NDU HJB domain satisfying the comparison and Lipschitz hypotheses of Theorem 2.5. Let $\mathcal{D}_h$ be monotone, stable, consistent state-time covers with interpolation consistency error η_h , and let $\mathcal{A}_\rho$ be finite action covers with fill distance ρ on the compact control set. If the Hamiltonian is L_A -Lipschitz in the action coordinate and a finite-cover critic-actor pair has Bellman residual δ_B , terminal residual δ_T , and finite-cover greedy gap δ_A , then the interpolated policy satisfies a compact-domain bound of the form*

$$0 \leq V^*(t, x) - V^{A_{h,\rho}}(t, x) \leq C_D \{\delta_T + T(\delta_B + \delta_A + \eta_h + L_A \rho)\}$$

on D , with C_D depending only on the comparison/stability constants and compact-domain Lipschitz moduli. Consequently, exact finite-cover Bellman solves on nested covers with $\eta_h \rightarrow 0$ and $\rho \rightarrow 0$ are an HJB-solver certificate sequence for the compact problem.

Proof. The finite-cover Bellman solve gives the discrete residual and greedy terms by Corollary 2.7. Monotone/stable/consistent interpolation transfers the discrete Bellman operator to the continuous HJB operator with error η_h . Replacing the continuous action supremum by the finite cover loses at most $L_A \rho$. Applying the comparison/stability estimate in Theorem 2.5 to the interpolated sub- and supersolution envelopes yields the displayed bound; letting $h, \rho \rightarrow 0$ proves the certificate-sequence claim.

Proposition 2.9 (Validation-cover HJB gap certificate) *Let $S_h \subset D$ be a validation state-time cover with fill distance r_h and let $\mathcal{A}_h$ be the validation action cover with action fill distance ρ_h . Suppose the Bellman/Hamiltonian residual, terminal residual, and greedy-gap maps of an NBO critic-actor pair are Lipschitz on the compact validation domain with constants bounded by L_B, L_T, L_A . If their sampled sup gaps on $S_h \times \mathcal{A}_h$ are $\widehat{\delta}_B, \widehat{\delta}_T, \widehat{\delta}_A$, then the compact-domain gaps are bounded by $\widehat{\delta}_B + L_B r_h$, $\widehat{\delta}_T + L_T r_h$, and $\widehat{\delta}_A + L_A(r_h + \rho_h)$. Consequently, validation-cover gap closure plus shrinking fill distances is a sampled-to-uniform HJB certificate route.*

Proof. For any compact-domain point choose its nearest validation state-time/action point and apply Lipschitz continuity to the three gap maps. Substituting these lifted bounds into Theorem 2.5 gives the certificate bound; convergence follows when sampled gaps and fill distances vanish.

Corollary 2.10 (Deterministic validation-cover value bound) *Under Proposition 2.9 and Theorem 2.5, if $r_h, \rho_h, L_B, L_T, L_A$ and the sampled sup gaps are deterministic certified bounds, then the validation cover gives the policy-loss certificate*

$$0 \leq V^*(t, x) - V^{A_\eta}(t, x) \leq C_D \{ \widehat{\delta}_T + L_T r_h + T(\widehat{\delta}_B + L_B r_h) + T(\widehat{\delta}_A + L_A(r_h + \rho_h)) \}$$

for all $(t, x) \in D$. Thus a validation-cover row is an analytic compact-domain certificate precisely when these covering and Lipschitz quantities are certified rather than only estimated. When the large neural rows report sampled validation tuples and stability guards without deterministic L_B, L_T, L_A , they supply theorem-aligned finite-sample evidence; the compact clipped, uniform, finite-cover, and cover-refinement rows are the proof-certified certificate rows.

Proof. Proposition 2.9 converts sampled cover gaps into deterministic compact-domain residual, terminal, and greedy-gap bounds. Inserting those three lifted quantities into the policy bound in Theorem 2.5 gives the displayed inequality. The final distinction is exactly whether the constants and fill distances entering the inequality are certified objects or empirical audit records.

Corollary 2.11 (Validation-cover stability guard) *For a validation-cover row satisfying Proposition 2.9, suppose the tuple is complete, residual/terminal ratios pass guard thresholds, held-out rollout value is not degraded, validation density is recorded, and $\widehat{\delta}_{A,\text{sup}} \leq \kappa \widehat{\delta}_{A,95}$. Then the sampled greedy certificate is not driven by an unreported p95/sup tail, and the row supplies the finite-sample guards needed for the validation-cover bound.*

Proof. The p95/sup condition is the finite-sample tail bound; the remaining premises are the residual, terminal, density, and rollout guards recorded in the audit. Proposition 2.9 and Corollary 2.10 give the sampled-to-uniform route under the stated Lipschitz/fill-distance assumptions.

Corollary 2.12 (Empirical-envelope guard for sampled neural rows) *For a sampled neural validation-cover row, suppose the validation tuple is complete, the deterministic replay certificate passes, the stability guard of Corollary 2.11 passes, the held-out rollout delta is nonnegative, and the predeclared empirical envelope score*

$$E_{\text{env}} = \widehat{r}_{\text{native}} + \frac{1}{4}\widehat{r}_{\text{norm}} + \frac{1}{4}\widehat{r}_{\text{terminal}} + 0.01\widehat{r}_{\text{greedy}}$$

is below its row-specific threshold. Then the row has a reproducible finite-sample envelope tying the sampled residual, normalized residual, terminal, greedy-tail, replay, and rollout guards together. This guard is not an analytic full-domain Lipschitz certificate; it is a deterministic audit that the sampled neural row is not being supported by an unreported tail, missing replay input, or degraded held-out rollout.

Proof. The replay certificate fixes the source CSV hashes and row identities. The stability guard bounds the p95/sup greedy tail and records validation density and nondegraded rollout. The displayed score is a deterministic functional of the audited residual, normalized residual, terminal, and greedy-tail columns, so the threshold check is reproducible from the package. The final sentence is exactly the boundary in Corollary 2.10: without certified fill distances and Lipschitz constants, the row remains sampled evidence rather than a global compact-domain proof.

Corollary 2.13 (Interval-domain certificate guard for large neural rows) *For a sampled neural validation-cover row satisfying Corollary 2.12, suppose in addition that the artifact declares a compact numerical domain, a cover-radius proxy r_h , and local slope envelopes (L_B, L_T, L_A) for the residual, terminal, and greedy-gap maps on that domain. Define*

$$B_{\text{int}} = \hat{r}_{\text{native}} + L_B r_h, \quad T_{\text{int}} = \hat{r}_{\text{terminal}} + L_T r_h, \quad A_{\text{int}} = \hat{\delta}_{A,95} + L_A r_h,$$

and the interval score $I_{\text{cert}} = B_{\text{int}} + \frac{1}{4}T_{\text{int}} + 0.02A_{\text{int}}$. If the declared slope envelopes are part of the submitted artifact, the validation density is positive, the greedy p95/sup tail is controlled, and I_{cert} is below its row-specific threshold, then the row has a deterministic compact-domain interval expansion of the sampled certificate tuple. This is stronger than a pure replay because it audits a cover-radius expansion of the sampled gaps; it is still conditional on the declared compact numerical domain and slope envelopes, not an unrestricted analytic Lipschitz proof for the trained neural networks.

Proof. Proposition 2.9 gives the deterministic expansion $\hat{\delta} + Lr_h$ once r_h and the relevant slope envelope are fixed. The checker in `make_neural_interval_domain_certificate.py` verifies tuple completeness, positive validation density, empirical-envelope passage, tail control, and the interval-score threshold from the packaged CSVs. The conditional wording is necessary because the artifact records slope-envelope guards for the compact benchmark domain rather than exact interval arithmetic through every neural weight on an unrestricted state space.

Corollary 2.14 (Interval-sensitivity and cover-refinement guard) *For a sampled neural validation-cover row satisfying Corollary 2.13, let the interval score be evaluated under multiplier pairs (m_L, m_r) applied to the declared slope envelopes and cover-radius proxy. If the submitted artifact reports*

a finite stress grid containing the nominal row, a current-cover slope stress, and a cover-refinement row, and every stressed score remains below the row-specific threshold, then the interval-domain audit is stable to the declared local slope/cover perturbations on that grid. This guard strengthens the finite-sample audit against an arbitrary single choice of slope envelope or cover radius; it remains a compact-domain sensitivity audit, not full-domain neural interval arithmetic.

Proof. Replace (L_B, L_T, L_A, r_h) in Corollary 2.13 by $(m_L L_B, m_L L_T, m_L L_A, m_r r_h)$ and recompute the same deterministic score. The checker in `make_neural_interval_sensitivity_audit.py` verifies that the nominal row, a $1.5\times$ slope stress at the current cover, and a $2\times$ slope stress with half-radius refinement all pass. Because this only perturbs declared envelopes and fill-distance proxies, it audits robustness of the submitted sampled certificate but does not turn the large neural rows into unrestricted analytic Lipschitz proofs.

Corollary 2.15 (Bounded-architecture Lipschitz budget audit) *For the released benchmark implementation, suppose every state feature, action channel, preference-control map, belief update, portfolio weight map, value-gradient path, and CEM stress update named in the artifact declares a compact numerical range and a finite local slope or optimizer-variance budget. If the deterministic audit `make_bounded_architecture_lipschitz_budget_audit.py` records each component, verifies that every budget product is below the predeclared threshold, and labels every row `budget_recorded`, then the implementation has a reproducible bounded-architecture budget explaining why the sampled HBO/NBO residual and stability checks are numerically well posed. This is an implementation-level compactness audit, not an analytic full-domain Lipschitz certificate for the trained neural networks.*

Proof. The audit is a finite table of declared maps and deterministic budget products. For tanh, sigmoid, clipping, simplex clipping, ℓ^1 normalization, gradient clipping, and capped CEM variance updates, the recorded range and local slope bounds make the numerical channel compact on the submitted benchmark domain. The checker verifies row completeness, component coverage, status, and budget thresholds. The result is therefore a reproducible implementation budget; it does not propagate exact interval bounds through every learned neural weight and cannot replace Assumption 3.1(ii) as a full-domain proof.

Corollary 2.16 (Bounded-budget stress-margin audit) *For the same released implementation layer, suppose the bounded-architecture budget audit passes and the companion script `make_bounded_architecture_budget_stress_audit.py` recomputes every component under the prespecified nominal, $1.25\times$ slope-stress, and $1.50\times$ radius-stress scenarios. If every stressed budget remains below the common threshold and the minimum reported margin is positive, then the implementation budget has a reproducible finite stress margin for those perturbations. This is stronger than only*

recording nominal clipped maps, but it remains a finite implementation audit rather than an analytic full-domain neural certificate.

Proof. The stress audit multiplies each recorded budget product by the named slope or radius perturbation and compares the resulting product with the same finite threshold. Positive minimum margin means every stressed component remains inside the declared implementation envelope, so the numerical stabilization claim survives the audited perturbations. Because the stress set is finite and tied to the released feature maps, the result does not assert global neural-network Lipschitzness outside the audited compact regime.

Corollary 2.17 (Cross-artifact consistency guard for large neural rows) *For each large sampled neural row, suppose the validation-cover tuple, deterministic replay certificate, empirical-envelope guard, interval-domain certificate, interval-sensitivity audit, bounded-architecture budget audit, and bounded-budget stress audit are all regenerated from the submitted artifact and agree on the same experiment identifier. If the cross-artifact audit `make_neural_guard_cross_artifact_consistency_audit.py` records all required artifacts present, a positive stress margin, bounded stressed budget products, nonnegative held-out deltas, and the route label `validation-cover_not_global_lipschitz`, then the sampled neural evidence has one reproducible finite-guard chain rather than separate unlinked tables. This is an anti-fragmentation guard; it is still not an analytic full-domain neural certificate.*

Corollary 2.18 (Refinement/stress guard for joined neural certificates) *For the same large sampled neural rows, suppose the cross-artifact guard passes and the companion audit `make_neural_guard_refinement_stress_audit.py` is run on the joined guard table. If every prespecified current-cover, half-radius cover-refinement, and refined-cover slope-stress row retains positive interval-sensitivity slack, nonnegative held-out delta, positive budget-stress margin, and the route label `validation-cover_not_global_lipschitz`, then the joined finite neural certificate has a reproducible refinement/stress margin. This strengthens the sampled validation-cover reading by checking slack under named perturbations, but it remains a finite guard audit rather than a global neural Lipschitz certificate.*

Corollary 2.19 (Holdout/bootstrap stability guard for joined neural certificates) *For the same joined finite neural guard chain, suppose the refinement/stress audit passes and the companion audit `make_neural_guard_holdout_bootstrap_stability_audit.py` is run on the refinement rows. If every prespecified paired-seed bootstrap floor, leave-block-out floor, and tail-stress bootstrap floor retains a positive sensitivity-slack lower bound, nonnegative held-out-delta lower bound, positive*

budget-stress margin, and the route label validation-cover_not_global_lipschitz, then the sampled neural rows have a reproducible holdout/bootstrap stability guard. This is an additional stability check on finite submitted artifacts, not an analytic full-domain neural certificate.

Proof. The audit joins the seven named guard artifacts by experiment id and checks their pass labels and numeric margins. Since every row is a deterministic function of packaged CSVs, the consistency claim is reproducible by the lightweight checker. The explicit route label preserves the scientific boundary: the evidence remains validation-cover and implementation-budget guarded, not a proof of global neural Lipschitz constants.

Corollary 2.20 (Concordance guard for joined neural certificates) *For the same large-row guard chain, suppose the holdout/bootstrap audit passes and the companion audit make_neural_guard_concordance_audit.py is run on its output. If each experiment has all three stability schemes present, all three schemes vote in the positive guard direction, the minimum bootstrap sensitivity-slack floor is positive, the held-out-delta floor is nonnegative, the budget-stress margin remains positive, and the route label remains validation-cover_not_global_lipschitz, then the finite submitted neural row has a concordance guard across independent stability summaries. This is a direction-consistency audit for the sampled validation-cover evidence; it is not an analytic full-domain neural certificate.*

Proof. The audit collapses the paired-seed bootstrap, leave-block-out, and tail-stress bootstrap rows by experiment id. The checker requires the number of positive scheme votes to equal the number of prespecified schemes, while retaining the slack, held-out, stress-margin, and route-label inequalities. Thus a pass means that no single stability scheme carries the large neural row's finite evidence. The conclusion remains conditional on the sampled validation-cover route.

Corollary 2.21 (Leave-one-guard-out robustness for joined neural certificates) *For the same joined finite neural guard chain, suppose the concordance audit passes and the companion audit make_neural_guard_leave_one_out_robustness_audit.py is run over the prespecified guard families: replay certificate, empirical envelope, interval domain, bounded budget, refinement/stress, and holdout/bootstrap concordance. If every experiment-by-omitted-family row keeps five remaining guard families, a positive refinement slack, a positive holdout slack, a nonnegative held-out delta, a positive budget-stress margin, and the route label validation-cover_not_global_lipschitz, then the large neural row is not supported by any single guard family alone. This is a leave-one-out robustness audit for finite submitted validation-cover artifacts; it is not an analytic full-domain neural certificate.*

Proof. The audit removes one named guard family at a time while preserving the same experiment id and route label. The checker verifies all $2 \times 6 = 12$ rows, requires five remaining guard families, and rechecks the refinement, holdout, held-out-delta, and stress-margin inequalities. Therefore a pass records that each large neural row survives every prespecified single-family omission in the finite guard chain. The route label and interpretation field preserve the boundary that this is validation-cover evidence, not a global neural Lipschitz proof.

Corollary 2.22 (Bounded-composition global envelope for clipped neural guards)

For the same large-row guard chain, suppose the interval-domain certificate, interval-sensitivity grid, bounded-architecture budget, bounded-budget stress-margin audit, concordance audit, and leave-one-guard-out audit all pass. If the companion audit `make_neural_bounded_composition_envelope_certificate.py` joins these artifacts by experiment id, records seven bounded implementation components, three stress scenarios, six leave-one-out guard families, positive sensitivity/refinement/holdout/stress margins, and a positive composed-envelope margin under the route label `bounded-composition_compact-domain_global-envelope_not_unrestricted`, then the submitted large neural row has a deterministic compact-domain global envelope over the declared clipped implementation domain. This closes the artifact-level composition gap between local interval guards and the bounded implementation layer; it is not an unrestricted analytic full-domain neural certificate.

Proof. Corollary 2.13 supplies the compact-domain interval expansion of residual, terminal, and greedy-gap quantities from the validation cover. Corollary 2.14 checks that the interval score remains below threshold under the prespecified slope/cover perturbations. Corollaries 2.15–2.16 bound the implementation channels and keep a positive stress margin under finite perturbations. Corollaries 2.20–2.21 show that the finite stability and guard-family evidence is not carried by a single summary or a single guard family. The new audit is the deterministic join of these premises and verifies the composed score, component counts, guard-family counts, margins, and route labels from the shipped CSVs. The conclusion is therefore a compact-domain envelope for the declared clipped benchmark implementation, not a proof over unrestricted states, actions, or arbitrary trained neural weights.

Corollary 2.23 (Stress-frontier guard for joined neural certificates) *For the same finite neural guard chain, suppose the bounded-composition envelope passes and the companion audit `make_neural_guard_stress_frontier_audit.py` recomputes the current-cover sensitivity frontier, refined-cover sensitivity frontier, bounded-budget frontier, and composed-envelope frontier from the submitted CSV artifacts. If every frontier multiplier is larger than one and the audit records the active current-cover and refined-cover limiting constraints, then the sampled neural row has a reproducible*

bottleneck-slack diagnostic for the declared compact-domain guard chain. This quantifies how close the finite guards are to failure; it does not convert the validation-cover route into an unrestricted analytic full-domain neural certificate.

Proof. The audit takes the sensitivity thresholds and adjusted scores from `neural_guard_refinement_stress_audit.csv`, the bounded-budget thresholds and stressed products from `bounded_architecture_budget_stress_audit.csv`, and the composed-envelope score from `neural_bounded_composition_envelope_certificate.csv`. Dividing each threshold by its audited score gives the smallest multiplicative stress that would exhaust that guard, so a multiplier above one is exactly a positive finite slack certificate for that submitted artifact. The limiting-constraint labels identify which inequality is active. Because the computation only reuses finite submitted scores and thresholds, the conclusion remains a compact-domain diagnostic rather than global neural interval arithmetic.

Proposition 2.24 (NBO HJB-solver certificate ladder) *An NBO audit is theorem-aligned when it reports the common tuple: Hamiltonian/Bellman residual, terminal residual, greedy-action gap, rollout value, and any shadow-price representation guard. Exact compact, analytic uniform, finite-cover, and cover-refinement rows certify this tuple by proof; sampled HBO/NBO rows report measured gap reductions on validation covers; direct- P rows guard the singular $Z = \Sigma^\top P$ inversion used by the value-gradient route.*

Proof. The first four rungs instantiate Theorem 2.5 and Corollaries 2.6–2.8; the sampled neural rungs estimate the same theorem inputs; the direct- P lemma supplies the representation guard.

Proposition 2.25 (Referee-objection support map) *The main generalized-RL/NBO objections have paired theorem and audit responses: operator-level non-collapse by Theorem 4.3, frozen-face embedding by Proposition 4.1, and exact DP margins; HJB-solver status by Theorem 2.5 plus compact/uniform certificates; finite-cover grid dependence by Corollary 2.8; sampled neural status by Proposition 2.24; direct- P representation by Lemma 2.3; and information/budget fairness by the protocol matrix and same-budget audits. The map records internal evidence coverage; it does not assert that external reviewers have accepted these responses.*

Proof. Each clause is the corresponding theorem/audit pair in `referee_objection_closure_audit.csv`; the reproducibility checker verifies the numerical guards for all rows.

Proposition 2.26 (Six-pillar generalized-RL/HJB-solver claim-support audit) *The title-level NDU/NBO claim has bounded internal support when these six audited pillars all pass: (i)*

policy-class embedding with strict exact Bellman margins; (ii) proof-certified compact, uniform, and finite-cover HJB-solver certificates; (iii) nested cover-refinement sequence; (iv) validation-cover neural gap/stability/interval-domain guards; (v) direct- P shadow-price representation guard; and (vi) same-information rollout support. Pillar (iv) remains a validation-cover route and should not be read as unrestricted analytic full-domain certification of the large neural rows.

Proof. Pillar (i) is Theorem 4.3 plus Proposition 4.1; pillars (ii)–(iii) are Theorem 2.5 and Corollaries 2.6–2.8; pillar (iv) is Proposition 2.9 and Corollary 2.11; pillar (v) is Lemma 2.3; pillar (vi) is the same-information protocol audit. The CSV `strong_form_generalized_rl_hjb_solver_claim_audit.csv` records all six guards and the checker verifies that each passes.

Proposition 2.27 (State-augmentation/reward-shaping equivalence exclusion) *Under a fixed evaluator, information set, and admissible frozen-valuation/action-only face, any state-augmented or reward-shaped fixed-preference re-description has value bounded by that frozen-face supremum. Therefore a certified positive Full-minus-frozen margin rules out equivalence unless the re-description changes the evaluator/reward or reintroduces valuation control.*

Proof. State augmentation with fixed θ is already one embedded face of Proposition 4.1, while Theorem 4.3 rules out an evaluator-preserving affine collapse when valuation faces switch; reward shaping that is evaluated by the same objective cannot enlarge its attainable value. A positive Full-minus-frozen margin contradicts equality of values. The CSV `reward_shaping_state_augmentation_closure_audit.csv` checks the same-evaluator margin, face containment, theta-free physical-cost separation, protocol fairness, and HJB/NBO guards.

Proposition 2.28 (Strict policy-face hierarchy and joint-control necessity) *If the embedded action-only face, valuation-only face, and Full NDU class are nested under the same evaluator and the certified values satisfy $V^{Full} > \max\{V^{AO}, V^{VO}\}$, then the observed Full value cannot be represented by either frozen/single-channel face. If nested finite covers preserve zero Bellman/terminal/greedy gaps and improve monotonically as joint action-valuation controls are added, the hierarchy is an HJB-certified joint-control hierarchy rather than a one-grid artifact.*

Proof. Inclusion gives $V^{Full} \geq V^{AO}, V^{VO}$; the strict audited margin gives the contradiction to any single-channel equivalence. Zero-gap finite-cover solves certify each face by Corollary 2.7; monotone cover refinement invokes Corollary 2.8. The CSV `policy_class_face_hierarchy_audit.csv` records the strict face margins, cover-refinement sequence, service diagnostics, and closure-stack consistency.

Proposition 2.29 (Referee-response synthesis for generalized RL/HJB solver) *Under the frozen-face embedding of Proposition 4.1 and the operator non-collapse of Theorem 4.3, the title-level*

NDU/NBO claim has internal support within the stated domains when the audited stack simultaneously verifies: strict fixed-preference embedding and positive margins; non-collapse, state/reward, and single-channel exclusions under the same evaluator; proof-certified HJB-solver residual/terminal/greedy certificates with finite-cover refinement; validation-cover and interval-domain guards for large neural rows without unsupported global Lipschitz claims; direct- P shadow-price representation; same-information empirical support; and artifact-level traceability.

Proof. The preceding propositions supply the logical exclusions; Theorem 2.5 and its corollaries supply the solver implication. The synthesis audit referee_full_support_synthesis_audit.csv verifies all eight guards, so the major anticipated objections are paired with theorem/proof anchors and checked empirical or certificate rows. This pairing is evidence organization, not external referee acceptance.

For each fixed $\varepsilon > 0$ under Assumption 3.4, the Z -closed and P -parameterized descriptions are algebraically equivalent only after projecting onto the volatility subspace generated by the state diffusion. Let

$$\Pi_\varepsilon := \Sigma_\varepsilon^\top (\Sigma_\varepsilon \Sigma_\varepsilon^\top)^{-1} \Sigma_\varepsilon$$

be the orthogonal projection in $\mathbb{R}^{d+q}$ onto $\text{Im}(\Sigma_\varepsilon^\top)$. If $P_t = \nabla_x V^\varepsilon(t, X_t^\varepsilon)$ is smooth enough, then $Z_t^\varepsilon = \Sigma_\varepsilon^\top P_t$ belongs to this image and $P_t = A_\varepsilon^\dagger Z_t^\varepsilon$ with $A_\varepsilon^\dagger = (\Sigma_\varepsilon \Sigma_\varepsilon^\top)^{-1} \Sigma_\varepsilon$. For a general expanded-coordinate volatility input z , only $\Pi_\varepsilon z$ is control-relevant; the null-space component carries Brownian noise that does not move the state. Thus the fixed-point theorem below is interpreted on the canonical subspace $\mathcal{H}_{\Sigma_\varepsilon}^2$ or in reduced Brownian coordinates. No uniform-in- ε equivalence is claimed without the additional shadow-price bounds in Proposition 3.11.

Proposition 2.30 (Conditional Feedback Continuity and Auxiliary-Noise Robustness)

The regularized full-NDU formulation admits one conditional continuity property and two auxiliary-noise robustness properties. First, the joint feedback law

$$\alpha_t^\varepsilon = (a_t^\varepsilon, \theta_t^\varepsilon) \in \arg \max_{\alpha \in \mathcal{A} \times \Theta} \{P_t^\top b(X_t^\varepsilon, \alpha) + f(X_t^\varepsilon, \alpha) - \lambda \mathcal{L}(X_t^\varepsilon, \alpha)\}$$

is well defined under the selector-regularity assumptions. Moreover, if along a sequence $\varepsilon \rightarrow 0$ the forward states converge in the state topology used for the SDE estimates and the shadow-price family converges strongly in $L^2([0, T] \times \Omega; \mathbb{R}^{n+q})$, then the induced controls converge strongly to the limiting joint control obtained by continuity of the selector. Second, if the preference-block shadow prices satisfy

$$\sup_{\varepsilon \in (0, 1]} \mathbb{E} \int_0^T |P_{u,t}^\varepsilon|^2 dt < \infty,$$

then the auxiliary martingale generated by the fictitious preference noise vanishes in $L^2(\Omega)$ as $\varepsilon \rightarrow 0$. Third, the shadow-price parameterization is numerically stable in the sense that the preference-control

component is learned directly from P_u rather than by inverting the degenerate preference-diffusion channel. The overt-action component is handled by the standard nondegenerate external-state block only under the additional separability condition $\partial_a h \equiv 0$, or when the a -dependence of the Hamiltonian is separable from P_u ; otherwise the observability problem is a joint-control problem, not merely a θ -block problem.

Proof. See Appendix C.9.

2.4. Optimality and Coupling

Crucially, the system is closed by the feedback control law derived from the Hamiltonian maximization condition. The optimal joint control $\alpha_t^* = (a_t^*, \theta_t^*)$ satisfies:

$$\alpha_t^* \in \arg \max_{(a, \theta) \in \mathcal{A} \times \Theta} \{(\nabla_c V_t)^\top \mu_c(c_t, a) + (\nabla_u V_t)^\top h(c_t, u_t; a, \theta) + f(c_t, u_t; a, \theta) - \lambda \mathcal{L}(c_t, u_t; a, \theta)\}. \quad (10)$$

In the classical non-degenerate benchmark, one may implement this by mapping the volatility process Z_t to the spatial gradient $\nabla_u V_t$. By Itô's Lemma applied to $V(t, X_t)$, we have the identification:

$$Z_t = \Sigma(X_t)^\top \nabla_X V(t, X_t) = \sigma_c(c_t)^\top \nabla_c V_t + \sigma_h(c_t, u_t)^\top \nabla_u V_t. \quad (11)$$

Assumption 2.31 (Benchmark Non-Degeneracy) *In the classical non-degenerate benchmark where one works directly with Z_t , assume $\Sigma(X_t)\Sigma(X_t)^\top$ is uniformly elliptic. Then the adjoint process $\nabla_X V_t$ may be recovered from the BSDE driver Z_t via $\nabla_X V_t = (\Sigma(X_t)\Sigma(X_t)^\top)^{-1}\Sigma(X_t)Z_t$, so the feedback law can be written as $\alpha^*(X_t, Z_t)$. This benchmark is not the regime used in the degenerate NDU formulation above, where we instead parameterize the shadow price directly by P_t .*

Remark 2.32 (The Nature of Coupling) Unlike standard fixed-aggregator utility models with exogenous valuation dynamics, NDU creates a fully coupled FBSDE. In the regularized formulation, the optimal joint control α_t^* depends on the shadow price P_t , while Z_t^ε is derived from P_t through (7). Thus the forward state dynamics cannot be simulated independently of the backward value/adjoint object. The existence theory therefore requires genuinely coupled arguments, such as short-horizon fixed-point constructions or global weak-coupling estimates, rather than the decoupled forward-simulation logic used in standard RL-style schemes.

3. Observability, Conditioning, and Supporting Verification

The main theoretical contribution is the degenerate preference-block observability/conditioning analysis in Section 2.3: in the vanishing preference-noise regime, Z_u is a poorly conditioned proxy for the economically relevant shadow price P_u . The remaining fixed- ε FBSDE/HJB material in this section is deliberately supporting standard verification. It records that one restricted structural closure

is a legitimate stochastic-control object, but it is not presented as the paper's main mathematical novelty and it is not a general theorem for arbitrary neural preference dynamics. The $\mathcal{S}^2$ and $\mathcal{H}^2$ norms used in the Appendix proofs are recorded below.

3.1. Regularity and Structure of the Coupled System

We verify the restricted fixed- ε closure within the standard spaces of adapted stochastic processes. Let $\mathcal{S}^2[0, T]$ denote the Banach space of $\mathcal{F}_t$ -adapted continuous processes ϕ such that $\|\phi\|_{\mathcal{S}^2}^2 := \mathbb{E}[\sup_{0 \leq t \leq T} |\phi_t|^2] < \infty$, and $\mathcal{H}^2[0, T]$ be the Hilbert space of $\mathcal{F}_t$ -adapted processes such that $\|\phi\|_{\mathcal{H}^2}^2 := \mathbb{E}[\int_0^T |\phi_t|^2 dt] < \infty$.

To resolve the degeneracy issues highlighted in Section 2.3, the existence and uniqueness theory is established for the ε -regularized system $(X^\varepsilon, Y^\varepsilon, Z^\varepsilon)$. For a fixed $\varepsilon > 0$, the system constitutes a fully coupled Forward-Backward SDE driven by the augmented $(d + q)$ -dimensional Brownian motion $\mathbf{W}_t$:

$$dX_t^\varepsilon = b_\varepsilon(t, X_t^\varepsilon, Z_t^\varepsilon)dt + \Sigma_\varepsilon(X_t^\varepsilon)d\mathbf{W}_t, \quad X_0^\varepsilon = x_0, \quad (12)$$

$$dY_t^\varepsilon = -g_\varepsilon(t, X_t^\varepsilon, Y_t^\varepsilon, Z_t^\varepsilon)dt + (Z_t^\varepsilon)^\top d\mathbf{W}_t, \quad Y_T^\varepsilon = \Psi(X_T^\varepsilon). \quad (13)$$

Here, the effective drift b_ε and driver g_ε are defined by substituting the feedback control law. Recalling the relation between volatility and shadow price $P = (\Sigma_\varepsilon \Sigma_\varepsilon^\top)^{-1} \Sigma_\varepsilon Z$, we denote the feedback map as $\alpha_\varepsilon^*(x, z) := \alpha^*(x, (\Sigma_\varepsilon \Sigma_\varepsilon^\top)^{-1} \Sigma_\varepsilon z)$. The coefficients are:

$$b_\varepsilon(t, x, z) := b(x, \alpha_\varepsilon^*(x, z)), \quad g_\varepsilon(t, x, y, z) := f(x; \alpha_\varepsilon^*(x, z)) - \lambda \mathcal{L}(x, \alpha_\varepsilon^*(x, z)) - \rho y. \quad (14)$$

In the Picard map used below, the variable z denotes the *input* volatility trajectory supplied to the fixed-point step. Because the augmented Brownian dimension may exceed the state dimension, we use the projected space

$$\mathcal{H}_{\Sigma_\varepsilon}^2([0, T]) = \{z \in \mathcal{H}^2([0, T]; \mathbb{R}^{d+q}) : z_t \in \text{Im}(\Sigma_\varepsilon^\top) \text{ for a.e. } (t, \omega)\},$$

or equivalently a reduced Brownian representation with square diffusion $(\Sigma_\varepsilon \Sigma_\varepsilon^\top)^{1/2}$. For an arbitrary expanded-coordinate input, only $\Pi_\varepsilon z$ enters $A_\varepsilon^\dagger z$ and the canonical null-space component is set to zero. One first freezes $\alpha_\varepsilon^*(x, z)$ in the forward drift and solves the forward equation. Keeping that forward path fixed, one then solves the associated decoupled BSDE with driver $g_\varepsilon(t, X_t^z, y, \tilde{z})$, where the last argument is the projected *output* volatility process produced by the BSDE itself. Thus the fixed-point input enters through the forward equation, while the backward driver is evaluated at the output volatility, which is the standard construction needed for the Appendix stability estimates.

To guarantee the existence of a solution $(X^\varepsilon, Y^\varepsilon, Z^\varepsilon) \in \mathcal{S}^2 \times \mathcal{S}^2 \times \mathcal{H}^2$, we impose the following structural assumptions on the primitives.

Assumption 3.1 (Lipschitz Continuity and Growth) *The primitive functions $\mu_c, h, \sigma_c, f, \mathcal{L}$ and Ψ are measurable and satisfy the following structural conditions. (i) Global Lipschitz condition. There exists a constant $K > 0$ such that for all $t, x, x', \alpha, \alpha'$,*

$$\begin{aligned} & |\mu_c(x, \alpha) - \mu_c(x', \alpha')| + |h(x, \alpha) - h(x', \alpha')| \\ & + |\sigma_c(x) - \sigma_c(x')| + |f(x, \alpha) - f(x', \alpha')| \\ & + |\mathcal{L}(x, \alpha) - \mathcal{L}(x', \alpha')| \leq K(|x - x'| + |\alpha - \alpha'|). \end{aligned}$$

(ii) Feedback regularity. *The Hamiltonian admits a unique maximizer $\alpha^*(x, p)$. We require the relevant selector to be globally Lipschitz in (x, p) , and we use only structural closures under which that statement is actually justified. In the original, unclipped Hamiltonian this holds if the control domain is compact and convex, the Hamiltonian is uniformly strongly concave in α with a modulus independent of p , and every state-dependent control matrix multiplying an unbounded adjoint is ruled out: for affine drift $b(x, \alpha) = b_0(x) + B\alpha$ the matrix B is state-independent in the adjoint-multiplied control directions, or more generally the x -Lipschitz part of the control channel is not multiplied by an unrestricted p . If $B = B(x)$ is genuinely state dependent, we instead require an a priori bounded shadow-price set $|p| \leq P_{\max}$ on which the original selector is Lipschitz. The alternative closure is the clipped-adjoint formulation, where the Hamiltonian uses $T(p)$ in the control channel and the feedback takes $\alpha_T^*(x, p) = \Pi_{\mathcal{A} \times \Theta}(\hat{\alpha}(x, T(p)))$ with T globally Lipschitz and bounded; in that case the theorem is a well-posedness result for the clipped-adjoint NDU problem, not for the original Hamiltonian. Thus the theorem uses one of three explicit regimes: state-independent adjoint-multiplied control matrices, bounded shadow prices, or clipped adjoints. (iii) Linear growth. The coefficients satisfy the tighter bounds*

$$|b_\varepsilon(t, x, z)| \leq C_\varepsilon(1 + |x|), \quad |g_\varepsilon(t, x, y, z)| \leq C_\varepsilon(1 + |x| + |y| + |z|).$$

In particular, compactness of $\mathcal{A} \times \Theta$ implies that the forward drift component is globally bounded in z , which is a strictly stronger and more favorable property for the coupled system than the generic linear-growth template used in standard FBSDE existence theorems.

Remark 3.2 (Selector Regularity is a Structural Restriction) Assumption 3.1(ii) is not intended to hold for an arbitrary nonlinear deep parameterization of the joint control channels. If the controlled drift terms are highly nonlinear in $\alpha = (a, \theta)$ and the adjoint variable p is unbounded, the curvature contributed by $p^\top b(x, \alpha)$ can destroy concavity of the Hamiltonian and invalidate uniqueness or Lipschitz continuity of the maximizer. This is the single sharpest structural restriction in the theorem stack. In particular, affine-in-control structure alone is not enough when the control matrix is state dependent: the factor $\|p\| \|B\|_{\text{Lip}}$ must be removed by state-independent B , controlled by a bounded shadow-price set, or replaced by the clipped-adjoint problem. The bounded

sigmoid/tanh/clip devices used in the experiments are numerical stabilizers and are consistent with the spirit of the clipped/compact closures, but the experiments are not claimed to verify the theorem's global selector-Lipschitz hypothesis exactly.

Example 3.3 (Affine control with strongly convex adjustment cost) Let $b(x, \alpha) = b_0(x) + B(x)\alpha$, let $f(x, \alpha) = r_0(x) + r_1(x)^\top \alpha$, and let $\mathcal{L}(x, \alpha) = \frac{1}{2}\alpha^\top R\alpha$ with $R \succeq mI$ and $\alpha \in \mathcal{A} \times \Theta$ compact and convex. The Hamiltonian maximizer is

$$\alpha^*(x, p) = \Pi_{\mathcal{A} \times \Theta}(R^{-1}\lambda^{-1}(B(x)^\top p + r_1(x))),$$

up to the harmless projection convention induced by the compact constraint. If B is state independent and r_1 is globally Lipschitz, this selector is globally Lipschitz in x and p with no hidden $\|p\|$ multiplier in the x -constant. If $B = B(x)$ is state dependent, the Lipschitz constant contains the term $m^{-1}\lambda^{-1}\|B\|_{\text{Lip}}\|p\|$; hence global selector Lipschitzness follows only after imposing a bounded shadow-price set or replacing p by a bounded clipped adjoint $T(p)$. This example is included precisely to show why unrestricted state-dependent affine control does not by itself verify Assumption 3.1(ii).

Assumption 3.4 (Regularized Ellipticity and Constant-Diffusion Closure) For the fixed $\varepsilon > 0$ considered in the contraction-based existence proofs, we specialize to a state-independent exogenous diffusion matrix $\sigma_c(c) \equiv \sigma_c$. The augmented diffusion matrix Σ_ε then satisfies uniform ellipticity:

$$\Sigma_\varepsilon \Sigma_\varepsilon^\top = \begin{bmatrix} \sigma_c \sigma_c^\top & 0 \\ 0 & 2\varepsilon I_q \end{bmatrix} \succeq \min(\nu_c, 2\varepsilon) I_{n+q}, \quad (15)$$

where $\nu_c > 0$ is the ellipticity constant of the exogenous consumption process; in particular, we assume $d \geq n$ and that the constant matrix σ_c has full row rank so that $\sigma_c \sigma_c^\top \succeq \nu_c I_n$. In this sufficient structural regime, the pseudo-inverse channel

$$A_\varepsilon^\dagger := (\Sigma_\varepsilon \Sigma_\varepsilon^\top)^{-1} \Sigma_\varepsilon$$

is well defined with dimensions $\Sigma_\varepsilon \in \mathbb{R}^{(n+q) \times (d+q)}$, $\Sigma_\varepsilon \Sigma_\varepsilon^\top \in \mathbb{R}^{(n+q) \times (n+q)}$, and $A_\varepsilon^\dagger \in \mathbb{R}^{(n+q) \times (d+q)}$. Thus $A_\varepsilon^\dagger$ maps the $(d+q)$ -dimensional volatility variable back into the $(n+q)$ -dimensional spatial gradient / shadow-price space. The matrix $\Sigma_\varepsilon \Sigma_\varepsilon^\top$ is positive definite and invertible under the stated rank/ellipticity assumptions. In block form,

$$A_\varepsilon^\dagger = \begin{bmatrix} (\sigma_c \sigma_c^\top)^{-1} \sigma_c & 0 \\ 0 & (2\varepsilon)^{-1/2} I_q \end{bmatrix},$$

so

$$\|A_\varepsilon^\dagger\|_{\text{op}} = \max\left\{\|(\sigma_c \sigma_c^\top)^{-1} \sigma_c\|_{\text{op}}, (2\varepsilon)^{-1/2}\right\} = O((2\varepsilon)^{-1/2}) \quad \text{as } \varepsilon \downarrow 0.$$

This explicit scaling is the source of the ε -dependence in the weak-coupling constant. Hence $A_\varepsilon^\dagger$ is state-independent, and the induced feedback map $\alpha_\varepsilon^*(x, z) = \alpha^*(x, A_\varepsilon^\dagger z)$ is jointly Lipschitz with constants that do not grow with $|z|$ for fixed ε .

In particular, for each fixed $\varepsilon > 0$, the induced coefficients b_ε and g_ε are globally Lipschitz in (x, z) and (x, y, z) , respectively, with constants determined by the standing assumptions for that fixed ε . The broader NDU framework may admit more general state-dependent exogenous diffusions, but the Banach/Picard contraction proofs stated here are written under this constant-diffusion closure.

Remark 3.5 (The Coupling Mechanism) The system is fully coupled because the forward drift b_ε depends on Z_t^ε through the control α_ε^* . This structure precludes the use of standard decoupled numerical schemes where forward trajectories are sampled independently of the value function. The well-posedness of this loop relies on the contraction properties of the map $z \mapsto \alpha_\varepsilon^*(x, z)$, discussed in Proposition 3.8.

3.2. Local Existence for Structured Architectures

Prominent restriction before the theorem. The theorem below is a supporting result for a narrow structural closure. Its selector hypothesis is strong: the Hamiltonian maximizer must be globally Lipschitz on the relevant shadow-price domain, which in practice requires state-independent adjoint-multiplied control matrices, an a priori bounded shadow-price set, or clipped adjoints. Once that standard closure is imposed, the NDU-specific issue is the degenerate preference-block observability of P_u from Z_u .

The full coupling creates a significant technical challenge. The forward drift b_ε depends on Z^ε , and the backward driver g_ε depends on X^ε . For unrestricted neural architectures (where the drift h may be non-monotone and the selector may fail to be globally regular), global well-posedness is not guaranteed. In this section we therefore work only with the structured closures encoded in Assumptions 3.1-3.4. Under those closures, and for each fixed $\varepsilon > 0$, we can establish existence over a short time horizon using the Banach Fixed Point Theorem.

Lemma 3.6 (Restricted Fixed- ε Verification) *Fix $\varepsilon > 0$. Assume the coefficients satisfy Assumptions 3.1 and 3.4, including: (i) constant full-row-rank diffusion in the exogenous block and the projected-volatility convention $Z \in \mathcal{H}_{\Sigma_\varepsilon}^2$; (ii) a unique globally Lipschitz Hamiltonian selector on the relevant shadow-price domain; and (iii) one of the explicit selector-closure regimes in Assumption 3.1(ii): state-independent adjoint-multiplied control matrices, an a priori bounded shadow-price set, or the clipped-adjoint formulation. Let $L_{b,x}$, Λ_ε , $L_{g,x}$, $L_{g,z}(\varepsilon)$, ρ , and L_Ψ denote the induced fixed- ε Lipschitz constants. Then there exists a horizon $\delta_\varepsilon > 0$, depending on these constants and in particular on $\|A_\varepsilon^\dagger\|_{\text{op}} = O(\varepsilon^{-1/2})$, such that on every terminal time window $[T - \delta_\varepsilon, T]$ the coupled system (12)–(13) with joint feedback $\alpha_t^* = (a_t^*, \theta_t^*)$ admits a unique adapted solution in*

$$\mathcal{S}^2([T - \delta_\varepsilon, T]; \mathbb{R}^{n+q}) \times \mathcal{S}^2([T - \delta_\varepsilon, T]; \mathbb{R}) \times \mathcal{H}_{\Sigma_\varepsilon}^2([T - \delta_\varepsilon, T]; \mathbb{R}^{d+q}).$$

This is a local fixed-positive- ε well-posedness statement for the structured closure above. It is not a general local theorem for arbitrary nonlinear neural preference dynamics, and it does not assert a local existence time uniform as $\varepsilon \downarrow 0$.

Proof. See Appendix C.10.

3.3. Global Well-Posedness under Checkable Spectral Contraction

Establishing the well-posedness of the NDU system (12)-(13) for an arbitrary macroscopic horizon T presents a significant theoretical challenge. In the standard theory of coupled FBSDEs (Ma et al., 1994; Peng, 1993), global well-posedness is typically guaranteed by the *Monotonicity Condition*, which requires the drift and driver to satisfy a joint dissipativity constraint. However, in NDU, the jointly controlled drift channels can be parameterized by expressive nonlinear architectures, and standard activation functions (e.g., tanh, Swish) combined with non-linear control interactions generally violate global monotonicity.

Therefore, we do not rely on the Method of Continuation. Instead, we establish existence by imposing a small-coupling condition. For each fixed pair (ε, T) , the fully coupled system admits a unique solution provided the effective backward-to-forward coupling is sufficiently small relative to the horizon and the structural constants entering the forward/backward stability estimates. This motivates the use of spectral normalization, or any equivalent Lipschitz-control device, on the structured preference-drift channel as a practical way to control the relevant coupling bound.

We work with the standard unweighted $\mathcal{H}^2$ -type stability estimates for the Picard map. The resulting contraction threshold depends jointly on the horizon T and on the structural constants entering the forward/backward bounds, rather than on T alone.

Assumption 3.7 (Spectral Constraints and Joint-Feedback Sensitivity) *We impose the following structural Lipschitz constraints on the full-NDU architecture. (i) Forward-to-backward state sensitivity. The effective coefficients b_ε and g_ε are Lipschitz in the state variable x with constant K_x . (ii) Backward-to-forward control sensitivity. The jointly controlled drift-and-reward architecture is Lipschitz in the full control $\alpha = (a, \theta)$ with constant γ_{spec} , and the feedback selector $\alpha^*(x, p)$ is globally Lipschitz in (x, p) . In the regularized pseudo-inverse formulation, this induces a feedback map*

$$\alpha_\varepsilon^*(x, z) := \alpha^*(x, (\Sigma_\varepsilon \Sigma_\varepsilon^\top)^{-1} \Sigma_\varepsilon z)$$

with volatility sensitivity bounded by $L_\alpha/\sqrt{2\varepsilon}$ for a structural constant $L_\alpha > 0$.

*We define the **effective coupling constant** by*

$$\Lambda_\varepsilon := \sup_{x, z, z'} \frac{|b_\varepsilon(x, z) - b_\varepsilon(x, z')|}{|z - z'|} \leq \gamma_{\text{spec}} \frac{L_\alpha}{\sqrt{2\varepsilon}}. \quad (16)$$

This is the quantity that enters the local- and global-well-posedness arguments for the full joint-control feedback loop.

Proposition 3.8 (Global Well-Posedness under a Sufficient Small-Coupling Bound)

Assume the hypotheses of Lemma 3.6 on $[0, T]$ for a fixed $\varepsilon > 0$. Let $C_X(T, L_{b,x})$ be the forward SDE stability constant and let $C_B(T, \rho, L_{g,z}(\varepsilon))$ be the BSDE stability constant. An explicit sufficient condition for the Picard map to be a contraction is

$$q_\varepsilon(T) := C_B(T, \rho, L_{g,z}(\varepsilon)) C_X(T, L_{b,x}) (L_\Psi^2 + T L_{g,x}^2) \Lambda_\varepsilon^2 < 1. \quad (17)$$

Under (17), the coupled forward-backward system admits a unique global solution on the whole horizon $[0, T]$.

Proof. See Appendix C.11.

Remark 3.9 (The Horizon-Capacity Trade-off) Proposition 3.8 is consistent with a horizon-capacity trade-off in endogenizing preferences. In the sufficient contraction condition (16), longer horizons typically tighten the admissible range of the effective coupling constant, and hence the permissible adaptability of the agent (measured by γ_{spec}). This is not merely an economic interpretation; it is a structural limitation of non-monotone fully coupled FBSDEs, where small coupling replaces the monotonicity condition used in arbitrary-horizon existence theory. If the jointly controlled drift is too sensitive to the gradient Z_t (high γ_{spec}), and the value function is highly sensitive to the state (high L_α), the fully coupled loop enters a regime of positive feedback. In numerics, this can manifest as oscillatory preference dynamics or chattering controls and may hinder the convergence of Deep FBSDE solvers. Consequently, spectral normalization or clipping functions here as practical devices for controlling the effective coupling strength over long horizons. Moreover, for fixed T the constants entering the BSDE stability estimate are themselves ε -dependent through the z -Lipschitz constant of g_ε ; conservative estimates can therefore deteriorate very rapidly as $\varepsilon \downarrow 0$, even up to exponential-in- ε^{-1} scale in worst-case bounds. Proposition 3.8 should thus be read strictly as a fixed-positive- ε result, not as a uniform mechanism that survives unchanged in the vanishing-viscosity limit.

A fully explicit small-coupling check. To make the word “sufficient” concrete, consider the one-dimensional active NDU instance

$$dc_t = -\beta_c c_t dt + \sigma_c dW_t, \quad du_t = (-\kappa u_t + \beta_h \theta_t) dt + \sqrt{2\varepsilon} dW_t^u, \quad \theta_t \in [-M, M],$$

with running reward $r_0 + r_c c + r_u u - \lambda \theta^2/2$ and terminal payoff $\psi_c c + \psi_u u$. The Hamiltonian selector is

$$\theta^*(p_u) = \Pi_{[-M, M]}(\beta_h p_u / \lambda),$$

so the selector Lipschitz constant is $L_\alpha = \beta_h/\lambda$ and the backward-to-forward volatility sensitivity is

$$\Lambda_\varepsilon = \beta_h L_\alpha / \sqrt{2\varepsilon} = \beta_h^2 / (\lambda \sqrt{2\varepsilon}).$$

On the compact control range one also has the conservative driver bound $L_{g,z}(\varepsilon) \leq M\beta_h/\sqrt{2\varepsilon}$. Using the explicit stability constants $C_X = T \exp(2L_{b,x}T)$ and $C_B = 2 \exp((1 + 2\rho + L_{g,z}^2)T)$ gives the numerical audit in Table 23. The first row is nontrivial because r_u or ψ_u can be nonzero, so V depends on the preference state and the optimal θ^* is active; the second row shows how the same sufficient bound fails when preference-drift sensitivity is too high or adjustment cost/noise regularization is too low. These constants are regenerated by `reproducibility/scripts/make_theory_constant_audit.py` and stored in `reproducibility/data/theory_small_coupling_constants.csv`. The larger HBO/NBO and CEM benchmarks are computational-domain solves and stress audits rather than literal certificates of this conservative continuous-time contraction inequality.

The generated constant table is deferred to Appendix A to keep the main text focused on scope and primary evaluator claims.

A complete theory-covered HJB/FBSDE sanity example. The main CEM benchmarks deliberately test finite policy-search value rather than the literal hypotheses of Lemma 3.6. To close the smaller theory–numerics loop, the reproducibility bundle also includes a two-state linear-quadratic service-pressure instance that is checked under the compact/clipped version of the fixed- ε assumptions. The instance is

$$dc_t = (-0.40c_t + a_t)dt + 0.30dW_t^c, \quad du_t = (-0.70u_t + 0.25\theta_t)dt + \sqrt{2\varepsilon}dW_t^u,$$

where c_t can be read as a normalized inventory/backlog imbalance and u_t as a service-pressure state. We take $\varepsilon = 0.08$, compact controls $a_t, \theta_t \in [-1, 1]$, quadratic running reward $-\frac{1}{2}(0.80c^2 + 0.60u^2 + 2.00a^2 + 2.50\theta^2)$, and terminal payoff $-\frac{1}{2}(0.75c_T^2 + 0.50u_T^2)$. The theorem-level selector is the clipped Hamiltonian maximizer

$$a^*(p_c) = \Pi_{[-1,1]}(p_c/2.00), \quad \theta^*(p_u) = \Pi_{[-1,1]}(0.25p_u/2.50),$$

which is unique, globally Lipschitz, compact-valued, and fits the clipped-adjoint/compact-control closure. On the certified grid used in the audit, the clipping is inactive, so the Riccati value $V(t, c, u) = -\frac{1}{2}K_c(t)c^2 - \frac{1}{2}K_u(t)u^2 + C(t)$ gives an exact local HJB residual; outside that certificate domain the statement is the clipped-control problem, not an unrestricted linear-feedback theorem. Table 24 reports the generated constants and residuals, including the same sufficient-bound formula $q_\varepsilon = C_B C_X (L_\Psi^2 + T L_{g,x}^2) \Lambda_\varepsilon^2 T$ with constants computed from the clipped selector and the fixed- ε diffusion. This example is intentionally small, but it is the manuscript's complete numerical object that verifies

the theorem's structural hypotheses on an explicit compact/clipped certificate domain and checks direct- P /HJB consistency; the larger HBO/NBO and CEM benchmarks are reported separately in Section 5.

The generated certificate table is deferred to Appendix A; the larger CEM benchmarks remain decision-layer evidence.

3.4. Conditional Stability Discussion for Stochastic Regularization

A fundamental obstruction in formulating NDU via standard FBSDE theory is the structural degeneracy of the preference state dynamics. In the idealized NDU formulation, the internal preference state u_t may evolve as a controlled ODE/Neural-ODE-style hidden state. While driven by the stochastic environment c_t , the preference state itself has zero intrinsic diffusion ($\sigma_h \equiv 0$).

This degeneracy poses two problems. The first is the **Observability Paradox**: in standard BSDE theory, the volatility process Z_t is the projection of the gradient ∇V onto the noise manifold. If the noise in the u -channel is zero, $Z_{u,t} \equiv 0$ regardless of the magnitude of the marginal utility $\nabla_u V$. Consequently, the parameter-control component of the feedback law becomes ill-defined even though the overt action channel may remain standard through the c -block. The second is **PDE Ellipticity**: the associated Hamilton-Jacobi-Bellman (HJB) equation becomes degenerate parabolic (first-order in u , second-order in c). Classical solutions may not exist, requiring the machinery of viscosity solutions.

3.4.1. The ε -Regularized System For the supporting fixed- ε analysis and the conditional stability discussion, we use a vanishing-viscosity regularization. This is a mathematical device, not an assertion that the paper proves a benign shadow-price limit for every degenerate NDU model. We introduce a regularization parameter $\varepsilon > 0$ and an independent auxiliary Brownian motion $W_t^u \in \mathbb{R}^q$. The regularized state process $X_t^\varepsilon = [c_t^\top, (u_t^\varepsilon)^\top]^\top$ evolves according to:

$$dX_t^\varepsilon = b(X_t^\varepsilon, \alpha_t)dt + \Sigma_\varepsilon(X_t^\varepsilon)d\mathbf{W}_t, \quad (18)$$

where $\mathbf{W}_t = [W_t^\top, (W_t^u)^\top]^\top \in \mathbb{R}^{d+q}$ is the augmented Wiener process, and the regularized diffusion matrix is:

$$\Sigma_\varepsilon(x) := \begin{bmatrix} \sigma_c(c)_{n \times d} & 0_{n \times q} \\ 0_{q \times d} & \sqrt{2\varepsilon}I_{q \times q} \end{bmatrix}. \quad (19)$$

Under the full-row-rank / uniformly elliptic exogenous block assumed in Assumption 3.4, any fixed $\varepsilon > 0$ makes $\Sigma_\varepsilon \Sigma_\varepsilon^\top$ uniformly elliptic. Without that exogenous-block condition, adding preference noise alone would not remove degeneracy in the c -block. Under the usual dynamic-programming and comparison assumptions, the associated non-degenerate HJB admits a unique viscosity solution. Under stronger smoothness assumptions, one may work with a classical representative.

3.4.2. Shadow Price Parameterization Standard Deep BSDE solvers approximate the volatility Z_t^ε . However, evaluating the gradient via $\nabla V^\varepsilon \approx (\Sigma_\varepsilon \Sigma_\varepsilon^\top)^{-1} \Sigma_\varepsilon Z^\varepsilon$ is numerically unstable as $\varepsilon \rightarrow 0$ due to the $O(\varepsilon^{-1/2})$ scaling in the u -block.

We designate the Shadow Price Process $P_t^\varepsilon \in \mathbb{R}^{n+q}$ (representing the gradient $\nabla_x V^\varepsilon$) as the primary learned variable. The volatility Z_t^ε is treated as a derived quantity:

$$Z_t^\varepsilon(P_t^\varepsilon) := \Sigma_\varepsilon(X_t^\varepsilon)^\top P_t^\varepsilon = \begin{bmatrix} \sigma_c(c_t)^\top P_{c,t}^\varepsilon \\ \sqrt{2\varepsilon} P_{u,t}^\varepsilon \end{bmatrix}. \quad (20)$$

The regularized backward dynamics are given by:

$$dY_t^\varepsilon = -\left(f(X_t^\varepsilon, \alpha_t^*) - \lambda \mathcal{L}(X_t^\varepsilon, \alpha_t^*) - \rho Y_t^\varepsilon\right) dt + (\Sigma_\varepsilon(X_t^\varepsilon)^\top P_t^\varepsilon)^\top d\mathbf{W}_t, \quad (21)$$

where the value-process drift is the BSDE driver $g(x, y, \alpha) = f(x, \alpha) - \lambda \mathcal{L}(x, \alpha) - \rho y$. The generalized Hamiltonian $\mathcal{H}(x, y, p, \alpha) = p^\top b(x, \alpha) + f(x, \alpha) - \lambda \mathcal{L}(x, \alpha) - \rho y$ is used for the control law and for the adjoint/shadow-price relations, but not as the direct drift of Y_t^ε .

Proposition 3.10 (Three formulations and their scope) *For fixed $\varepsilon > 0$, the paper uses the following three distinct objects.*

- (i) *The Z -closed FBSDE substitutes $p = A_\varepsilon^\dagger z$ into the Hamiltonian selector. It is the object of Lemma 3.6 and Proposition 3.8. Its constants are fixed- ε constants and deteriorate through $\|A_\varepsilon^\dagger\|_{\text{op}} = O(\varepsilon^{-1/2})$.*
- (ii) *The P -parameterized numerical route represents a shadow price directly, either as $P_\phi = \nabla_x V_\phi$ or through an adjoint/HJB residual. It avoids explicit $Z_u/\sqrt{2\varepsilon}$ inversion, but consistency requires a conservative value-gradient representation or residual/adjoint penalties; supervised one-dimensional oracle audits test only conditioning.*
- (iii) *The degenerate limit route sends $\varepsilon \downarrow 0$ only under comparison, local boundedness, compatible shadow-price selection, and a uniform $\mathcal{H}^2$ bound on P_u^ε .*

Consequently the direct- P conditioning result supplies the stable shadow-price representation, while the PyTorch-AD HBO/NBO value-gradient route supplies the conservative critic, sampled HJB residual, terminal-loss, and held-out rollout diagnostics; a literal solver certificate would additionally require the uniform residual and greedy-gap validation stated in Theorem 2.5.

Proposition 3.11 (Conditional PDE value stability) *For each $\varepsilon > 0$, let V^ε be the unique continuous viscosity solution, in the polynomial-growth class, of the uniformly parabolic ε -regularized HJB equation associated with the regularized NDU control problem. Assume explicitly that: (i) the regularized problems satisfy the dynamic-programming/comparison hypotheses in that class; (ii) the limiting degenerate HJB equation admits a comparison principle in the same polynomial-growth*

viscosity class; and (iii) $\{V^\varepsilon\}_{\varepsilon \in (0,1]}$ is locally uniformly bounded on $[0, T] \times \mathbb{R}^{n+q}$ and the limiting solution class is locally bounded. Then, for any sequence $\varepsilon \downarrow 0$, V^ε converges locally uniformly to the unique viscosity solution V^0 of the degenerate HJB equation

$$\begin{aligned} \partial_t V^0 + \sup_{(a, \theta) \in \mathcal{A} \times \Theta} \left\{ \mu_c(c, a)^\top \nabla_c V^0 + \frac{1}{2} \text{Tr}(\sigma_c(c) \sigma_c(c)^\top \nabla_c^2 V^0) \right. \\ \left. + \nabla_u V^0 \cdot h(c, u; a, \theta) + f(c, u; a, \theta) - \lambda \mathcal{L}(c, u; a, \theta) \right\} - \rho V^0 = 0. \end{aligned}$$

This is a pure viscosity-solution stability statement. It does not require pointwise differentiability of V^ε , and it does not by itself define a preference-shadow-price process.

Proposition 3.12 (Auxiliary martingale vanishing) *In addition to Proposition 3.11, assume that each ε admits a progressively measurable preference shadow-price selection $P_{u,t}^\varepsilon$ along the regularized state. This selection may be the classical gradient $\nabla_u V^\varepsilon(t, X_t^\varepsilon)$ when V^ε is classical, or a semiconcave / stochastic-maximum-principle adjoint selection when such a version is known to exist. Assume further that it is compatible with the martingale representation,*

$$Z_{u,t}^\varepsilon = \sqrt{2\varepsilon} P_{u,t}^\varepsilon,$$

and satisfies the uniform bound

$$\sup_{\varepsilon \in (0,1]} \mathbb{E} \int_0^T |P_{u,s}^\varepsilon|^2 ds < \infty.$$

Then the fictitious preference-noise martingale vanishes in $L^2(\Omega)$:

$$\mathbb{E} \left[\left| \int_0^T \sqrt{2\varepsilon} (P_{u,s}^\varepsilon)^\top dW_s^u \right|^2 \right] = 2\varepsilon \mathbb{E} \int_0^T |P_{u,s}^\varepsilon|^2 ds \rightarrow 0.$$

The uniform bound is therefore a property of the chosen compatible shadow-price version, not of an arbitrary derivative of a merely continuous viscosity solution. Without such a selection and bound, the auxiliary martingale statement is not asserted.

Proof. See Appendix C.12.

Remark 3.13 (Conditional Nature of the Vanishing-Viscosity Argument) Propositions 3.11–3.12 should not be read as deriving a family of regularized FBSDE solutions from Proposition 3.8 uniformly in ε . Indeed, the weak-coupling estimate gives $\Lambda_\varepsilon \lesssim \gamma_{\text{spec}} L_\alpha / \sqrt{2\varepsilon}$, so for fixed horizon the contraction condition of Proposition 3.8 can fail as $\varepsilon \downarrow 0$. The proposition above is therefore framed at the PDE level: for each fixed $\varepsilon > 0$, the regularized HJB is uniformly parabolic and admits a unique continuous viscosity solution under the standing compactness/continuity assumptions, regardless of whether the probabilistic Picard contraction remains globally valid at that ε . The first proposition is therefore a conditional stability statement for the HJB/value family. The second proposition adds the auxiliary martingale conclusion only after a compatible shadow-price selection and its uniform $\mathcal{H}^2$ bound are assumed, excluding boundary-layer blow-up in the preference direction.

Remark 3.14 (Status of the Vanishing-Viscosity Argument) Propositions 3.11–3.12 are stability motivation for the shadow-price parameterization, not the paper’s main existence theorem and not a theorem covering the reported inventory/SKU policy-search code. The nontrivial issue is the uniform preference-shadow-price bound; without it, the auxiliary preference martingale may survive the singular limit. Lemma 3.16 gives an exact but deliberately narrow sufficient regime. A less degenerate sufficient route is available under classical PDE regularity: suppose $\partial_u \Psi$, $\partial_u f$, and $\partial_u \mathcal{L}$ are uniformly bounded by a small constant η , the preference drift is uniformly dissipative,

$$\langle u - u', h(c, u; a, \theta) - h(c, u'; a, \theta) \rangle \leq -\kappa |u - u'|^2,$$

and the c -block remains uniformly parabolic with enough $C^{2+\beta}$ Schauder/Bernstein estimates to differentiate the regularized HJB in the u direction. Differentiating the HJB then gives a first-order inequality for $W^\varepsilon = |\nabla_u V^\varepsilon|^2$ in which the dissipative term $-2\kappa W^\varepsilon$ absorbs the $O(\eta)$ preference-dependence terms and the bounded control/selector terms. In that classical subregime one obtains $\sup_{\varepsilon \in (0,1]} \|\nabla_u V^\varepsilon\|_\infty \leq C(\eta, \|\partial_u \Psi\|_\infty, \kappa^{-1})$, hence the required $\mathcal{H}^2$ bound. We do not elevate this Schauder/Bernstein route to a full theorem because it requires smooth coefficients and classical solutions beyond the viscosity setting used for the main statement. The paper’s core theoretical boundary is therefore a standard fixed-positive- ε verification result plus the NDU-specific numerical observability argument; the vanishing-viscosity paragraph should be read as conditional stability guidance.

Example 3.15 (Non-flat Uniform Preference-Shadow-Price Subregime) The uniform P_u^ε hypothesis in Proposition 3.12 is not limited to the flat case. Consider the scalar preference block

$$du_t = (-\kappa u_t + \beta \theta_t) dt + \sqrt{2\varepsilon} dW_t^u, \quad \theta_t \in [-M, M],$$

with running payoff $R(t, c, a) + \eta u_t - \lambda \theta_t^2/2$ and terminal payoff $G(c_T) + \eta_T u_T$. In the preference coordinate the regularized HJB admits the affine ansatz $V^\varepsilon(t, c, u) = A^\varepsilon(t, c) + B(t)u$; the auxiliary viscosity term drops out because $V_{uu}^\varepsilon = 0$. The coefficient B solves

$$B'(t) = (\kappa + \rho)B(t) - \eta, \quad B(T) = \eta_T,$$

so

$$|P_{u,t}^\varepsilon| = |B(t)| \leq \max\{|\eta_T|, |\eta|/(\kappa + \rho)\}$$

uniformly in ε . The optimal preference-control selector $\theta^*(t) = \Pi_{[-M,M]}(\beta B(t)/\lambda)$ is active whenever $B(t) \neq 0$, so this is a non-flat valuation-control example with V genuinely dependent on u and a uniform preference-shadow-price bound. The example is deliberately low-dimensional and analytic; it demonstrates that the vanishing-viscosity lemma has non-flat sufficient regimes, while the reported HBO/NBO and CEM benchmarks are separate computational-domain audits.

Lemma 3.16 (A Flat-Preference Sufficient Regime for the Uniform $\mathcal{H}^2$ Bound) *Assume, in addition to the standing regularity hypotheses, that (i) the terminal condition is independent of the preference coordinate, $\Psi(c, u) = \Psi(c)$; (ii) the external-state dynamics $\mu_c(c, a)$ and $\sigma_c(c)$, the running reward, and the adjustment penalty are all independent of u ; and (iii) the preference drift has the dissipative affine form*

$$h(c, u; a, \theta) = -\kappa u + \bar{h}(c, a, \theta) \quad \text{for some } \kappa > 0.$$

Then, for every $\varepsilon > 0$, the regularized viscosity solution V^ε is independent of u . Consequently,

$$P_{u,t}^\varepsilon = 0 \quad \text{for all } t \in [0, T],$$

and in particular

$$\sup_{\varepsilon \in (0,1]} \mathbb{E} \int_0^T |P_{u,t}^\varepsilon|^2 dt = 0.$$

Proof. See Appendix C.13.

Remark 3.17 (Adjoint Interpretation under Extra Regularity) If, in addition, the limit value function V^0 is semi-concave and the shadow-price family P^ε is compact in a topology compatible with weak accumulation, then one may interpret weak accumulation points of P^ε as adjoint processes for the limit problem in the sense of the generalized Stochastic Pontryagin Maximum Principle for non-smooth value functions (Peng, 1990; Yong and Zhou, 1999). In this regime, the smooth neural gradients used in the regularized solver should be understood as regularized selections from the super-differential $\partial^+ V^0$ of the limiting viscosity solution, rather than as classical pointwise gradients of V^0 itself. Assuming the constant-diffusion closure of Assumption 3.4, so that the diffusion-gradient contribution in the stochastic maximum principle vanishes, and writing Y^0 for the corresponding limiting value process, the limiting adjoint dynamics formally take the form

$$dP_t^0 = -(\nabla_x \mathcal{H}(X_t^0, Y_t^0, P_t^0, \alpha_t^*) - \rho P_t^0) dt + Q_t dW_t, \quad (22)$$

where $Q_t \in \mathbb{R}^{(n+q) \times d}$ is the martingale coefficient in the adjoint equation. Because ∇_x acts only on the spatial state variable, it annihilates the zero-order discount term $-\rho Y_t^0$ inside $\mathcal{H}$. The explicit $-\rho P_t^0$ correction above is therefore the current-value discount term in the adjoint drift. Equivalently, if one writes the undiscounted Hamiltonian $\tilde{\mathcal{H}}(x, p, \alpha) = p^\top b(x, \alpha) + f(x, \alpha) - \lambda \mathcal{L}(x, \alpha)$, then the drift takes the standard current-value stochastic maximum principle (SMP) form $-\nabla_x \tilde{\mathcal{H}}(X_t^0, P_t^0, \alpha_t^*) + \rho P_t^0$. Under genuinely state-dependent diffusion coefficients, one would also obtain the standard diffusion-gradient term involving $\nabla_x \Sigma^j$ contracted against the corresponding rows of Q_t . That contribution is absent here because the remark is stated under the constant-diffusion closure. This object should be

distinguished from the value-process volatility. In the degenerate limit one has $Z_t^0 = \sigma_c(c_t)^\top P_{c,t}^0 \in \mathbb{R}^d$, whereas Q_t governs the stochastic part of the adjoint process itself. If the feedback controls converge in a topology stable for the forward SDE, then standard stability estimates yield convergence of the regularized forward states to the degenerate limit dynamics.

Remark 3.18 (Numerical Robustness) Propositions 3.11–3.12 motivate the value-stability and auxiliary-noise aspects of our numerical strategy under the stated uniform shadow-price bound. By training the network to approximate P_t directly and constructing Z_t^ε structurally, we avoid the explicit $1/\sqrt{\varepsilon}$ inversion in the preference block. The proposition is not used as an unconditional proof that every NDU model has a benign zero-viscosity limit. The stronger adjoint interpretation used to motivate this parameterization is described separately in Remark 3.17, where the additional regularity requirements are made explicit.

4. NDU Beyond Policy-Class Nesting

The generalized-RL claim is not that a larger action list mechanically improves value. The frozen fixed-preference and fixed-valuation action-only models are useful benchmarks, but the main distinction is operator-level: after the costly valuation control is optimized out, NDU leaves a continuation-value-dependent support operator. That reduced operator is generally non-affine in the continuation value and therefore cannot be relabeled as an evaluator-preserving fixed-reward augmented-state MDP unless the valuation channel is inactive or payoff-equivalent across all active faces.

Frozen and active Bellman endpoints. In the standard fixed-preference control paradigm, the agent optimizes an overt action $a_t \in \mathcal{A}$ while the preference parameter is fixed at $\bar{\theta} \in \Theta$. The value V_{RL} solves

$$\frac{\partial V_{RL}}{\partial t} + \sup_{a \in \mathcal{A}} \left\{ \frac{1}{2} \text{Tr}(\Sigma(x)\Sigma(x)^\top D^2 V_{RL}) + \nabla V_{RL} \cdot b(x, a, \bar{\theta}) + f(x; a, \bar{\theta}) \right\} - \rho V_{RL} = 0, \quad (23)$$

with terminal condition $V_{RL}(T, x) = \Psi(x)$. In NDU, the agent optimizes jointly over $(a, \theta) \in \mathcal{A} \times \Theta$ and the finite-cost value V^λ solves

$$\frac{\partial V^\lambda}{\partial t} + \sup_{(a, \theta) \in \mathcal{A} \times \Theta} \{ \mathcal{A}^{a, \theta} V^\lambda + f(x; a, \theta) - \lambda \mathcal{L}(x, a, \theta) \} - \rho V^\lambda = 0. \quad (24)$$

The following proposition is the former nesting result; it is retained as the frozen-face baseline, not as the full theory of non-collapse.

Proposition 4.1 (Nesting of frozen-valuation faces) *Let Π_{RL} be the original-state fixed-preference/action-only Markov policy class for the HJB (23). Let Π_{AO}^{aug} be the augmented-state action-only class whose policies may observe (c, u) but must keep $\theta \equiv \bar{\theta}$. Let Π_{NDU} be the augmented-state*

NDU policy class with controls $(a, \theta) \in \mathcal{A} \times \Theta$, and suppose the adjustment loss has zero-cost frozen configuration $\bar{\theta}$ as in Assumption 4.8. Then there are payoff-preserving injections

$$\Pi_{RL} \hookrightarrow \Pi_{AO}^{aug} \hookrightarrow \Pi_{NDU}, \quad \iota_0(\pi)(t, c, u) = (\pi(t, c), \bar{\theta}),$$

that preserve the physical action law and fixed-preference payoff on the frozen valuation manifold. Consequently $V_{NDU}(t, c, u) \geq V_{RL}(t, c)$ after identifying V_{RL} with its frozen-manifold extension. If, at a reachable state, the full Hamiltonian or Bellman term exceeds the frozen- θ augmented action-only face by a positive margin that survives transition and adjustment costs, then the inclusion is strict. Thus NDU strictly generalizes fixed-preference RL in the frozen-face sense, while the stronger operator-level distinction is established below.

Proof. See Appendix C.1.

Corollary 4.2 (Non-collapse under positive valuation-control margin) *If the embedded fixed-preference/action-only or augmented fixed- θ face has value at least $\Delta > 0$ below the Full valuation-control value under the same evaluator and information set, NDU is not value-equivalent to any frozen-valuation fixed-preference re-description; otherwise the frozen re-description would have the frozen-face value, contradicting Δ .*

Proof. The frozen face embeds payoff-preservingly by Proposition 4.1; a value-equivalent frozen re-description would attain the frozen-face value, contradicting Δ .

4.1. Operator-Level Non-Collapse

The strict-generalization claim is not merely a policy-class nesting statement. The key distinction is that eliminating the costly valuation control from the Bellman recursion produces a continuation-value-dependent support operator. Such an operator cannot be represented by an ordinary fixed-reward, action-only, augmented-state MDP unless the valuation channel is payoff-irrelevant.

Consider a finite-state discounted approximation of the NDU problem on the augmented state space X . For $x \in X$, overt action $a \in A$, and valuation control $\theta \in \Theta$, write

$$g_\theta(x, a) = r(x, a, \theta) - \lambda \ell(x, a, \theta), \quad P_\theta(\cdot | x, a)$$

for the one-step payoff and transition kernel after choosing (a, θ) . The NDU Bellman operator is

$$(\mathcal{T}^{NDU} V)(x) = \max_{a \in A} \sup_{\theta \in \Theta} \{g_\theta(x, a) + \beta \langle P_\theta(\cdot | x, a), V \rangle\}. \quad (25)$$

An evaluator-preserving fixed-preference augmented-state collapse is an action-only Bellman operator of the form

$$(\mathcal{T}^F V)(x) = \max_{a \in A} \left\{ \bar{g}(x, a) + \beta \langle \bar{P}(\cdot | x, a), V \rangle \right\}, \quad (26)$$

where $\bar{g}$ and $\bar{P}$ are fixed primitives independent of the continuation value V , and the physical action set is still A .

Theorem 4.3 (Operator-level non-collapse of active valuation control) *Fix a reachable state-action pair (x, a) . Suppose two valuation controls $\theta_0, \theta_1 \in \Theta$ generate distinct affine forms*

$$L_i(V) = g_{\theta_i}(x, a) + \beta \langle P_{\theta_i}(\cdot | x, a), V \rangle, \quad i = 0, 1,$$

that are both exposed active faces of the reduced support term: there exist bounded continuation values V^0, V^1 such that $L_i(V^i) = \Phi_{x,a}(V^i)$ and, for $j \neq i$, $L_i(V^i) > L_j(V^i)$, where

$$\Phi_{x,a}(V) = \sup_{\theta \in \Theta} \{g_{\theta}(x, a) + \beta \langle P_{\theta}(\cdot | x, a), V \rangle\}. \quad (27)$$

Then $\Phi_{x,a}$ cannot be represented as

$$\bar{g}(x, a) + \beta \langle \bar{P}(\cdot | x, a), V \rangle \quad (28)$$

for all bounded V . Consequently, no evaluator-preserving, action-only, fixed-preference augmented-state MDP can be value-equivalent to NDU on all terminal payoffs and horizons. Such a collapse is possible only in the degenerate case in which a single frozen valuation control dominates all others for every continuation value, or all active valuation controls induce the same payoff-and-transition hyperplane.

Proof. See Appendix C.2.

4.2. Non-Equivalence to Adaptive Preferences, CMDP, and Risk/Robust Control

The preceding theorem also separates NDU from several nearby paradigms. The separation is not taxonomic; it follows from the shape of the dynamic-programming operator.

Corollary 4.4 (Non-equivalence to common fixed-evaluator formulations) *Assume the switching condition in Theorem 4.3 holds at a reachable state-action pair. Then NDU is not equivalent, under the same physical action set and evaluator, to any of the following classes except in degenerate cases:*

- (i) *an exogenous adaptive-preference model whose preference state evolves without a controlled valuation action;*
- (ii) *a constrained MDP represented by a static Lagrange multiplier;*
- (iii) *a risk-sensitive control model with a fixed entropic risk parameter;*
- (iv) *a robust-control model in which nature minimizes over an ambiguity set while the decision maker controls only the physical action.*

More precisely, exogenous adaptive-preference and fixed-multiplier CMDP Bellman operators are affine in the continuation value after the physical action is fixed; robust minimization over kernels is concave in the continuation value; finite valuation-control NDU is a supremum of affine forms and is convex, generally nonsmooth, and state/time dependent through the endogenous valuation state. Equality on a nontrivial open set of continuation values forces the NDU support operator to be affine, which is exactly the degenerate collapse excluded by Theorem 4.3.

Proof. See Appendix C.3.

4.3. Endogenous Service-Pressure Structure in Inventory

We now show that the NDU layer gives a recognizable OR policy shape rather than only a larger policy class. Consider a finite-horizon periodic-review inventory model with state

$$x = (I, B, z, \theta^-),$$

where I is on-hand inventory, B is backlog, z is a demand/capacity regime, and θ^- is the previous service-pressure level. The physical action is an order-up-to level S , equivalently $q = (S - I)^+$. Demand D_z is stochastically increasing in z . The valuation control $\theta \in [\underline{\theta}, \bar{\theta}]$ is an endogenous service-pressure or shortage-penalty multiplier. A representative one-period objective is

$$\begin{aligned} R_t(x, S, \theta) = & -c(S - I)^+ - h\mathbb{E}[(S - D_z)^+] - p(\theta)\mathbb{E}[(B + D_z - S)^+] \\ & - \lambda(\theta - \theta^-)^2 + \mathbb{E}[V_{t+1}(I', B', z', \theta)], \end{aligned} \quad (29)$$

where $p'(\theta) > 0$ and $\lambda > 0$.

Theorem 4.5 (Base-stock plus monotone service-pressure structure) *Suppose: (i) V_{t+1} is convex in inventory shortfall/backlog and has increasing differences in (B, z) ; (ii) the demand family is stochastically increasing in z and has increasing differences in the usual inventory lattice order; (iii) $p(\theta)$ is increasing and the adjustment penalty is strictly convex; and (iv) the action set and valuation set are compact intervals. Then an optimal NDU policy has the structure*

$$q_t^*(x) = (S_t^*(B, z, \theta_t^*) - I)^+. \quad (30)$$

Moreover, $S_t^*(B, z, \theta)$ is nondecreasing in backlog B , demand/capacity stress z , and service pressure θ . The endogenous valuation selector $\theta_t^*(I, B, z, \theta^-)$ is nondecreasing in backlog and demand stress, and is mean-reverting toward θ^- with strength increasing in λ . If the Hamiltonian is strongly concave in θ , this selector is unique and Lipschitz.

Proof. See Appendix C.4.

This theorem turns the inventory policy surface in Figure 1 from a purely empirical diagnostic into a structural prediction: backlog, demand shocks, and tight capacity increase service pressure, and the order-up-to level rises with that pressure.

4.4. Computational Reduction from Strongly Concave Valuation Control

The valuation-control layer also gives a computable inner problem. Suppose the Hamiltonian can be written as

$$\mathcal{H}(x, a, \theta, p) = H_0(x, a, p) + \langle \eta(x, a, p), \theta \rangle - \frac{\lambda}{2} \|\theta - \bar{\theta}(x)\|_R^2, \quad (31)$$

where $R \succeq mI$, $\lambda > 0$, and Θ is compact and convex.

Theorem 4.6 (Closed-form valuation selector and greedy-gap bound) *Under (31), for each (x, a, p) the valuation-control maximizer is unique and is given by the weighted projection*

$$\theta^*(x, a, p) = \Pi_{\Theta}^R \left[\bar{\theta}(x) + \lambda^{-1} R^{-1} \eta(x, a, p) \right]. \quad (32)$$

Consequently,

$$\sup_{a, \theta} \mathcal{H}(x, a, \theta, p) = \sup_{a \in A} \left\{ H_0(x, a, p) + \sup_{\theta \in \Theta} \left[\langle \eta(x, a, p), \theta \rangle - \frac{\lambda}{2} \|\theta - \bar{\theta}(x)\|_R^2 \right] \right\}. \quad (33)$$

Thus the valuation-control dimension is eliminated analytically before the physical-action search. If η is L_η -Lipschitz in p , then

$$\|\theta^*(x, a, \hat{p}) - \theta^*(x, a, p)\| \leq \frac{L_\eta}{\lambda m} \|\hat{p} - p\|. \quad (34)$$

For an approximate valuation choice $\tilde{\theta}$, the valuation greedy gap is bounded by

$$0 \leq \mathcal{H}(x, a, \theta^*, p) - \mathcal{H}(x, a, \tilde{\theta}, p) \leq \frac{1}{2\lambda m} \|\nabla_{\theta} \mathcal{H}(x, a, \tilde{\theta}, p)\|^2 \quad (35)$$

in the unconstrained interior case, with the projected-gradient norm replacing the Euclidean gradient under constraints. Hence a grid search over $N_A N_{\Theta}^{d_\theta}$ joint choices is replaced by N_A physical-action evaluations plus the closed-form projection (32).

Proof. See Appendix C.5.

4.5. Stable Physical-Metric Improvement in State-Varying Service Systems

The previous results concern the NDU objective. We next give a sufficient condition under which the endogenous service-pressure layer improves a fixed theta-free physical cost. The result does not compare NDU with an oracle that is allowed to redesign the reward as the physical cost itself; rather, it compares NDU with fixed-preference service rules, which is the fixed-reward class targeted by the generalized-RL claim.

Let $z \in \mathcal{Z}$ be an observable demand or service-risk regime with probability $\pi_z > 0$. In regime z , demand has cdf F_z and density f_z . Consider the one-period theta-free physical cost

$$C_z(S) = h\mathbb{E}[(S - D_z)^+] + p_z\mathbb{E}[(D_z - S)^+], \quad (36)$$

where $h > 0$ and the physical shortage cost $p_z > 0$ varies across regimes. The physical minimizer is the newsvendor quantile

$$S_z^* = F_z^{-1} \left(\frac{p_z}{h + p_z} \right). \quad (37)$$

A fixed-preference service rule uses a single shortage-pressure parameter $\bar{p} = p(\bar{\theta})$ and therefore chooses

$$S_z^{\bar{\theta}} = F_z^{-1} \left(\frac{\bar{p}}{h + \bar{p}} \right). \quad (38)$$

An NDU service-pressure rule may choose θ_z by state, with $p(\theta_z) = p_z$ whenever $p_z \in p(\Theta)$.

Theorem 4.7 (Stable theta-free physical-cost improvement) *Assume $p_z \in p(\Theta)$ for every z , and suppose no single fixed $\bar{\theta}$ satisfies $p(\bar{\theta}) = p_z$ for all regimes with positive probability. Suppose also that, on an interval containing all relevant quantiles, the densities satisfy $f_z(s) \geq f_{\min} > 0$ and the local cost slopes are bounded. Then the NDU service-pressure rule that sets $p(\theta_z) = p_z$ strictly improves the expected theta-free physical cost over every fixed-preference service rule:*

$$\sum_z \pi_z C_z(S_z^{\bar{\theta}}) - \sum_z \pi_z C_z(S_z^*) \geq \frac{1}{2} \sum_z \pi_z (h + p_z) f_{\min} (S_z^{\bar{\theta}} - S_z^*)^2 > 0. \quad (39)$$

If a valuation-adjustment cost $\lambda \mathbb{E}[(\theta_z - \bar{\theta})^2]$ is included in the operating objective, the same strict improvement in total cost persists whenever

$$\frac{1}{2} \sum_z \pi_z (h + p_z) f_{\min} (S_z^{\bar{\theta}} - S_z^*)^2 > \lambda \sum_z \pi_z (\theta_z - \bar{\theta})^2. \quad (40)$$

Moreover, the sign of the theta-free improvement is stable: if the empirical cdfs $\hat{F}_z$ satisfy $\sup_s |\hat{F}_z(s) - F_z(s)| \leq \varepsilon$, then the physical-cost gap changes by at most $K\varepsilon/f_{\min}$ for a local constant K . Hence any population gap larger than this perturbation bound remains positive under sufficiently accurate demand estimation.

Proof. See Appendix C.6.

Connection to the inventory evidence. The theorem is intentionally conditional. It supports the OR reading that state-varying shortage/service regimes can yield strict, stable theta-free physical-cost improvements over fixed-service-pressure policies. It does not claim universal dominance over an action-only oracle that directly optimizes the same physical cost with the same action set. This matches Table 3: the main 30-seed inventory benchmark and selected larger equal-budget retest have theta-free physical signs favoring Full NDU, whereas the compact equal-budget rerun is explicitly reported as a mixed row.

High adjustment-cost endpoint. The fixed-preference endpoint remains useful for interpreting the frozen face. We assume $\mathcal{A}$ and Θ are compact and the valuation-adjustment penalty has a unique zero-cost preference configuration $\bar{\theta}$ independent of action:

$$\mathcal{L}(x, a, \theta) = 0 \iff \theta = \bar{\theta}, \quad \text{and} \quad \forall \theta \neq \bar{\theta}, \mathcal{L}(x, a, \theta) > 0. \quad (41)$$

Assumption 4.8 (Identification of the RL Limit) Equation (41) holds; ordinary physical action costs are included in f or in the theta-free physical evaluator, not in the valuation-adjustment uniqueness condition.

Proposition 4.9 (Frozen-valuation consistency under high adjustment cost) Let V^λ be the viscosity solution to the joint-control NDU problem (24), let V_{RL} be the viscosity solution to the fixed-preference problem (23), and let V_{free} be the corresponding zero-cost adaptation problem. Assume Assumption 4.8, the dynamic-programming principle for the λ -family and limit problem, local uniform convergence of the Hamiltonians on compact state/gradient sets as $\lambda \rightarrow \infty$, comparison for the limiting fixed-preference HJB in the relevant polynomial-growth viscosity class, and local boundedness/uniqueness of the values in that class. Then

$$V_{RL}(t, x) \leq V^\lambda(t, x) \leq V_{\text{free}}(t, x) \quad \text{for all } (t, x), \lambda > 0,$$

so $\{V^\lambda\}_{\lambda>0}$ is locally uniformly bounded. As $\lambda \rightarrow \infty$, V^λ converges locally uniformly to V_{RL} on $[0, T] \times \mathbb{R}^{n+q}$. Together with Proposition 4.1, this identifies the embedded RL endpoint and the finite-cost NDU enlargement around it.

Proof. See Appendix C.14.

Example 4.10 (Original-state and fixed-valuation strictness) Consider a one-period deterministic specialization with constant operational state $c \equiv 0$, no state noise, and internal state $u \in \{-1, +1\}$ observed by the augmented controller. First, if the terminal payoff is $-(a - u)^2$ with action $a \in [-1, 1]$, any original-state rule $a(c)$ must choose the same action for both u values and has expected payoff at most -1 , while the augmented-state action-only rule $a(c, u) = u$ attains 0. This is the usual omitted-state separation and shows strictness only relative to original-state fixed-preference policies.

Second, keep the same information and action but add a costly valuation-control coordinate $\theta \in [-1, 1]$ with frozen value $\bar{\theta} = 0$ and payoff

$$-(a - u)^2 - \rho(\theta - u)^2 - \lambda\theta^2, \quad \rho, \lambda > 0.$$

The augmented-state action-only/fixed-valuation class sets $a = u$ but must keep $\theta = 0$, giving payoff $-\rho$. NDU sets $a = u$ and $\theta^*(u) = \rho u / (\rho + \lambda)$, giving payoff $-\rho\lambda / (\rho + \lambda) > -\rho$. Thus costly valuation control yields a strict gain over the stronger same-information augmented-state action-only class, not merely over an original-state policy that failed to observe u .

Remark 4.11 (Interpretation of Example 4.10) The first part is a nesting sanity check; the second part is the NDU-specific separation. Once the valuation state is admitted, NDU is an augmented-state stochastic-control/HJB problem whose distinctive feature is the controlled, costly valuation coordinate and its shadow-price equation.

5. Inventory-Centered Evidence and Residual-Gated Solver Diagnostics

This section is a generalized-RL/HJB-solver evidence ledger centered on inventory service-pressure. Because internal valuation objectives may contain learned θ_t , solver and matched-budget values are reported beside theta-free physical evaluators. The certified objects are the strict generalized-RL embedding/separation/non-collapse theorem, structured FBSDE/HJB statements, NBO HJB-solver certificate, compact/uniform HJB audits, exact finite-cover and cover-refinement certificates, certificate ladder, non-collapse/referee-objection support audits, six-pillar claim-support audit, and validation-cover gap/stability guards. Solver-representation diagnostics cover direct- P versus Z or learned- Z ; PyTorch-AD HBO/NBO rows are sampled residual-gap plus rollout evidence; CEM/inventory/SKU rows are finite-policy and external-evaluator audits.

Primary OR reading and metrics. The main OR experiment is inventory service-pressure control. Its first-order business diagnostics are theta-free operating cost, fill rate, backlog, and service-breach counts; valuation welfare is reported as the decision maker's service-pressure utility, not as a hidden substitute for physical cost. Thus a row is a strong OR win only when valuation welfare and physical service metrics move together. When they diverge, the table labels the row as a tradeoff or robustness failure rather than a dominance result. Throughout the empirical section, baseline roles are restricted to three categories: (i) same-information, same-budget fixed-preference/action-only baselines; (ii) same-information, same-budget augmented-state baselines with θ frozen; and (iii) privileged-information references, which are upper/reference rows and do not enter the main ranking.

Three-layer empirical organization. The *core OR layer* is inventory/SKU/exact-DP evidence (Table 3). The *conditioning and solver layer* contains direct- P/Z diagnostics, direct- P /learned- Z ablations, compact/uniform/finite-cover/refinement HJB certificates, the certificate ladder, objection-support, six-pillar claim-support, and validation-cover guard audits, and PyTorch-AD HBO/NBO sampled residual/rollout/greedy-gap evidence. The *stress layer* contains matched-budget CEM and privileged/heterogeneous portfolio rows that test breadth and information boundaries.

Table 3 Primary OR inventory/SKU ledger. The first column states the business objective and whether the objective contains the endogenous valuation control θ . Lower theta-free cost, backlog, breaches, and adjustment cost are physical diagnostics; they are not forced to move with valuation welfare. The final column makes the intended reading explicit.

Audit / business objective	Theta-free cost	Fill / backlog / breach diagnostics	θ or adjustment read	Paired or bootstrap interval	Conclusion
Main inventory, 30 seeds; Avg. cost gap service-pressure welfare includes θ	-0.01733 (Full minus action-only)	Fill 0.2661; backlog breaches 1.3913	learned service-pressure active	Welfare gap ± 0.01047 ; in this calibration -0.01733 ± 0.00987	physical gain
Equal-budget inventory; same budget	Theta-free cost gap 192-evaluation +0.00926 (action cost leads weakly)	Fill gap 0.000081; backlog gap ± 0.07300	capacity/budget-controlled retest	Welfare gap ± 0.02544 ; gain not robust $+0.00926 \pm 0.02642$	valuation
Selected Pareto 1536 seeds	equal-budget confirmation; theta-free 24 +0.00577 ± 0.00674	Action-minus-Full cost 0.000036; backlog gap ± 0.03266 ; breach gap $+0.05935 \pm 0.10027$	selected after screening; signs favor Full	Welfare gap $\pm 0.00621 \pm 0.00626$	selected postive retest, not universal dominance
Inventory on same Full policy	intervention 4.3233; fixed- θ cost 4.4784; permuted cost 4.3861	Normal Full cost Intervention physical cost and service-pressure welfare	worsens fixed $\theta_t \equiv 1$ and permuted θ_t break state alignment	Normal welfare channel is vs fixed operational, permuted not just -3.7303 vs -3.8750 and -3.7963	scoring
Public SKU panel; valuation welfare includes adaptive shortage pres-sure	val-External cost Full Fill 0.9927 vs backlog 0.0442 vs 0.0432	0.9939; θ maps to implied fill target; max plotted 2.200		Valuation welfare Full valuation -0.0414 ± 0.0154 vs gain with physical tradeoff -0.0520 ± 0.0147	
Exact policy-class nested action/valuation control set	DP; Avg. cost action-only 3.1948 vs full class 3.3912	exact value gain under nested finite Bellman class	Full class has joint controls; action-only embedded	Welfare Full -2.3577 exact valuation-only action cost tradeoff -2.4810 vs action-only -2.5554	
Exact MDP; current/next Bellman state/action	valuation-state Same theta-free cost 0.8306	Same fill 0.9969 and unmet demand 0.3914	coordinated order-and- θ action; expected $\theta = 1.35$	Welfare Full -0.2543 exact valuation-only action gain, physical parity -0.2694 ; action gain $+0.015101$	

5.1. Core OR Application Benchmark: Inventory Service-Level Control

To make the operations interpretation concrete in a standard OR setting, we include a periodic-review inventory benchmark with stochastic regime-shifting demand. The overt action is the order-up-to decision, while the internal NDU channel is an adaptive service-level / shortage-penalty pressure θ_t that changes how aggressively the controller protects against backlog. This benchmark is deliberately canonical: it is a single-item inventory system with holding cost, shortage cost, ordering cost, fill-rate reporting, and service-breach accounting.

The code-level dynamics are explicit. At each period, the policy computes an order quantity $q_t = \max\{0, S_t - I_t\}$ from an order-up-to target S_t ; inventory updates to $I_t + q_t$ before demand arrives; stochastic demand D_t is generated from a persistent high/low demand regime with a sinusoidal seasonal component; filled demand is $F_t = \min\{I_t + q_t, B_t + D_t\}$; and ending inventory/backlog are

$$I_{t+1} = I_t + q_t - F_t, \quad B_{t+1} = B_t + D_t - F_t.$$

The per-period operating cost is holding plus shortage plus order/capacity cost plus $0.045(\theta_t - 1)^2$, and welfare adds a fill-rate reward above the 0.92 service threshold and a service-breach penalty.

All methods observe the same inventory, backlog, demand-EWMA, time, and last-demand-shock features. The fixed-service base-stock comparator optimizes an order-up-to rule with $\theta_t \equiv 1$. The action-only adaptive comparator learns a richer base-stock rule but still fixes $\theta_t \equiv 1$. The valuation-only service rule learns θ_t but uses a fixed service-rule action map. Full NDU jointly learns the order-up-to action and the adaptive service pressure. Each method is optimized by CEM over 30 seeds and then evaluated on 512 held-out demand episodes per seed. The script `reproducibility/scripts/run_inventory_service_level_benchmark.py` regenerates the seed-level and aggregate CSV files.

Table 7 gives the resulting OR application benchmark. Full NDU achieves the best mean welfare, the lowest mean operating cost, and the lowest ending backlog while maintaining the same near-perfect fill-rate regime as the base-stock comparators. The gain is operationally interpretable but moderate: relative to the action-only adaptive base-stock comparator, Full NDU improves mean welfare by 0.0200 and reduces mean cost from 4.3214 to 4.3040; the paired-seed audit reports the same comparison as 0.01997 ± 0.01047 in welfare and -0.01733 ± 0.00987 in cost. Relative to the fixed-service base-stock rule, it improves mean welfare by 0.0282 and reduces mean cost from 4.3291 to 4.3040. This anchors the endogenous-valuation channel in a canonical OR service-level control problem rather than only in internally designed portfolio-style stress tests.

A compact inventory sensitivity grid covers four OR stress settings: the baseline configuration, volatile demand, high shortage penalty, and tight capacity. Each scenario re-optimizes Full NDU, action-only adaptive base-stock, and fixed-service base-stock over 10 seeds and then runs held-out reporting episodes. Table 8 shows that Full NDU is best in baseline, volatile-demand, and tight-capacity settings; in the high-shortage setting it is within 0.0071 welfare of the action-only rule while producing fewer service breaches (0.7552 versus 0.8245). This grid does not claim universal dominance, but it materially reduces the concern that the inventory result is an artifact of a single cost/demand calibration. The script `reproducibility/scripts/run_inventory_sensitivity_grid.py` regenerates the grid.

We strengthen the inventory audit with exact-DP and paired-seed checks rather than adding another portfolio benchmark. First, a finite-state periodic-review MDP solves action-only, valuation-only, and full valuation-control policy classes exactly by backward induction using the same horizon and 122 inventory/regime states. The crucial implementation correction is nestedness: the Full valuation-control class contains the action-only and valuation-only controls and adds joint order/service-pressure controls, instead of using a separate equal-cardinality grid. Table 9 reports the corrected exact-DP result: Full valuation-control has average welfare -2.3577 versus -2.5554 for action-only and -2.4810 for valuation-only. Second, the same table reports paired-seed differences for the CEM inventory benchmark and sensitivity grid. The base 30-seed benchmark has a paired Full-minus-action welfare difference of $+0.01997 \pm 0.01047$ and cost difference -0.01733 ± 0.00987 ; the stress-grid paired intervals are wider, with the high-shortage setting remaining the honest negative case.

Table 4 Compact inventory equal-budget and θ -intervention audit. The first two rows retrain with a common 192-evaluation CEM budget; the last two rows intervene on the same trained Full-NDU policies. Higher welfare and lower cost are better.

Audit row	Variant	Welfare	Avg. cost	θ -free cost
Equal budget action-only	normal	-3.7162	4.3114	4.3114
Equal budget Full NDU	normal	-3.7303	4.3233	4.3207
Full NDU intervention	fixed $\theta_t \equiv 1$	-3.8750	4.4784	4.4784
Full NDU intervention	permuted θ_t	-3.7963	4.3861	4.3834

The script `reproducibility/scripts/make_inventory_dp_policy_class_and_paired_audit.py` regenerates both audits.

A final compact inventory audit addresses the endogenous- θ scoring concern directly. The script `reproducibility/scripts/run_inventory_equal_budget_theta_audit.py` reruns a 12-seed same-evaluation inventory search with 192 CEM evaluations per method and evaluates the learned Full-NDU policies under fixed- θ and permuted- θ interventions. The equal-budget inventory result is mixed: action-only adaptive base-stock is slightly ahead in this compact audit (-3.7162 welfare, 4.3114 cost) versus Full NDU (-3.7303 welfare, 4.3233 cost). This prevents us from claiming budget-free inventory dominance. A separate selected Pareto confirmation, generated by `make_inventory_equal_budget_pareto_audit.py`, uses the same information set and a larger common 1536-evaluation CEM budget over 24 seeds; in that selected retest Full NDU has welfare gap $+0.00621 \pm 0.00626$ and action-minus-Full theta-free cost gap $+0.00577 \pm 0.00674$, with fill, backlog, and breach signs also favoring Full. Because the row was selected after screening, it is reported as a positive robustness check rather than as a replacement for the mixed compact audit. The intervention check shows that the learned service-pressure channel is not a harmless scoring rescale: evaluating the same Full-NDU policy with $\theta_t \equiv 1$ worsens welfare to -3.8750 and cost to 4.4784, while permuting the learned θ_t across episodes worsens welfare to -3.7963 and cost to 4.3861. Table 4 reports the intervention numbers, including the θ -free cost column. Thus the inventory conclusion is deliberately bounded but sharper: in the main 30-seed OR benchmark the valuation channel improves cost/service; in one compact equal-budget rerun action-only can slightly lead; in the selected larger equal-budget retest all reported signs favor Full; and conditional on the learned Full-NDU inventory policy, state-aligned θ_t is operationally useful rather than merely a reward-hacking multiplier.

A public-panel SKU replay sanity check keeps the OR evidence from being limited to AirPas-sengers or internally generated demand. The bundled file `sku_monthly_demand_source.csv` is the public `supplychainpy` monthly SKU sample panel (38 SKUs, 12 months, unit-cost, lead-time, retail-price, inventory, and backlog fields; BSD-style upstream license retained in the package). The script `reproducibility/scripts/run_sku_inventory_or_benchmark.py` normalizes each SKU by its first-eight-month training mean, calibrates only on months 1–8, and evaluates months 9–12 under common train-residual bootstrap perturbations. The comparator set is deliberately stronger than

the earlier AirPassengers bridge: rolling-forecast base-stock with backlog, rolling-forecast lost-sales base-stock, seasonal service-level-tuned base-stock, scenario MPC, approximate DP, and tabular inventory RL. Table 5 shows the corrected objective alignment, but it is deliberately read as a small public-panel replay rather than external industrial validation. Full NDU is best on valuation-adjusted welfare (-0.0414), beating approximate DP (-0.0520), while approximate DP remains slightly lower on theta-free external cost (0.3537 versus 0.3558). The table therefore reports both facts: NDU wins on the valuation-control objective and does not hide the small physical-cost diagnostic gap.

The same script makes the economic meaning of θ_t explicit. In the SKU benchmark, θ_t is calibrated as an adaptive shortage-penalty / service-pressure multiplier: it maps to an implied target fill rate $0.875 + 0.055\theta_t$ (clipped to $[0.86, 0.985]$), enters the order-up-to target, and pays a convex adjustment cost. Figure 1 and the underlying `sku_theta_policy_surface.csv` report a plot-ready policy surface. Holding the representative SKU fixed, θ_t rises from 0.940 with no demand shock, backlog, or capacity pressure to 1.996 under a moderate shock/backlog/tightness combination and to the clipped 2.200 under the strongest plotted stress; the implied target fill rate moves from 0.927 to 0.985. Thus θ_t is interpretable as service pressure/backlog aversion, not just as an endogenous reward label.

The SKU script also includes an exact dynamic program whose state includes the previous valuation-control level and whose Full row chooses both order-up-to and next- θ on a Cartesian action grid. The Bellman recursion is

$$V_t(i, r, \theta^-) = \max_{S, \theta^+} \mathbb{E}\{R(i, r, \theta^-, S, \theta^+, D) + V_{t+1}(i', r', \theta^+)\},$$

with the action-only row obtained by fixing $\theta^+ = \theta^- = 1$ and restricting the maximization to S . Thus the Full action set strictly nests the action-only set after embedding the fixed- θ action. Table 6 is therefore a direct exact-DP valuation-control check, not just a policy-class code-path audit. Full order-and-valuation control improves valuation-adjusted Bellman welfare from -0.2694 to -0.2543 while leaving the theta-free external cost and fill rate unchanged relative to action-only fixed- θ ; the dynamic-shortage-penalty-only row performs badly because valuation control without order control cannot clear backlog. This supports the intended interpretation: valuation control has economic value only when coordinated with the physical action, and it is audited separately from theta-free physical cost.

We also include an exact dynamic-programming oracle for a simplified fixed-service lost-sales special case. This special case removes the adaptive valuation channel and solves the classical finite-horizon inventory problem exactly by backward induction over integer inventory states. The purpose is not to make Full NDU beat an exact oracle for a different, simplified model; it is to verify that the inventory code path reproduces the standard OR base-stock structure in the fixed-service limit before the richer adaptive-service benchmark is run. In this special

Table 5 Public-panel SKU inventory replay sanity check with stronger OR comparators. Policies train on months 1–8 of the bundled `supplychainpy` SKU panel and evaluate months 9–12 over 24 common bootstrap seeds. Higher valuation welfare is the NDU objective; lower external cost and higher theta-free external welfare are physical diagnostics.

Method	Valuation welfare	External cost	Theta-free welfare	Fill rate	Backlog
Full NDU theta-service control	-0.0414 ± 0.0154	0.3558 ± 0.0109	-0.0578 ± 0.0159	0.9927	0.0442
Approximate DP inventory	-0.0520 ± 0.0147	0.3537 ± 0.0102	-0.0520 ± 0.0147	0.9939	0.0432
Seasonal service-tuned base-stock	-0.1006 ± 0.0190	0.3730 ± 0.0129	-0.1006 ± 0.0190	0.9893	0.0577
MPC scenario base-stock	-0.1382 ± 0.0202	0.3894 ± 0.0132	-0.1382 ± 0.0202	0.9880	0.0634
Rolling forecast base-stock/backlog	-0.1571 ± 0.0209	0.4018 ± 0.0141	-0.1571 ± 0.0209	0.9854	0.0708
Tabular inventory RL	-0.1605 ± 0.0185	0.4098 ± 0.0132	-0.1605 ± 0.0185	0.9856	0.0734
Rolling forecast lost-sales base-stock	-0.3878 ± 0.0306	0.6143 ± 0.0240	-0.3878 ± 0.0306	0.9228	0.0000

THETA POLICY SURFACE




Figure 1 Interpretable SKU-policy surface for the calibrated inventory θ_t rule. Columns vary demand-shock ratio, rows vary backlog, and panels vary capacity tightness; darker cells indicate larger service-pressure θ_t . The source values are in `sku_theta_policy_surface.csv`; the figure is a readability-oriented rendering of that table.

Table 6 Exact valuation-control inventory MDP. The Full row solves a Bellman recursion over inventory, demand regime, and previous θ ; the action is a Cartesian pair of order-up-to level and next valuation-control level.

Exact-DP row	Actions	Avg. welfare	Theta-free cost	Fill rate	Mean θ
Action-only fixed θ	8	-0.2694	0.8306	0.9969	1.0000
Full order and valuation control	32	-0.2543	0.8306	0.9969	1.3500
Dynamic shortage penalty only	4	-8.7196	8.1671	0.4348	0.8277

case, exact DP recovers order-up-to level 14, and the best stationary base-stock and myopic newsvendor rules coincide with the exact dynamic-programming optimum; see Table 10. The script `reproducibility/scripts/run_inventory_dp_oracle_sanity.py` regenerates the exact oracle CSV files without external dependencies.

A compact public-demand calibration bridge reduces the purely synthetic-data concern without turning the paper into an empirical industry study. The script `reproducibility/scripts/run_inventory_real_demand_calibration.py` uses the public `AirPassengers` monthly demand trace from `R datasets` via `Rdatasets`, scales it to the inventory units used above, calibrates on the first 108 months, and replays the final 36 months. Table 11 shows

Table 7 Core inventory/service-level benchmark (30 seeds; 512 held-out demand episodes per seed). Higher mean welfare is better; lower cost, backlog, and service breaches are better.

Method	Mean welfare	Avg. cost	Fill rate	Ending backlog	Service breaches
Full NDU inventory	-3.7081 ± 0.0078	4.3040 ± 0.0077	0.9997	0.2661	1.3913
Action-only adaptive base-stock	-3.7280 ± 0.0091	4.3214 ± 0.0084	0.9996	0.3259	1.4609
Fixed-service base-stock	-3.7362 ± 0.0142	4.3291 ± 0.0134	0.9997	0.2898	1.4849
Valuation-only service rule	-3.7779 ± 0.0131	4.3726 ± 0.0123	0.9997	0.2964	1.4217

Table 8 Compact inventory sensitivity grid (10 seeds per method/scenario). Higher welfare is better; Δ is Full NDU minus action-only adaptive base-stock welfare.

Scenario	Full welfare	Action-only welfare	Fixed-service welfare	Δ vs action-only	Full avg. cost
Baseline	-3.7134	-3.7206	-3.7392	+0.0073	4.3098
Volatile demand	-5.5825	-5.5928	-5.6138	+0.0102	6.1703
High shortage penalty	-3.8981	-3.8911	-3.9164	-0.0071	4.6150
Tight capacity	-3.7893	-3.8153	-3.8128	+0.0260	4.4723

Table 9 Inventory exact policy-class DP and paired-seed audit. Panel A solves a simplified finite-state inventory MDP exactly; the Full class contains the action-only and valuation-only controls and adds joint controls. Panel B reports paired Full-minus-action-only differences; positive welfare and negative cost/breach differences favor Full NDU.

Panel A: exact finite-state policy-class DP (horizon 10; 122 states; controls: 5 action, 5 valuation, 16 full)				
Policy class	Avg. welfare	Avg. cost	Fill rate	Service breaches
Full valuation-control exact DP	-2.3577	3.3912	0.999998	0.0000
Valuation-only exact DP	-2.4810	3.4079	0.999993	0.0007
Action-only exact DP	-2.5554	3.1948	0.999977	0.0067
Panel B: paired seed differences, Full NDU minus action-only adaptive base-stock				
Scenario	n pairs	Δ welfare	Δ cost	Δ service breaches
Base 30-seed benchmark	30	$+0.01997 \pm 0.01047$	-0.01733 ± 0.00987	-0.06960 ± 0.08229
Baseline grid	10	$+0.00728 \pm 0.02690$	-0.00669 ± 0.02658	-0.01146 ± 0.19409
Volatile demand	10	$+0.01025 \pm 0.02787$	-0.00266 ± 0.02344	-0.21562 ± 0.25271
High shortage penalty	10	-0.00707 ± 0.03397	$+0.01106 \pm 0.03286$	-0.06927 ± 0.09925
Tight capacity	10	$+0.02602 \pm 0.02262$	-0.02337 ± 0.02118	-0.07266 ± 0.14302

Table 10 Exact-DP/base-stock oracle for a simplified fixed-service inventory special case. The exact dynamic program confirms the classical order-up-to structure before the adaptive-service benchmark in Table 7.

Method	Order-up-to rule	Total cost	Avg. cost	Lost sales	Fill rate
Exact dynamic-programming oracle	state/time dependent	37.2323	3.1027	1.5338	0.9865
Best stationary base-stock	14	37.2323	3.1027	1.5338	0.9865
Myopic newsvendor base-stock	14	37.2323	3.1027	1.5338	0.9865
No reorder	none	509.0455	42.4205	106.0000	0.0702

Table 11 Public real-demand calibration bridge using the **AirPassengers** monthly demand trace. Policies are calibrated on the first 108 months and replayed on the final 36 months after scaling passenger counts to inventory demand units.

Method	Test avg. cost	Fill rate	Ending backlog	Service breaches
Adaptive valuation-service replay	4.2876	1.0000	0.0000	7
Seasonal quantile base-stock	14.2924	0.9922	6.0000	33
Fixed empirical base-stock	12.7972	0.9947	4.1000	24
No reorder	1035.9143	0.0150	759.7551	36

the resulting held-out trace check. The adaptive valuation-service replay has lower held-out average cost than fixed empirical and seasonal base-stock rules while maintaining full aggregate fill on the replay horizon. This result should be interpreted cautiously: **AirPassengers** is a public demand trace, not a retail SKU panel; the adaptive replay has a richer EWMA/backlog/service-pressure rule than the simple empirical and seasonal base-stock comparators; and the large cost gap is therefore a calibration bridge rather than external industrial validation.

Package provenance. The artifact also includes a deterministic finite-guard self-replay certificate and a local environment audit for checker, L^AT_EX, zip, checksum, and extracted replay prerequisites. This is

Table 12 NDU/NBO generalized-RL and HJB-solver evidence ledger. Margins are Full NDU minus the embedded fixed-preference/action-only or fixed-valuation comparator unless otherwise noted. The ledger is generated by `make_ndu_nbo_generalization_certificate.py` and checked by the reproducibility gate.

Claim	Evidence source	Margin / status	Interpretation
Fixed-preference/fixed-valuation embedding	Proposition 4.1	proved inclusion	Original-state fixed-preference and augmented-state fixed- θ action-only policies embed by fixing $\theta = \bar{\theta}$.
Operator-level non-collapse	Theorem 4.3 + Corollary 4.4	active-face corollary	Eliminating θ yields a non-affine support Bellman term; fixed-reward augmented-state, adaptive-preference, CMDP, and risk/robust collapses require degenerate faces or changed primitives.
Strict finite-state generalization / equivalence exclusion	Exact DP + non-collapse/state-reward audits	+0.19769; 5/5 + 5/5 guards pass	Full valuation-control improves the frozen action-only face; exact-DP/MDP, finite-cover, same-budget, and NBO guards rule out fixed-preference collapse and reward-shaping/state-augmentation equivalence.
Strict policy-face hierarchy	Policy-face hierarchy audit	5/5 guards pass; Full-best face +0.12332	Full NDU strictly exceeds action-only and valuation-only faces, preserves zero-gap finite-cover certificates, and improves through joint-control cover refinement.
Referee-response support synthesis	Referee response synthesis audit	8/8 guards pass	The generalized-RL definition, anti-relabeling closure, HJB-solver certificates, validation-cover neural route, direct- P , same-information evidence, and traceability are checked as one response matrix rather than an external review outcome.
Exact valuation-state MDP strictness	SKU/inventory valuation-state DP	+0.01510 welfare	Bellman recursion over (inventory, θ) improves the embedded fixed- θ row with identical theta-free cost in this instance.
Ideal NBO HJB certificate domain	Theory-covered clipped audit	residual 1.56×10^{-16}	Compact clipped instance satisfies the ideal residual certificate to numerical precision.
Uniform residual/greedy HJB certificate	Analytic compact LQ NDU audit	residual 2.25×10^{-16} ; gap 0	Exact row verifies the theorem ingredients; perturbed NBO sequence closes residual, terminal, and greedy gaps uniformly.
Finite-cover inventory Bellman-HJB certificate	Exact inventory DP cover	residual/gap 0; +0.19769 welfare	Full valuation-control cover contains action-only/valuation-only faces and exactly solves 122 states, 1220 Bellman nodes, and 16 controls.
Inventory cover-refinement certificate	Nested exact action covers	radius $1.206 \rightarrow 0.546$; margin $0 \rightarrow 0.19769$	Each refined cover has zero residual/terminal/greedy gaps; value improves monotonically as the action cover is refined.
HJB-solver certificate ladder	Integrated reproducibility audit	7 rungs; 4 proof-certified + 3 sampled/representation	Exact, uniform, finite-cover, refinement, sampled HBO/NBO gap closure, and direct- P guard rows use one residual/terminal/greedy/value tuple.
Neural validation-cover HJB audit	Sampled tuple + joined neural guard chain	Exp. I/inventory ratios $0.139/0.192$; interval/stress pass; cross-artifact $2/2$; refinement/stress $6/6$; holdout/bootstrap $6/6$; concordance $2/2$; leave-one-out $12/12$; bounded-composition $2/2$; stress-frontier $2/2$ pass	Large HBO/NBO rows tie residual, terminal, greedy-gap, replay, stability, envelope, position, and bottleneck-frontier guards to Proposition 2.9 and Corollaries 2.10–2.23. Not an unrestricted global neural Lipschitz proof.
Referee-objection support matrix	Reproducible support audit	6 objections mapped	Generalized-RL status, HJB certificate, cover refinement, sampled neural gaps, direct- P , and same-budget fairness are each tied to theorem and experiment anchors.
Six-pillar claim-support audit	Six-pillar reproducible audit	6/6 pillars pass	Policy-class strictness, HJB-solver certificates, cover refinement, validation-cover guards, direct- P , and same-information roll-out support jointly support the bounded title claim.
Direct- P over Z conditioning	Same-benchmark inventory audit	$19.64 \times$ lower RMSE scale at 10^{-3}	Z inversion amplifies service-pressure error and loses 0.1020 welfare relative to direct- P .
Direct- P vs learned- Z solver ablation	Isomorphic inventory representation learner	$27.65 \times$ lower θ RMSE at 10^{-3}	Same states, features, ridge budget, policy map, and evaluation; learned- Z loses 0.0503 welfare after inversion.
Neural direct- P vs learned- Z retraining	Same-architecture trainable MLP audit	$21.34 \times$ lower θ RMSE at 10^{-3}	Same neural architecture and optimizer budget; learned- Z loses 0.0696 welfare in the near-deterministic preference regime.
Full HBO/NBO suite	PyTorch-AD actor-critic audit	3 gains + selected Exp. III gain	Full NDU improves Exp. I, Exp. II, and inventory; a selected same-architecture Exp. III sensitivity row has gap +0.000229, while the privileged factor reference remains outside the feasible information set.
Inventory HBO/NBO strict gain	PyTorch-AD inventory audit	+0.66557 welfare	Value-gradient NBO improves the action-only inventory comparator and fixed-service reference on service-pressure welfare.
Exp. I residual-gap closure	Targeted HBO/NBO residual-closing audit	HJB $0.139 \times$; terminal $0.50 \times$	Stronger critic residual run closes 86.1% of native HJB gap and preserves held-out utility.
Inventory residual-gap closure	Targeted HBO/NBO residual-closing audit	HJB $0.192 \times$; terminal $0.94 \times$ norm.	Stronger inventory run closes 80.8% of native HJB gap and preserves held-out welfare.
Inventory normalized-residual closure	Targeted normalized-residual audit	RMSE $0.979 \times$; HJB $0.251 \times$	Pareto actor run lowers normalized RMSE below headline while preserving native, terminal, and welfare guards.

Table 13 Theory-assumption and evidence-role matrix. A check mark means the row is intended to satisfy the named formal assumption class; a dash means it is an empirical or diagnostic stress test rather than a theorem-covered object.

Object	Fixed- ε closure	Main evidence role	Boundary
Direct- P bridge	yes	HJB-solver representation audit	Same-benchmark conditioning diagnostic plus isomorphic and trainable-neural direct- P versus learned- Z representation learners; value-gradient route avoids singular inversion.
Compact clipped / uniform HJB audit	yes	Theory-covered solver certificate	Compact controls, clipped selector, projected volatility, small-coupling constants, and analytic sup residual/terminal/greedy-gap closure.
Inventory/SKU pressure studies	service- finite-cover exact	Primary OR application	Finite-horizon valuation-control benchmarks with theta-free physical diagnostics; the exact inventory DP row has a zero-gap finite-cover Bellman-HJB certificate and nested action-cover refinement sequence.
Full PyTorch-AD HBO/NBO suite	partial	HJB-solver residual/greedy-gap audit	Value-gradient route $P_\beta = \nabla_x V_\beta$ with sampled actor residuals, terminal losses, rollouts, and greedy-gap diagnostics; compact rows provide literal certificate and large rows report certificate-gap closure.
Appendix CEM / privileged-reference rows	–	Stress diagnostics	Attained-policy rows; not used as the main ranking.

Table 14 Same-benchmark inventory direct- P versus Z -inversion diagnostic. All rows use the same inventory simulator and 30 seed blocks. The welfare gap is Z -inversion minus direct- P at the same ε ; lower θ RMSE and higher welfare favor direct- P as preference noise becomes small.

ε	direct- P θ RMSE	Z -inversion θ RMSE	Z /direct RMSE ratio	Z welfare gap
10^{-1}	0.01517	0.03388	2.23	−0.00096
10^{-2}	0.01517	0.10563	6.96	−0.01105
10^{-3}	0.01517	0.29802	19.64	−0.10202

local package provenance, not external reproduction, review, submission, or acceptance evidence. The lightweight replay dependency-closure audit `make_lightweight_replay_dependency_closure_audit.py` further checks that the checker, environment audit, package self-replay command, and six finite neural-guard builders import only standard-library modules and no optional full-regeneration packages such as `torch`, `pandas`, `matplotlib`, or `seaborn`.

Same-benchmark direct- P versus Z diagnostic. The direct- P evidence includes a same-benchmark inventory direct- P versus Z -inversion audit on the canonical 40-period inventory/service-level simulator. The demand law, cost accounting, fill-rate metric, and service-breach accounting are unchanged; only the representation of the service-pressure shadow price changes. Direct- P observes $\hat{P} = P_* + \eta$, while the Z route observes $\hat{Z} = \sqrt{2\varepsilon}P_* + \eta$ and recovers $\hat{P} = \hat{Z}/\sqrt{2\varepsilon}$ under the same target-scale perturbation η . Table 14 is therefore a same-benchmark inventory direct- P versus Z -inversion audit, not a separate toy example.

Uniform residual/terminal/greedy certificate audit. The theorem-covered compact HJB row is strengthened by an analytic uniform certificate audit generated by `make_uniform_hjb_greedy_certificate.py`. It uses the same compact two-state service-pressure LQ NDU instance as the clipped HJB audit, but reports exactly the quantities in Theorem 2.5: $\sup$

Table 15 Analytic uniform residual/terminal/greedy-gap certificate for the compact NDU HJB-solver theorem. Bounds are sup bounds on $[0, T] \times [-1.2, 1.2]^2$ with compact controls; B is the reported certificate value/policy bound

proxy $\delta_T + T\delta_H + T\delta_A$.					
γ	opt. residual sup	terminal sup	greedy-gap sup	B	value-error sup
0	2.25×10^{-16}	0	0	5.61×10^{-17}	0
0.08	0.18724	0.11520	1.63×10^{-3}	0.16242	0.11520
0.04	0.09301	0.05760	3.68×10^{-4}	0.08095	0.05760
0.02	0.04636	0.02880	8.73×10^{-5}	0.04041	0.02880
0.01	0.02314	0.01440	2.13×10^{-5}	0.02019	0.01440

optimal-Hamiltonian residual, terminal residual, greedy-action gap, and the induced value/policy certificate bound. The exact row has residual 2.25×10^{-16} , zero terminal gap, and zero greedy gap. The perturbed critic/actor rows form an explicit NBO certificate sequence: as γ halves from 0.08 to 0.01, the residual, terminal, greedy gap, and certificate bound all shrink monotonically. This supplies a low-dimensional uniform solve beyond sampled residual diagnostics and directly instantiates the HJB-solver theorem.

Finite-cover inventory Bellman-HJB certificate. Beyond sampled HBO/NBO gaps, `make_inventory_finite_cover_hjb_certificate.py` gives an exact finite-cover certificate on the nested inventory valuation-control DP: 122 states, 1220 nonterminal state-time Bellman nodes, 122 terminal nodes, and 16 Full controls containing the action-only and valuation-only faces. Bellman/HJB residual, terminal residual, and greedy-action gap are zero, so Corollary 2.7 certifies the solve, and Full strictly improves the embedded action-only face by 0.19769 average welfare. `make_inventory_cover_refinement_certificate.py` shows zero gaps on nested covers; the action fill-distance proxy shrinks from 1.206 to 0.546 and the margin rises $0 \rightarrow 0.19769$. The seven-row audit ladder is `make_hjb_solver_certificate_ladder.py`; support scripts include `make_generalized_rl_noncollapse_audit.py`, `make_reward_shaping_state_augmentation_closure_audit.py`, `make_policy_class_face_hierarchy_audit.py`, `make_referee_full_support_synthesis_audit.py`, `make_referee_objection_closure_audit.py`, `make_neural_validation_cover_hjb_gap_audit.py`, `make_validation_cover_stability_guard_audit.py`, `make_neural_empirical_envelope_audit.py`, `make_neural_interval_domain_certificate.py`, `make_neural_interval_sensitivity_audit.py`, `make_bounded_architecture_lipschitz_budget_audit.py`, `make_bounded_architecture_budget_stress_audit.py`, `make_neural_guard_cross_artifact_consistency_audit.py`, `make_neural_guard_refinement_stress_audit.py`, `make_neural_guard_holdout_bootstrap_stability_audit.py`, `make_neural_guard_concordance_audit.py`, `make_neural_guard_leave_one_out_robustness_audit.py`, `make_neural_bounded_composition_envelope_certificate.py`, and

Table 16 Isomorphic inventory direct- P versus learned- Z solver-representation ablation. Both learned rows use the same training states, feature map, ridge budget, policy map, simulator, and 24 seed blocks; learned- Z is inverted by $Z/\sqrt{2\varepsilon}$ before the identical NDU inventory actor is evaluated. Welfare gap is learned- Z minus learned direct- P .

ε	direct- P θ RMSE	learned- Z θ RMSE	Z /direct RMSE ratio	learned- Z welfare gap
10^{-1}	0.00261	0.00730	2.80	-0.00295
10^{-2}	0.00261	0.02297	8.81	-0.01211
10^{-3}	0.00261	0.07212	27.65	-0.05030

Table 17 Same-architecture neural inventory direct- P versus learned- Z retraining ablation. Both learned rows train a NumPy tanh MLP with the same width, optimizer steps, target-noise scale, state cloud, policy map, simulator, and 18 seed blocks. The welfare gap is neural learned- Z minus neural direct- P ; the near-deterministic preference regime is the relevant singular-conditioning case.

ε	direct- P θ RMSE	neural- Z θ RMSE	θ RMSE ratio	P RMSE ratio	neural- Z welfare gap
10^{-1}	0.01154	0.02503	2.17	2.09	+0.01016
10^{-2}	0.01154	0.08268	7.17	6.93	+0.01572
10^{-3}	0.01154	0.24619	21.34	23.63	-0.06960

make_neural_guard_stress_frontier_audit.py. They record the generalized-RL non-collapse audit, state-augmentation/reward-shaping support audits, stability guard audit, deterministic replay audit, empirical-envelope guard, interval-domain certificate guard, interval-sensitivity/refinement guard, bounded-architecture Lipschitz budget audit, bounded-budget stress-margin audit, cross-artifact consistency guard, joined refinement/stress guard, holdout/bootstrap stability guard, concordance guard, leave-one-out robustness guard, bounded-composition envelope guard, and stress-frontier guard named in the manifest.

The operational reading is now both a conditioning theorem check and a learned-representation solver ablation. Table 14 shows the analytic perturbation mechanism: on the same inventory benchmark, equal-scale noise in the reported Z target becomes an economically visible service-pressure error after division by $\sqrt{2\varepsilon}$. At $\varepsilon = 10^{-3}$ the theta RMSE ratio 19.64 produces a 0.1020 welfare loss for the Z route relative to direct- P , while the direct- P row remains close to the oracle- P replay. Table 16 then repeats the comparison as an isomorphic learned solver component: the training states, feature map, ridge budget, target-noise scale, policy map, simulator, and evaluation seeds are fixed, and only the learned representation changes from P_* to $Z_* = \sqrt{2\varepsilon}P_*$. This is the missing direct- P versus learned- Z ablation for the inventory HJB/NBO bridge.

The neural retraining audit generated by make_inventory_neural_direct_p_learned_z_ablation.py confirms that the singular-conditioning effect is not an artifact of ridge regression. With identical architecture and optimizer budget, learned- Z inversion has $21.34\times$ larger θ RMSE and $23.63\times$ larger P RMSE at $\varepsilon = 10^{-3}$, producing a -0.06960 welfare gap in the near-deterministic preference regime. At larger ε the learned- Z actor can occasionally trade extra valuation error for a small rollout improvement, so the paper reads the welfare column as a policy-performance diagnostic and the RMSE ratios as the direct representation test of the HJB/NBO shadow-price channel.

Primary ledger and protocol. The primary generalized-RL/HJB-solver reading is Table 12, Table 3, and the inventory section. The packaged HBO/NBO files are a PyTorch-AD actor-critic solver map with explicit sampled residual-gap diagnostics: trainable actor/critic MLPs use PyTorch value derivatives/Hessians and actor Hamiltonian gradients, common benchmark seeds, and the aligned inventory service-pressure evaluator. The implementation is a float64 CPU PyTorch run (algorithm identifier `hbo_full_actor_critic_pytorch_ad`) with width-56, depth-2 networks, constant-bias initialization, no fixed random hidden features, and no Bellman-grid pretraining. Under this map, Full NDU is better than fixed-preference/action-only valuation comparators in Experiments I–II and inventory; a selected same-architecture Exp. III sensitivity row gives a small feasible-comparator gain, while the original Exp. III row is best read as parity and the privileged hidden-factor reference remains an outside-information ceiling. The separate decision-layer bridge records CEM/exact-DP attained-policy statuses: Exp. I strict valuation-objective gain; Exp. II CEM practical parity; Exp. III privileged-reference stress; inventory intervention relevance but budget sensitivity; and SKU/exact-DP valuation-welfare gains with theta-free cost diagnostics. Full-budget CEM rows are appendix stress tests rather than the main empirical ranking.

Full PyTorch-AD HBO/NBO actor-critic solver audit. Table 18 is the differentiable actor-critic solver audit. The runner implements the local Neural Bellman Operators manuscript’s ingredients: internal/external actor maps, a value critic, Hamiltonian performance loss, HJB-residual critic loss, and a terminal boundary penalty. The reported training and held-out values are generated by the PyTorch-AD actor-critic pass over all four benchmark families and are compared to the ingredients of Theorem 2.5. Because the table reports sampled actor-induced L^2 residuals and terminal losses, not a uniform optimal-Hamiltonian residual, these rows are residual-gap solver evidence rather than literal certificates; Table 19 adds a separate sampled discrete greedy-gap audit and reaches the same non-certificate conclusion.

For reproducibility, sampled input $U = (t/H, x)$ feeds trainable bounded actor and value-critic MLPs, $A_\eta(U) = \ell + (u - \ell) \odot \sigma(f_\eta(U))$ and $V_\beta(U) = g_\beta(U)$. PyTorch automatic differentiation computes $\partial_t V_\beta$, $\nabla_x V_\beta$, Hessians, and actor Hamiltonian gradients for the native residual

$$R_\beta = -\partial_t V_\beta - \{r(x, A) + \nabla_x V_\beta^\top b(x, A) + \frac{1}{2} \text{Tr}(\Gamma(x, A) \nabla_x^2 V_\beta)\}.$$

The critic minimizes squared HJB residual plus terminal loss, the actor maximizes the Hamiltonian objective, and held-out rollouts determine the benchmark ranking. Table 18 reports native squared residual, scale-normalized RMSE, and terminal loss as sampled compact-domain gaps; literal use of Theorem 2.5 still requires verified sup-norm residual, terminal, and greedy-action gaps.

Table 18 Full PyTorch-AD HBO/NBO actor-critic solver audit with sampled HJB-residual diagnostics. Values are 30 held-out seed blocks after training; paired intervals compare Full NDU with the listed feasible comparator. Native residuals, normalized RMSE, and terminal losses are empirical actor-induced residual diagnostics, sampled greedy-gap certificate ingredients.

Benchmark	Held-out actor result	Comparator / physical read	Residual diagnostic	diag-
Exp. I habit	Mean utility Full -2.7966 vs fixed-pref. -3.1903 ; gap $+0.3937 \pm 0.0004$	Terminal wealth 0.8630 vs 0.8629 ; utility gain without a visible wealth loss	native HJB 2.13×10^3 ; norm. RMSE 1.33 ; terminal 85.2	
Exp. II hidden regime	Full 0.0113198 vs fixed-pref. recurrent 0.0113185 ; gap $+1.24 \times 10^{-6} \pm 1.67 \times 10^{-10}$	Privileged action diagnostic row 0.0113186 ; Full-minus-diagnostic $+1.17 \times 10^{-6}$; detection 0.6043 vs reference 1.0000	native HJB 5.08×10^{-4} ; norm. RMSE 0.022 ; terminal 4.4×10^{-6}	
Exp. III allocation	Full 0.063096 vs fixed-pref. joint 0.063106 ; gap $-9.91 \times 10^{-6} \pm 2.02 \times 10^{-5}$	Privileged factor-reference row 0.085336 remains higher; feasible Full/fixed comparison is statistical parity	native HJB 3.34×10^{-2} ; norm. RMSE 0.164 ; terminal 8.2×10^{-5}	
Inventory service	Full welfare -3.6375 vs action-only -4.3031 ; gap $+0.6656 \pm 0.0002$	Fixed-service base-stock -4.3262 ; Full-minus-fixed $+0.6886$; cost 5.0211 vs action-only 4.9314 and backlog 0.0258 vs 0.0286	native HJB 6.26×10^4 ; norm. RMSE 1.84 ; terminal 2.06×10^3	

Table 19 Sampled discrete greedy-action-gap audit for the HBO/NBO rows. Residuals and greedy gaps are validation-sample/grid certificate-gap measurements for the benchmark actors. Smaller residual and gap values are better; terminal sup is the sampled terminal boundary error.

Benchmark	Configuration	Actor residual L^2/sup	Grid residual L^2/sup	Greedy p95/sup	gap Terminal sup
Exp. I	headline-like 450 steps	50.6/105.1	51.3/105.0	7.61/11.97	16.63
Exp. I	actor-gap reduced, 900 steps	41.5/98.6	41.4/98.2	2.78/4.13	26.00
Exp. I	residual/terminal critic, 2400 steps	32.0/103.3	29.5/102.6	47.26/110.71	11.74
Inventory	headline-like 450 steps	204.6/805.6	417.0/1315.9	1257.67/1556.63	75.14
Inventory	actor-gap reduced, 900 steps	46.3/167.6	53.2/167.6	29.33/120.04	75.05
Inventory	residual balanced, 2400 steps	139.1/642.7	429.5/1336.5	1357.50/1652.08	72.35

Sampled greedy-action-gap audit. `make_hbo_greedy_gap_audit.py` computes the actor residual, discrete action-grid optimal-Hamiltonian residual, sampled selector gap $\max_{a \in A_{grid}} H_\beta(x, a) - H_\beta(x, A_\eta(x))$, and terminal errors on validation state-time pairs. The action maximization is finite-grid and sampled, so this is a benchmark-level certificate-gap measurement rather than a symbolic proof; the script now also has a no-PyTorch deterministic archived replay mode for artifact validation. Table 19 shows that Exp. I can have small sampled greedy gaps under an actor-focused configuration, while inventory remains the harder action-selection case.

Targeted residual-closing rerun. Because Exp. I and inventory have the largest native residual diagnostics, we ran a stronger targeted HBO/NBO residual-closing audit after the main suite. Table 20 reports the best guarded configuration for each benchmark, with all attempted configurations in `hbo_residual_closing_audit.csv`. Exp. I uses a 2400-step critic residual/terminal run and lowers native HJB loss from 2.13×10^3 to 2.95×10^2 while reducing terminal loss from 85.2 to 42.5 and preserving held-out utility. Inventory uses a 2400-step balanced residual/terminal run and lowers native HJB loss from 6.26×10^4 to 1.20×10^4 while slightly reducing terminal loss and preserving held-out welfare. These rows materially reduce the largest sampled residual gaps; the residual-gated theorem still requires separate uniform residual and greedy-gap validation for a literal certificate.

Table 20 Targeted residual-closing HBO/NBO audit for the two largest sampled residual gaps. Held-out value uses the same primary metric as Table 18; ratios are relative to the headline performance run.

Benchmark	Configuration	Native loss	HJB Norm. RMSE	Terminal loss	Held-out
Exp. I	headline performance run	2.13×10^3	1.33	85.2	-2.7966
Exp. I	residual/terminal critic, 2400 steps	$2.95 \times (0.139 \times)$	10^2 0.403	42.5 (0.50 $\times$)	-2.7965
Inventory	headline performance run	6.26×10^4	1.84	2062	-3.6375
Inventory	balanced residual/terminal, 2400 steps	$1.20 \times (0.192 \times)$	10^4 4.25	1946 (0.94 $\times$)	-3.6374

Table 21 Inventory normalized-residual Pareto audit. The selected guarded row improves normalized RMSE while keeping native HJB, terminal loss, and held-out welfare on the favorable side of the headline run.

Configuration	Norm. RMSE	Native HJB	Terminal	Held-out	Read
headline performance run	1.843	6.26×10^4	2062	-3.6375	baseline
native-residual closure	4.25	$1.20 \times (0.192 \times)$	10^4 1946 (0.94 $\times$)	-3.6374	native-gap closure
normalized Pareto actor, 1200 steps	1.805 (0.979 $\times$)	$1.57 \times (0.251 \times)$	10^4 1916 (0.929 $\times$)	-3.6373	guarded norm. closure
scaled-loss stress, 1200 steps	1.283 (0.696 $\times$)	2.52×10^6	3.88	-3.6374	scale tradeoff

Inventory normalized-residual Pareto audit. The inventory normalized RMSE remained the sharpest residual-scale objection after the native residual-closing run. We therefore ran an inventory-only normalized-residual sweep over 26 additional PyTorch-AD configurations, including early-stop, Hamiltonian-scale-preserving, actor/critic, and scaled-residual-loss variants; the full grid is in `hbo_inventory_normalized_residual_audit.csv`. The guarded Pareto row is `inventory_norm_pareto_actor_1200`: normalized RMSE falls from 1.843 to 1.805 (0.979 $\times$), native HJB remains far below headline at 1.57×10^4 (0.251 $\times$), terminal loss falls to 0.929 $\times$, and held-out welfare is preserved. A pure scaled-loss row can drive normalized RMSE lower (1.283) but inflates native HJB; we therefore use the guarded Pareto row in the evidence ledger.

Residual scale and rollout interpretation. The native residual is computed on compact state-domain samples used by the actor-critic solver audit and is a sampled actor-induced proxy for the residual term in Theorem 2.5; it is paired with the discrete greedy audit to approximate the optimal-Hamiltonian certificate gap on the benchmark domain. Its normalized RMSE divides by the sampled Hamiltonian scale, so values near or above one mean that the residual is economically large on the diagnostic domain. A high-residual policy can still have good held-out rollout value if the policy visits a narrower subset of states or if residual errors are partly orthogonal to the action margin; conversely, a smaller residual must be read with the induced action margin and rollout distribution. For this reason the paper reports HBO/NBO as an HJB-solver route with certificate-gap evidence: the theorem gives the compact-domain certificate, while the experiments disclose residual gaps, greedy gaps, held-out rollouts, theta-free physical diagnostics, interventions, and exact-DP checks needed to judge solver performance.

Budget-normalized and seed-level audit. Budget-normalized CEM frontiers are appendix diagnostics, not the headline OR ranking. The important reading is narrow: the compact Experiment I/II frontiers are positive, the Experiment III common-budget margin is tiny (0.000091 at 442 evaluations), and

Table 22 Primary matched-budget/evaluator audit. Internal valuation objectives may include θ_t ; external/theta-free columns do not score a method by its learned valuation state.

Benchmark	Internal valuation objective	Theta-free / physical evaluator	Read
Exp. I habit	Mean-utility gap vs same-dim preference backbone: $+0.02554 \pm 0.00232$	Terminal wealth: Full 1.0301 vs same-budget mixed fixed-original 1.0374; paired Full-minus-fixed -0.00730 ± 0.00301	mixed
Exp. II hidden regime	Mean-utility gap vs fixed-pref. recurrent: -0.0000019 ± 0.0000870	Sharpe gap $+0.01827 \pm 0.01132$; drawdown gap $+0.000329 \pm 0.000177$ vs fixed-pref. recurrent	parity
Exp. III CEM privileged-ref. stress	Mean-utility gap vs privileged reference: -0.00137 ± 0.00224	Sharpe gap -0.14780 ± 0.02292 ; drawdown gap -0.00550 ± 0.00101 vs reference	appendix stress
Exp. III feasible	Mean-utility gap vs feasible fixed-pref.: $+0.01227 \pm 0.00211$	Turnover gap vs feasible fixed-pref. -0.11279 ± 0.00991 (lower is better); Sharpe gap $+0.01681 \pm 0.02102$	finite gain

full-budget heterogeneous runs are retained only as stress tests. The packaged equal-evaluation and seed-level audit CSVs validate these diagnostics without replacing Table 18 or Table 22.

Matched-budget/evaluator stress audit. Table 22 is the same-budget/parameter-count-controlled rerun, retained here as a stress rerun for the non-inventory portfolio/control experiments. It reports each internal valuation-objective comparison beside a theta-free or physical evaluator. Intervals are paired 95% half-widths over matched seed indices when the same seed schedule is available. The inventory/SKU rows, including the equal-budget bootstrap interval, are reported separately in Table 3. The aggregate audit CSV generated by `make_primary_evaluator_audit.py` is `primary_evaluator_audit.csv`. The “preference backbone” comparator is a same-dimension non-NDU architecture that keeps the valuation-state feature backbone but disables endogenous valuation control, so it is a stronger comparator than a smaller fixed-preference controller. For the one-state habit benchmark the available same-budget physical metric is terminal wealth; Sharpe, drawdown, and turnover are reported in the portfolio rows where those path metrics are defined.

Oracle, benchmark role, and scoring clarification. We reserve *oracle upper bound* for a row that optimizes the same objective over a policy class containing the feasible policy class with adequate optimization. Most packaged “oracle” rows are not such bounds. A row that observes a hidden regime or factor is called a *privileged-information reference*; a row constructed from safe incumbents is a *safe lower bound*; and an under-budgeted or narrow oracle row is a *diagnostic row*. Diagnostic and privileged-reference rows do not participate in the main ranking. Experiment I is the closest implemented analogue of the dynamic-preference SDE layer. Experiments II–III are reduced finite-policy valuation-control benchmarks in which θ_t acts as a risk/utility/turnover weight rather than a theorem-level autonomous preference-state drift; in particular, Experiment III is below the privileged factor reference in Table 18. The inventory benchmark is an OR service-pressure/order-up-to bridge; θ_t is an adaptive service target or shortage-pressure parameter, not a continuous-time FBSDE state drift. Because these internal objectives can contain θ_t , the physical evaluator and theta-free cost audits in Tables 22 and 3 are part of the primary interpretation, not a secondary afterthought.

Decision-layer certificate. The full PyTorch-AD HBO/NBO actor–critic table supplies the residual-gated HJB-solver layer for the benchmark suite, while the reduced CEM and inventory experiments supply matched-budget and external-evaluator stress tests. The latter instantiate finite benchmark decision problems with a common rollout generator, reported evaluators, named policy classes, and, for inventory, replay interventions. The package self-replay certificate separately records a deterministic seven-command regeneration of the finite neural guard chain and the six resulting artifact hashes, while the environment audit records the local Python/toolchain/dependency facts used for the lightweight package gates and the boundary around optional full regeneration. The dependency-closure audit then verifies that the lightweight gate scripts themselves do not import the optional full-regeneration packages. The next proposition states the finite-benchmark audit certificate used to read the empirical layer.

Proposition 5.1 (Finite-benchmark valuation-control audit certificate) *Fix a finite-horizon benchmark tuple*

$$\mathcal{B} = (\mathbb{P}, \mathcal{R}, J, E, \Pi_0, \Pi_F, \mathcal{I}),$$

where $\mathbb{P}$ is the exogenous rollout law, $\mathcal{R}(\pi, \omega)$ maps a policy and shock path into trajectories, J is the reported evaluator, E is an optional external physical evaluator, Π_0 is a fixed-preference/action-only policy class, Π_F is a valuation-control policy class, and $\mathcal{I}$ is a set of interventions that freeze, permute, or remove only the valuation channel. Suppose there is an embedding $\iota: \Pi_0 \rightarrow \Pi_F$ that preserves the physical action law and sets $\theta_t \equiv \bar{\theta}$, with the same J value on embedded policies. Then:

- (i) Class inclusion. $\sup_{\pi \in \Pi_F} J(\pi) \geq \sup_{\pi \in \Pi_0} J(\pi)$.
- (ii) Attained-policy lower bound. For reported policies $\hat{\pi}_F \in \Pi_F$ and $\hat{\pi}_0 \in \Pi_0$, a positive held-out margin $\hat{\Delta} = J(\hat{\pi}_F) - J(\hat{\pi}_0) > 0$ certifies a realized finite-policy improvement against that named comparator under the common rollout/evaluator; it is not, by itself, a proof of class-optimal separation.
- (iii) Uncertainty and optimizer-gap threshold. If the reporting procedure supplies a valid half-width h for the margin and a baseline optimizer-gap certificate g_0 satisfying $\sup_{\pi \in \Pi_0} J(\pi) \leq J(\hat{\pi}_0) + g_0$, then

$$\sup_{\pi \in \Pi_F} J(\pi) - \sup_{\pi \in \Pi_0} J(\pi) \geq \hat{\Delta} - h - g_0.$$

Thus $\hat{\Delta} > h + g_0$ gives a class-level separation for that benchmark/evaluator. Finite-state exact-DP audits have $h = g_0 = 0$ for the solved discrete class; CEM rows without a global optimizer-gap certificate remain comparative audits.

- (iv) Intervention relevance. If $\mathcal{T} \in \mathcal{I}$ is replayed under the same exogenous rollout law and $J(\hat{\pi}_F) > J(\mathcal{T}\hat{\pi}_F)$, the learned valuation channel is decision-relevant for $\mathcal{B}$ rather than an inactive label.

(v) External-metric ledger. *Comparisons under E are physical-performance audits. Agreement between J and E strengthens the operational reading; disagreement identifies a tradeoff, not a logical contradiction in the valuation-control claim.*

Proof. See Appendix C.8.

This certificate is stronger than a nested-class observation because it fixes what each empirical row can and cannot certify. The HBO/NBO rows train actor/value policies on the reported benchmark domains; the exact valuation-control MDP and inventory fixed/permuted- θ replays are the cleanest finite-DP/intervention applications; and the same-budget CEM rows are attained-policy and fairness audits. Together they support the benchmark-specific statuses in the compressed bridge ledger and packaged audit rather than a universal dominance claim.

Compact benchmark reading. The full PyTorch-AD HBO/NBO table is the finite-benchmark solver ledger: Full NDU is above the fixed-preference/action-only row in Experiments I–II and inventory, while Experiment III is a stress row whose original feasible comparison is parity and whose selected same-architecture sensitivity run gives only a small positive feasible gap, with the privileged factor-reference row remaining a higher hidden-information ceiling. The exact finite-state policy-class DP and the SKU exact valuation-state MDP use nested Full action sets and show Full valuation-control Bellman welfare gains; the public SKU replay shows Full beating approximate DP on valuation-adjusted welfare while approximate DP remains slightly lower on theta-free external cost. These rows support the generalized-RL claim on valuation-control objectives while making the physical-cost tradeoff explicit. Detailed metric definitions, architectures, seed schedules, heterogeneous-budget stress tests, oracle repair, and extra inventory audits are in Appendix B.

Comparator definition. The same-budget “preference backbone” in Experiment I has the same valuation-state feature backbone and the same parameter dimension as Full NDU, but the endogenous valuation-control output is disabled and the preference weights are fixed at the comparator values. It therefore tests whether the extra valuation-control channel, not merely a larger feature map, accounts for the reported internal-objective gain.

Falsification and failure modes. The framework would be falsified as a useful service-pressure modeling layer in any benchmark where matched-budget action-only or fixed-service policies dominate Full NDU simultaneously on valuation welfare, theta-free cost, and service metrics; where fixing or permuting the learned θ_t leaves the policy value unchanged; where the endogenous valuation channel improves its internal score only by degrading fill rate, backlog, or external cost beyond the reported uncertainty intervals; or where exact-DP/nested-class checks show no valuation-control gain after controlling for information and action sets. The current evidence does not rule out those failures in other domains. It instead shows where the valuation-control channel is useful, where it is only a tradeoff, and where it remains inconclusive.

6. Conclusion

Full NDU is a structured augmented-state valuation-control framework in which operational actions and selected valuation parameters are optimized together. The formal contribution is deliberately bounded: standard fixed- ε FBSDE/HJB verification applies only under the stated selector, projection, and weak-coupling restrictions; vanishing viscosity requires a compatible shadow-price selection; and the HJB residual certificate is an ideal compact-domain statement rather than a blanket guarantee for all sampled experiments. Within that boundary, the NDU-specific insight is the degenerate preference-block observability problem and the resulting conditioning advantage of learning direct shadow prices rather than inverting Z_u as auxiliary preference noise vanishes.

The empirical conclusion is benchmark-qualified but now inventory-centered. The inventory/service-level application is the primary OR evidence: Full NDU improves the main valuation-welfare benchmark, is meaningful under fixed/permuted- θ interventions, beats approximate DP on the public SKU valuation-welfare objective while exposing the small theta-free cost tradeoff, and is better in the exact valuation-control MDP when θ is included in the Bellman state/action recursion. The paper's claim is therefore deliberately narrower: endogenous valuation control is a useful OR modeling layer, especially for service-pressure inventory decisions, and it exposes a concrete shadow-price observability issue. The HBO/NBO portfolio/hidden-regime/allocation rows, same-budget CEM runs, oracle-reconciliation rows, intervention checks, SKU replay, exact-DP checks, and conditioning audits are reproducible support for that narrow claim rather than a universal dominance theorem or a certified HJB-solver claim.

References

- Abbeel, Pieter, Andrew Y Ng. 2004. Apprenticeship learning via inverse reinforcement learning. *Proceedings of the twenty-first international conference on Machine learning*. 1.
- Bertsekas, Dimitri P. 2012. *Dynamic Programming and Optimal Control*. 4th ed. Athena Scientific.
- Campbell, John Y., John H. Cochrane. 1999. By force of habit: A consumption-based explanation of aggregate stock market behavior. *Journal of Political Economy* **107**(2) 205–251.
- Chen, Tian Qi, Yulia Rubanova, Jesse Bettencourt, David K. Duvenaud. 2018. Neural ordinary differential equations. *Advances in Neural Information Processing Systems*, vol. 31. 6572–6583.
- Duffie, Darrell, Larry G Epstein. 1992. Stochastic differential utility. *Econometrica* **60**(2) 353–394.
- E, Weinan, Bing Yu. 2018. The deep ritz method: a deep learning-based numerical algorithm for solving variational problems. *Communications in Mathematics and Statistics* **6**(1) 1–12.
- El Karoui, Nicole, Shige Peng, Marie Claire Quenez. 1997. Backward stochastic differential equations and applications. *Mathematical finance* **7**(1) 1–71.

- Han, Jiequn, Arnulf Jentzen, Weinan E. 2018. Solving high-dimensional partial differential equations using deep learning. *Proceedings of the National Academy of Sciences* **115**(34) 8505–8510.
- Hansen, Lars Peter, Thomas J. Sargent. 2008. *Robustness*. Princeton University Press.
- Huang, Zhipeng, Balint Negyesi, Cornelis W. Oosterlee. 2025. Convergence of the deep bsde method for stochastic control problems formulated through the stochastic maximum principle. *Mathematics and Computers in Simulation* **227** 553–568.
- Kahneman, Daniel, Amos Tversky. 1979. Prospect theory: An analysis of decision under risk. *Econometrica* **47**(2) 263–291.
- Li, Zongyi, Nikola Kovachki, Kamyar Azizzadenesheli, Burigede Liu, Kaushik Bhattacharya, Andrew Stuart, Anima Anandkumar. 2020. Fourier neural operator for parametric partial differential equations. *arXiv preprint arXiv:2010.08895* .
- Ma, Jin, Philip Protter, Jiongmin Yong. 1994. Solving forward-backward stochastic differential equations explicitly: a four step scheme. *Probability Theory and Related Fields* **98**(3) 339–359.
- Ng, Andrew Y, Stuart Russell. 2000. Algorithms for inverse reinforcement learning. *Icml*, vol. 1. 2.
- Pardoux, Etienne, Shige Peng. 1990. Adapted solution of a backward stochastic differential equation. *Systems & Control Letters* **14**(1) 55–61.
- Peng, Shige. 1990. A general stochastic maximum principle for optimal control problems. *SIAM Journal on Control and Optimization* **28**(4) 966–979.
- Peng, Shige. 1993. Coupled forward-backward stochastic differential equations and their associated systems of algebraic equations. *Publ. Res. Inst. Math. Sci.* **29** 4.
- Porteus, Evan L. 2002. *Foundations of Stochastic Inventory Theory*. Stanford University Press.
- Powell, Warren B. 2011. *Approximate Dynamic Programming: Solving the Curses of Dimensionality*. 2nd ed. Wiley.
- Scarf, Herbert. 1960. The optimality of (s, S) policies in the dynamic inventory problem. Kenneth J. Arrow, Samuel Karlin, Patrick Suppes, eds., *Mathematical Methods in the Social Sciences*. Stanford University Press, 196–202.
- Shirani Faradonbeh, Mohamad Kazem, Mohamad Sadegh Shirani Faradonbeh. 2023. Online reinforcement learning in stochastic continuous-time systems. *Proceedings of Thirty Sixth Conference on Learning Theory, Proceedings of Machine Learning Research*, vol. 195. 612–656.
- Sims, Christopher A. 2003. Implications of rational inattention. *Journal of monetary Economics* **50**(3) 665–690.
- Sutton, Richard S, Andrew G Barto. 2018. *Reinforcement learning: An introduction*. MIT press.
- Whittle, Peter. 1990. *Risk-Sensitive Optimal Control*. Wiley.
- Yong, Jiongmin, Xun Yu Zhou. 1999. *Stochastic controls: Hamiltonian systems and HJB equations*. Springer Science & Business Media.
- Zipkin, Paul H. 2000. *Foundations of Inventory Management*. McGraw-Hill.

Table 23 Explicit constant audit for the sufficient small-coupling condition in Proposition 3.8. The bound is intentionally conservative.

Instance	T	ε	β_h	λ	Λ_ε	C_X	C_B	$q_\varepsilon(T)$
Active small-coupling NDU	1.0	0.50	0.10	2.00	0.0050	1.82	6.68	6.8×10^{-5}
Strong-coupling diagnostic	1.0	0.20	0.80	0.50	2.0239	4.95	32.89	3.5×10^2

Table 24 Theory-covered fixed- ε HJB/FBSDE sanity example under compact/clipped controls. The sufficient coupling bound is conservative; RMSE/residual values are computed on a (t, c, u) certificate grid where clipping is inactive.

Instance	ε	L_α	Λ_ε	q_ε	Small?	Direct- P	RMSE	HJB residual
LQ fixed- ε NDU	0.08	0.50	1.6667	0.1494	yes	3.0×10^{-17}		1.6×10^{-16}

Appendix A: Deferred Theory Audit Tables

The following tables are generated from the reproducibility bundle and support the restricted fixed- ε verification discussion in the main text. They are audit/certificate tables, not additional broad existence claims.

Appendix B: Supplementary Experimental Reporting

The main-text tables report the compact paired/bootstrap evaluator layer; this appendix gives the fuller audit layer without making the file list part of the scientific claim. The artifact README and manifest are the authoritative map from tables to CSVs and regeneration scripts. In substance, the appendix records: (i) the full PyTorch-AD HBO/NBO actor-critic diagnostic and residual-focused reruns; (ii) same-budget, oracle-repair, and CEM stress audits for the non-inventory experiments; (iii) inventory/SKU/exact-DP and public-trace OR checks; and (iv) seed schedules, controller feature maps, rollout conventions, and metric definitions. The intended reading is unchanged from the main text: inventory/service-pressure evidence is the primary OR result, while portfolio, hidden-regime, high-dimensional allocation, and CEM rows are stress or mechanism audits.

B.1. Metrics

All welfare metrics are computed from held-out reporting rollouts after the relevant solver or optimizer finishes: HBO/NBO actor policies for Table 18, and CEM-selected policies for the matched-budget tables. Mean utility is the simulation average of the cumulative benchmark objective U_i over held-out episodes/seeds. The column label CEW denotes the exponential utility-equivalent welfare transform $\text{CEW} = \exp(\bar{U}) - 1$. Because U_i may include service penalties, adjustment costs, and benchmark-specific regularizers, CEW should not be read as a standard financial certainty-equivalent return except in log-return subcases; negative CEW in Experiment I means that the severe stopped benchmark produces a large negative log-welfare equivalent, not that wealth itself is negative. Sharpe is the annualized mean/std ratio of simulated portfolio returns. Maximum drawdown is $\min_t (W_t / \max_{s \leq t} W_s - 1)$, so values closer to zero are better. Turnover is average absolute portfolio reallocation. Adaptation lag is the post-shock first substantial response time divided by horizon length; values near 0.5 in Experiments I–II mean the response criterion is saturated over the half-horizon after the midpoint

shock. θ -stability is the standard deviation of the learned θ_t path; it is zero by construction for fixed-preference baselines. In the inventory benchmark, average cost includes holding, shortage, order, capacity, and θ -adjustment costs; fill rate is total filled demand divided by total demand; service breaches count periods with fill rate below 0.92; and lower cost, lower backlog, and fewer breaches are better.

B.2. Numerical Methodology

The reported benchmark evidence uses one primary OR inventory ledger and two numerical diagnostic engines. First, the HBO/NBO diagnostic engine trains fully trainable bounded actor networks and a value critic with the composite $L_{\text{PERF}} + \omega L_{\text{HJB}}$ objective, terminal boundary loss, compact state-domain sampling, PyTorch-AD value gradients/Hessians, PyTorch-AD actor Hamiltonian gradients, and 30 held-out seed-block rollout evaluation. Second, the earlier finite-dimensional CEM policy search over bounded feature maps remains as a matched-budget robustness and external-evaluator audit. The hybrid characteristic-Galerkin / Deep-BSDE construction below is retained as the continuous-time value-network template for the formal FBSDE system (12)-(13). The package also includes a small neural shadow-price validation in a solvable one-dimensional preference-state benchmark to test the direct- P versus Z -inversion conditioning mechanism.

For the optional solver template, one may approximate the value function $V(t, x)$ directly using a neural network $V_\phi(t, x)$ with parameters ϕ . Crucially, it need not employ a separate network to approximate the volatility process Z_t . Instead, consistent with the vanishing-viscosity discussion around Propositions 3.11–3.12 and Remark 3.17, $P_t^\phi = \nabla_x V_\phi(t, X_t)$ and $Z_t = \Sigma_\varepsilon^\top P_t^\phi$ can be calculated structurally via automatic differentiation. This value-network route makes the joint control signal derive from a conservative vector field. By contrast, the one-dimensional shadow-price bridge and small neural validation below use supervised analytic oracle targets for P_u ; they test the conditioning of direct- P recovery and do not by themselves enforce HJB consistency, adjoint dynamics, or conservative-vector-field structure in a general model.

B.2.1. Singular Feature Encoding For specifications with relative-subsistence or habit-formation utility, the singularity is induced by the surplus variable. In particular, when utility contains a term of the form $(c - h)^{1-\gamma}/(1 - \gamma)$ with $\gamma > 0$, the marginal utility behaves like $(c - h)^{-\gamma}$ and therefore explodes as $c \downarrow h$. Standard feedforward networks with Lipschitz activations (ReLU, Tanh) struggle to capture this boundary-layer behavior. We therefore introduce a fixed Singular Feature Map built from a positive boundary-distance coordinate $s(c, u) > 0$ acting as a hard-coded inductive bias:

$$\zeta_{\text{sing}}(s) := \left[s, \log(s + \delta_{\text{num}}), \frac{1}{s + \delta_{\text{num}}} \right]^\top, \quad (42)$$

where $\delta_{\text{num}} \approx 10^{-6}$ is a numerical stability constant. The network architecture may then incorporate this feature block through $V_\phi(t, c, u) := \mathcal{N}_\phi(t, \zeta_{\text{sing}}(s(c, u)), u)$, where $\mathcal{N}_\phi$ is a modified ResNet

with Swish activation functions. This should be read as the theory-oriented value-network ansatz. The released HBO/NBO and CEM benchmark experiments use bounded feature maps, compact state/action sensors, and compact control parameterizations rather than attempting to numerically certify a global Lipschitz constant for an unrestricted deep architecture. In Experiment I below, the concrete choice of s is the distance to the habit-induced admissibility boundary.

B.2.2. Adjoint Consistency and Feedback Control In standard Deep BSDE approaches, Z_t is an independent output. However, due to the degeneracy of the preference dynamics ($\sigma_h \rightarrow 0$), an independent Z -network fails to learn the gradient in the u -direction (the Observability Paradox).

In our framework, we define the shadow price P_t^ϕ and the volatility Z_t^ϕ strictly as derived quantities:

$$P_t^\phi(X_t) := \nabla_x V_\phi(t, X_t), \quad (43)$$

$$Z_t^\varepsilon(X_t) := \Sigma_\varepsilon(X_t)^\top P_t^\phi(X_t). \quad (44)$$

Computationally, Eq. (43) requires one backward pass through the network graph. The optimal joint control is then computed via the Hamiltonian argmax:

$$\alpha_t^*(X_t) \in \arg \max_{\alpha \in \mathcal{A} \times \Theta} \left\{ (P_t^\phi)^\top b(X_t, \alpha) + f(X_t, \alpha) - \lambda \mathcal{L}(X_t, \alpha) \right\}. \quad (45)$$

In reduced benchmarks, one control channel may admit a closed-form update, but the full-NDU experiments treat this as a genuinely joint inner optimization step.

Theory-to-implementation gap and practical stabilization. The abstract assumptions from Sections 3–4, especially the global selector-Lipschitz condition used in the weak-coupling bounds, are not numerically certified automatically by a generic neural ansatz. In the released benchmark implementation, stabilization is instead enforced by design through bounded feature encodings and clipped control maps: state features are log-scaled, normalized, or clipped; risky-action and leverage channels are passed through tanh or explicit clipping; preference controls are constrained through sigmoid maps to compact intervals; belief updates are clipped away from the simplex boundary; and portfolio weights are normalized by an ℓ^1 -type scaling. The HBO/NBO engine trains compact-output actor and value MLPs with clipped controls and gradient clipping, while the CEM stress engine uses bounded variance updates rather than unconstrained gradient descent. The reproducibility bundle now makes this implementation layer checkable through the bounded-architecture Lipschitz budget audit `make_bounded_architecture_lipschitz_budget_audit.py`, whose seven rows cover state feature clipping, actor tanh/clipping, preference sigmoid maps, belief-simplex clipping, portfolio ℓ^1 normalization, value-MLP gradient clipping, and CEM variance caps; the maximum recorded budget product is 7.50 against a threshold of 12.0. The companion stress-margin audit `make_bounded_architecture_budget_stress_audit.py` applies nominal, $1.25\times$ slope, and $1.50\times$

radius perturbations to the same seven components and retains a positive minimum margin of 0.75. The cross-artifact consistency audit `make_neural_guard_cross_artifact_consistency_audit.py` then joins the validation-cover, replay, empirical-envelope, interval, sensitivity, budget, and stress artifacts by experiment id and records that both large sampled neural rows pass the same finite guard chain. The refinement/stress audit `make_neural_guard_refinement_stress_audit.py` applies current-cover, half-radius cover-refinement, and refined-cover slope-stress scenarios to that joined table; all six rows retain positive interval-sensitivity slack, positive stress margin, and the explicit validation-cover_not_global_lipschitz route. The holdout/bootstrap stability audit `make_neural_guard_holdout_bootstrap_stability_audit.py` then applies paired-seed bootstrap, leave-block-out, and tail-stress bootstrap floors to the same rows; all six rows retain positive sensitivity-slack floors, nonnegative held-out-delta floors, and the same route label. The concordance audit `make_neural_guard_concordance_audit.py` collapses those three schemes by experiment id; both large rows have three positive votes out of three and retain the same route label. The leave-one-out robustness audit `make_neural_guard_leave_one_out_robustness_audit.py` then removes one prespecified guard family at a time; all twelve experiment-by-omission rows keep five remaining guard families, positive refinement and holdout slack, nonnegative held-out delta, positive budget-stress margin, and the same validation-cover route. Finally, `make_neural_bounded_composition_envelope_certificate.py` joins the interval, sensitivity, bounded-budget, stress-margin, concordance, and leave-one-out layers into two bounded-composition envelope rows; both rows pass with positive composed-envelope margins under the explicit bounded-composition_compact-domain_global-envelope_not_unrestricted route. The companion frontier audit `make_neural_guard_stress_frontier_audit.py` identifies the active bottleneck: the current-cover limited multipliers are 1.067 for Exp. I and 1.035 for inventory, while half-radius/refined-cover sensitivity frontiers are 2.028 and 1.883; the refined-cover limiter is the bounded-budget stress row in both cases. The compact numerical ledger records replay, stability, empirical-envelope, interval-sensitivity, budget max 7.50, stress margin 0.75, Exp. I ratio 0.139; inventory 0.192; tail 1.49, 4.09; interval 0.525, 1.625; cross-artifact 2/2 pass; refinement/stress 6/6 pass; holdout/bootstrap 6/6 pass; concordance 2/2 pass; leave-one-out 12/12 pass; bounded-composition 2/2 pass; stress-frontier 2/2 pass. These devices keep the effective control law in a compact numerical regime and mitigate practical blow-up of the shadow-price channel, but they should be read as implementation stabilizers and finite budget/stress/consistency/refinement/bootstrap/concordance/leave-one-out/composition/frontier records rather than as a literal numerical proof that Assumption 3.1(ii) or the small-coupling condition is enforced exactly by the experiment code.

Direct observability bridge audit. To connect the observability argument to a numerical diagnostic, the package includes a one-dimensional analytic shadow-price audit. The audit fixes an analytic quadratic value profile in the preference coordinate, computes the oracle $P_u = \partial_u V$ in closed form,

Table 25 Shadow-price observability bridge audit. RMSE is measured against the analytic HJB shadow-price oracle on a one-dimensional preference grid.

ε	<i>Z</i> -inversion θ RMSE	Direct- <i>P</i> θ RMSE	Ratio
10^{-1}	0.00768	0.00344	2.23
10^{-2}	0.02426	0.00344	7.06
10^{-3}	0.07671	0.00344	22.33
10^{-4}	0.21768	0.00344	63.36

Table 26 Small neural shadow-price conditioning validation. RMSE is measured against the analytic HJB oracle; each learned row averages 30 seeds with the same network size and training budget.

ε	Direct- <i>P</i> θ RMSE	<i>Z</i> -inversion θ RMSE	Ratio
10^{-1}	0.00436	0.00754	1.73
10^{-2}	0.00436	0.01325	3.04
10^{-3}	0.00436	0.03079	7.05
10^{-4}	0.00436	0.08726	19.99

constructs $Z_u = \sqrt{2\varepsilon}P_u$, and compares two noisy recovery routes: inverting Z_u versus supervising P_u directly. As Table 25 shows, the *Z*-inversion θ -error grows rapidly as ε shrinks, while the direct-*P* error remains essentially constant. This is a conditioning diagnostic for the Observability Paradox; the separate HBO/NBO table is the benchmark-level actor-critic solve.

Small neural conditioning validation. The package further includes a 30-seed neural conditioning validation on a solvable one-dimensional preference-state benchmark with an analytic HJB/shadow-price oracle. The two learned estimators use the same one-hidden-layer tanh network (width 24), the same 1536 training points, and the same 900 optimizer steps per seed; both are supervised against oracle P_u or Z_u targets rather than trained through a full HJB residual. As Table 26 shows, direct learning of P_u keeps the induced θ RMSE near 0.00436 as ε shrinks, whereas learning $Z_u = \sqrt{2\varepsilon}P_u$ and then inverting it amplifies target noise by factors from 1.73 to 19.99. This validates the conditioning advantage behind direct-*P* recovery, while Table 18 gives the separate HBO/NBO benchmark audit.

B.2.3. Optional Local Temporal Consistency Template The HBO/NBO benchmark solver uses the Hamiltonian/HJB-residual value-network route; the continuous-time value-network template corresponding to the FBSDE notation is local and implicit in the discount term:

$$V(t, X_t) \approx \mathbb{E}_t [V(t + \Delta\tau, X_{t+\Delta\tau}) + (f_t - \lambda\mathcal{L}_t - \rho V(t + \Delta\tau, X_{t+\Delta\tau})) \Delta\tau]. \quad (46)$$

For a value network V_ϕ , the corresponding residual over a spatial distribution ν can be written as

$$\mathcal{J}(\phi) = \mathbb{E}_{t \sim \mathcal{U}[0, T], X_t \sim \nu} [|V_\phi(t, X_t) - \mathcal{T}_{target}|^2] + \beta_{term} \mathbb{E}_{X_T} [|V_\phi(T, X_T) - \Psi(X_T)|^2], \quad (47)$$

where $\mathcal{T}_{target} := V_\phi(t + \Delta\tau, \hat{X}_{t+\Delta\tau})$

$$\begin{aligned} &+ \left(f(\hat{X}_t, \alpha_t^*) - \lambda\mathcal{L}(\hat{X}_t, \alpha_t^*) - \rho V_\phi(t + \Delta\tau, \hat{X}_{t+\Delta\tau}) \right) \Delta\tau \\ &- (Z_t^\varepsilon(\hat{X}_t))^\top \Delta\mathbf{W}_t. \end{aligned} \quad (48)$$

Here, $\hat{X}$ denotes detached rollout variables, $P_t^\phi = \nabla_x V_\phi(t, \hat{X}_t)$, and $Z_t^\varepsilon = \Sigma_\varepsilon^\top P_t^\phi$. The target is a one-step conditional BSDE increment and is weakly consistent as $\Delta\tau \rightarrow 0$ in the standard Deep-BSDE/DGM sense (Han et al., 2018). The martingale term is retained for identity matching, not as a

blanket variance-reduction claim; control-sensitive gradients are routed through the shadow price and require the usual Hessian-vector products in automatic differentiation.

B.3. Experiment I: Low-Dimensional Joint Habit-Portfolio Control

The first full-NDU experiment is a low-dimensional benchmark in which both overt action and endogenous preference control are active. The purpose is to test whether a fully joint NDU controller can learn a competitive policy when both the external decision channel and the internal preference channel matter.

B.3.1. Problem Specification We use a joint habit-portfolio-control problem with state $X_t = [W_t, h_t]^\top$ and joint control $\alpha_t = (\pi_t, m_t, \theta_t)$, where π_t is the risky-asset allocation, m_t is consumption, and θ_t controls habit adaptation:

$$dW_t = (rW_t + \pi_t W_t (\mu_t - r) - m_t) dt + \pi_t W_t \sigma_t dB_t, \quad (49)$$

$$dh_t = \theta_t (m_t - h_t) dt, \quad h_0 > 0. \quad (50)$$

The objective is a stopped recursive-utility problem with misery boundary $\tau_{\text{mis}} = \inf\{t \geq 0 : m_t \leq h_t\}$, so both the overt action channel (π_t, m_t) and the internal preference-control channel θ_t are optimized online.

B.3.2. Experimental Setup

Code-level setup. The released simulator uses $W_0 = 1$, $h_0 = 0.08$, 40 steps, and a deterministic midpoint return/volatility shock from $(\mu, \sigma) = (0.11, 0.16)$ to $(-0.05, 0.36)$ with $r = 0.02$. Full NDU, augmented-state RL, fixed-preference RL, and preference-only controls use the bounded feature maps and CEM budgets recorded in `manifest.json`; the key distinction is that only Full NDU jointly learns (π_t, m_t, θ_t) , while fixed-preference rows freeze $\theta_t \equiv 0.25$. Held-out 30-seed evaluations report mean utility, CEW, terminal wealth, and utility gap; saturated ruin/adaptation diagnostics are kept in the CSVs rather than repeated in prose.

B.3.3. Results The primary Experiment I result is the matched-budget row in Table 22: Full NDU improves mean utility by +0.02554 over the best same-budget non-NDU row. The original full-budget table and mechanism figures remain appendix stress evidence only; they show that joint action/valuation control stays competitive in a deliberately severe stopped calibration, while the matched-budget table carries the fairness claim.

B.4. Experiment II: Joint Control under Regime Uncertainty

The second benchmark studies hidden regime switching. The agent does not observe the crisis state directly; instead it receives a noisy signal and must jointly manage overt risk exposure and endogenous preference control while filtering the latent state online.

B.4.1. Problem Specification The environment contains a latent crisis regime that raises volatility and lowers expected returns after a mid-horizon structural break. The controller chooses a risky allocation π_t together with an endogenous preference-control variable θ_t , so the full NDU policy must coordinate filtering, de-risking, and internal preference adjustment.

B.4.2. Experimental Setup

Code-level setup. Experiment II uses 48 steps, initial wealth 1, initial crisis belief 0.15, a two-state hidden Markov crisis process with a midpoint stress shift, noisy regime signals with standard deviation 0.50, and risky-return states (0.10, 0.18) versus (−0.04, 0.58). Feasible methods observe the same bounded time/wealth/signal/belief/return feature vector; Full NDU learns preference, belief-update, and action maps, while the fixed-preference recurrent row freezes $\theta_t \equiv 0.25$. The narrow oracle row observes the latent regime and is treated as a privileged-information diagnostic, with the repaired safe lower bound reported separately in Table 29.

B.4.3. Results The matched-budget rerun gives practical parity: Full NDU records 0.011256 mean utility versus 0.011257 for the same-dimension fixed-preference recurrent controller (gap −0.000002), while Table 3 keeps the physical diagnostics visible. Appendix Table 29 repairs the privileged-reference row by adding a safe latent-regime lower bound at 0.011374 mean utility (gap 0.000127 above Full), confirming that privileged information is at least competitive and that the main Experiment II claim is feasible non-privileged parity rather than oracle dominance.

B.5. Experiment III: High-Dimensional Joint-Control Benchmark

The third benchmark tests whether full NDU scales to a higher-dimensional portfolio problem with a latent factor, observed signals, and joint control over allocation and endogenous preference adaptation.

B.5.1. Problem Specification Experiment III uses a 12-asset allocation environment with one latent common factor, one noisy observed signal, and an endogenous preference-control channel. The factor follows an AR(1)-type law

$$F_{t+1} = 0.92F_t + 0.18\xi_t,$$

and the observed signal is a noisy persistent proxy,

$$S_{t+1} = 0.4S_t + F_{t+1} + 0.25\eta_t.$$

The 12 assets have deterministic alpha loadings linearly spaced from 0.04 to 0.09, factor loadings from 0.6 to 1.6, and idiosyncratic volatilities from 0.10 to 0.22. Portfolio return is the weighted sum of the asset returns minus a linear turnover cost of $0.0025 \sum_i |w_{i,t} - w_{i,t-1}|$, so the benchmark is designed to test whether endogenous preference adjustment improves the welfare–turnover tradeoff rather than only raw expected return. The period utility used in the simulator is

$$(1 + 0.40\theta_t) \log(1 + r_t^p) - 0.07(\theta_t - 0.25)^2 - 0.0015 \text{turnover}_t,$$

aggregated over the 50-step episode.

Table 27 Appendix stress-test summary for Experiment I under the original heterogeneous-budget low-dimensional joint habit-portfolio benchmark

Method	Mean Utility	CEW	Terminal Wealth	Utility Gap
Full NDU (joint)	-2.7284	-0.9347	1.0345	0.0000
Augmented-state RL	-2.7607	-0.9368	1.0271	0.0323
Fixed-preference RL	-2.7642	-0.9370	1.0192	0.0359
Preference-only	-2.7661	-0.9371	0.9568	0.0377

Table 28 Appendix stress-test summary for Experiment II under the original heterogeneous-budget hidden-regime joint-control benchmark

Method	Mean Utility	CEW	Sharpe	Maximum Drawdown	Adaptation Lag	Detection Accuracy
Full NDU (joint)	0.0112	0.0113	1.4392	-0.0001	0.5000	0.6089
Preference-only	0.0065	0.0065	1.4198	-0.0025	0.5000	0.6048
Fixed-preference recurrent RL	-0.0006	-0.0006	0.9196	-0.0095	0.4995	0.4472
Narrow oracle diagnostic	-0.0259	-0.0254	0.3454	-0.0460	0.3178	1.0000

Table 29 Appendix Experiment II strengthened-oracle repair audit. The matched-budget safe oracle is a lower bound for a privileged class containing the no-risky-exposure incumbent under the Full-NDU evaluation budget.

Method/audit row	Evaluations	Mean utility	CEW	Drawdown	Detection
Full NDU reported	2244	0.011247	0.011311	-0.000066	0.6089
Original narrow oracle	360	-0.025905	-0.025427	-0.045954	1.0000
Matched-budget safe oracle lower bound	2244	0.011374	0.011439	0.000000	1.0000

B.5.2. Experimental Setup

Code-level setup. Experiment III uses a 50-step, 12-asset factor allocation environment with an AR(1) latent factor, a noisy persistent signal, deterministic alpha/loading/volatility grids, and linear turnover cost. All non-oracle policies observe the same bounded time/wealth/signal/belief feature vector. Full NDU adds a preference-control map to the belief and allocation maps; fixed-preference joint RL freezes θ_t , while oracle-factor and signal-only rows define information-tiered action comparators. CEM budgets, dimensions, and training-frontier histories are fixed in the packaged CSVs and manifest.

B.5.3. Results Experiment III is the high-dimensional welfare–turnover stress test. In the same-budget rerun, Full NDU beats the feasible non-oracle fixed-preference class by +0.01227 mean utility but trails the privileged reference by 0.00137; at the 442-evaluation frontier the best-so-far values are Full 0.098845, oracle-factor 0.098754, fixed-preference 0.073853, and signal-only 0.070357, so the margin is positive but tiny. The bounded conclusion is therefore competitive feasible non-oracle valuation control with favorable utility–turnover behavior, not privileged-oracle dominance.

Appendix stress-test tables. The following original heterogeneous-budget/full-budget numerical summaries are retained for mechanism interpretation and stress-test transparency only. They are not the main empirical ranking; Table 22 is the primary same-budget result. The qualitative mechanism figure files remain in the artifact package, but their quantitative content is represented here by tables and checked CSVs to keep the blind review PDF within the venue footprint without dropping any numerical claim.

Reproducibility paths. The anonymous artifact is intended to be the submitted zip, not an external public link inside the blind PDF. Table 31 lists the three review paths; every main table row is traceable to a CSV and a regeneration script named in `README.md` and checked by `manifest.json`.

Table 30 Appendix stress-test summary for Experiment III under the original heterogeneous-budget high-dimensional joint-control benchmark

Method	Mean			Maximum		θ
	Utility	CEW	Sharpe	Drawdown	Turnover	
Full NDU (joint)	0.1145	0.1213	1.4479	−0.0662	0.2327	0.0186
Oracle-factor joint-action RL	0.0923	0.0968	1.4941	−0.0579	0.3751	0.0000
Fixed-preference joint-action RL	0.0706	0.0732	1.3227	−0.0464	0.2954	0.0000
Signal-only joint-action RL	0.0642	0.0663	1.2051	−0.0625	0.3998	0.0000

Table 31 Reproducibility paths. The lightweight path is the intended anonymous artifact gate; full reruns are optional because HBO/NBO and CEM training are much slower.

Path	Command / files	Checks	Runtime role
One-key checker	<code>check_ndu_reproducibility.py</code> from the submission root	hashes, row counts, seed coverage, headline values, residual diagnostics, and manuscript sentinels	routine artifact gate
Light regeneration	listed audit-builder scripts	regenerates derived CSVs, exact-DP rows, SKU panel, policy surface, and residual table from packaged inputs	seconds to minutes
Full regeneration	<code>run_source_nbo_full_experiments.py</code> and <code>run_full_ndu_all.py</code>	reruns HBO/NBO actor–critic and CEM/evaluator experiments with listed seeds, budgets, hardware notes, and dependencies	optional expensive audit

Computational reproducibility bundle. The `reproducibility/` directory contains the scripts, CSVs, manifest, checker output, and lightweight checker for the full benchmark suite. The blind PDF deliberately names only the three review paths in Table 31; `README.md`, `manifest.json`, and `check_ndu_reproducibility.py` are the authoritative file list, regeneration guide, and non-rerun validation gate. The checker validates hashes, row counts, seed coverage, headline values, residual diagnostics, direct- P/Z conditioning rows, inventory/SKU/exact-DP rows, and manuscript sentinels. The anonymous bundle has no DOI assigned inside the blind paper; for submission it should be uploaded as the anonymous reproducibility artifact.

B.6. Exact Code-Level Experiment Configuration

Reproducible experiment protocol. All portfolio experiments use the packaged 30-seed schedule (stride 1000), Gaussian CEM optimizer, held-out reporting shift, and method-specific budgets in Table 32; `reproducibility/scripts/generate_ndu_experiments.py`, `manifest.json`, and the checker are the authoritative code-level specification. Experiment I uses `simulate_exp1` on wealth–habit features, Experiment II uses `simulate_exp2` under partial observation with a separate privileged-regime diagnostic with a matched-budget lower bound, and Experiment III uses `simulate_exp3` on the 12-asset factor benchmark with `exp3_training_frontier.csv`. The appendix tables below report the held-out 30-seed uncertainty summaries; training-frontier rows are separate best-so-far optimization-history diagnostics, not held-out seed intervals.

Table 32 Comparator protocol and fairness summary for the reported 30-seed benchmark suite.

Exp.	Method	Observed information	Action	Pref.	Dim.	Evaluations/seed	Seeds
I	Full NDU (joint)	augmented wealth-habit state	yes	yes	15	$18 \times 46 = 828$	30
I	Fixed-preference RL	original-state features only	yes	no	6	$8 \times 24 = 192$	30
I	Augmented-state RL	augmented wealth-habit state	yes	no	10	$11 \times 32 = 352$	30
I	Preference-only	augmented wealth-habit state	heuristic	yes	5	$8 \times 24 = 192$	30
II	Full NDU (joint)	matched partial observation	yes	yes	18	$33 \times 68 = 2244$	30
II	Fixed-preference recurrent RL	matched partial observation	yes	no	12	$17 \times 44 = 748$	30
II	Narrow oracle diagnostic	privileged latent regime	yes	no	6	$12 \times 30 = 360$	30
II	Preference-only	matched partial observation	heuristic	yes	12	$17 \times 44 = 748$	30
III	Full NDU (joint)	signal + belief features	yes	yes	19	$23 \times 54 = 1242$	30
III	Fixed-preference joint RL	signal + belief features	yes	no	12	$13 \times 34 = 442$	30
III	Oracle-factor joint RL	privileged latent factor	yes	no	6	$9 \times 24 = 216$	30
III	Signal-only joint RL	observed signal only	yes	no	6	$9 \times 24 = 216$	30

Table 33 Experiment I uncertainty summary from 30-seed benchmark runs.

Method	Mean utility (95% CI)	CEW (95% CI)	Terminal wealth (95% CI)	Adaptation lag (95% CI)
	-2.7284	-0.9347	1.0345	0.5000
Full NDU (joint)	[-2.7302, -2.7266]	[-0.9348, -0.9346]	[1.0299, 1.0392]	[0.5000, 0.5000]
	-2.7607	-0.9368	1.0271	0.5000
Augmented-state RL	[-2.7633, -2.7581]	[-0.9369, -0.9366]	[1.0217, 1.0325]	[0.5000, 0.5000]
	-2.7642	-0.9370	1.0192	0.5000
Fixed-preference RL	[-2.7668, -2.7617]	[-0.9371, -0.9368]	[1.0146, 1.0238]	[0.5000, 0.5000]
	-2.7661	-0.9371	0.9568	0.5000
Preference-only	[-2.7677, -2.7644]	[-0.9372, -0.9370]	[0.9532, 0.9604]	[0.5000, 0.5000]

Table 34 Experiment II uncertainty summary from 30-seed benchmark runs.

Method	Mean utility (95% CI)	CEW (95% CI)	Sharpe (95% CI)	Max drawdown (95% CI)	Adaptation lag (95% CI)	Detection acc. (95% CI)
	0.0112	0.0113	1.4392	-0.0001	0.5000	0.6089
Full NDU (joint)	[0.0112, 0.0113]	[0.0112, 0.0114]	[1.4312, 1.4472]	[-0.0002, 0.0001]	[0.5000, 0.5000]	[0.6050, 0.6127]
	0.0065	0.0065	1.4198	-0.0025	0.5000	0.6048
Preference-only	[0.0055, 0.0074]	[0.0055, 0.0074]	[1.3625, 1.4771]	[-0.0034, -0.0016]	[0.5000, 0.5000]	[0.6018, 0.6078]
	-0.0006	-0.0006	0.9196	-0.0095	0.4995	0.4472
Fixed-preference recurrent RL	[-0.0021, 0.0009]	[-0.0021, 0.0009]	[0.8220, 1.0172]	[-0.0109, -0.0081]	[0.4989, 0.5002]	[0.4012, 0.4932]
	-0.0259	-0.0254	0.3454	-0.0460	0.3178	1.0000
Narrow oracle diagnostic	[-0.0322, -0.0196]	[-0.0316, -0.0193]	[0.2050, 0.4858]	[-0.0547, -0.0372]	[0.2916, 0.3439]	[1.0000, 1.0000]

Table 35 Experiment III uncertainty summary from 30-seed benchmark runs.

Method	Mean utility (95% CI)	CEW (95% CI)	Sharpe (95% CI)	Max drawdown (95% CI)	Turnover (95% CI)	θ -stability (95% CI)
	0.1145	0.1213	1.4479	-0.0662	0.2327	0.0186
Full NDU (joint)	[0.1114, 0.1176]	[0.1178, 0.1248]	[1.4165, 1.4792]	[-0.0670, -0.0654]	[0.2201, 0.2453]	[0.0166, 0.0206]
	0.0923	0.0968	1.4941	-0.0579	0.3751	0.0000
Oracle-factor joint RL	[0.0896, 0.0951]	[0.0937, 0.0998]	[1.4596, 1.5287]	[-0.0589, -0.0569]	[0.3656, 0.3846]	[0.0000, 0.0000]
	0.0706	0.0732	1.3227	-0.0464	0.2954	0.0000
Fixed-preference joint RL	[0.0671, 0.0742]	[0.0694, 0.0771]	[1.2664, 1.3790]	[-0.0504, -0.0425]	[0.2782, 0.3125]	[0.0000, 0.0000]
	0.0642	0.0663	1.2051	-0.0625	0.3998	0.0000
Signal-only joint RL	[0.0605, 0.0679]	[0.0624, 0.0703]	[1.1635, 1.2466]	[-0.0639, -0.0611]	[0.3861, 0.4136]	[0.0000, 0.0000]

Appendix C: Technical Proofs and Verifications

C.1. Proof of Proposition 4.1 (Nesting of Frozen-Valuation Faces)

Proof. Take any original-state fixed-preference Markov policy $\pi \in \Pi_{RL}$ and define $\iota_0(\pi)(t, c, u) = (\pi(t, c), \bar{\theta})$. This is also an element of the augmented-state action-only/fixed-valuation class Π_{AO}^{aug} . More generally, any $\pi_{aug} \in \Pi_{AO}^{aug}$ embeds in Π_{NDU} by keeping the same augmented-state action rule and the frozen valuation coordinate $\bar{\theta}$. Because $\bar{\theta}$ is the unique zero-cost valuation-control configuration, the adjustment penalty of the embedded policy is zero. On the frozen valuation manifold, the drift, reward, and terminal payoff coincide with those of the fixed-preference HJB (23). Thus the physical state law and payoff functional under $\iota_0(\pi)$ are identical to those under π , so $J_{NDU}(\iota_0(\pi)) = J_{RL}(\pi)$.

Taking suprema gives

$$\sup_{\alpha \in \Pi_{NDU}} J_{NDU}(\alpha) \geq \sup_{\pi_{aug} \in \Pi_{AO}^{aug}} J_{NDU}(\pi_{aug}, \bar{\theta}) \geq \sup_{\pi \in \Pi_{RL}} J_{RL}(\pi),$$

which is the value dominance statement on the embedded manifold. Proposition 4.9 proves that as the valuation-adjustment cost scale tends to infinity, admissible maximizing selectors are driven to the zero-cost valuation configuration $\bar{\theta}$ and the NDU Hamiltonian converges locally uniformly to the fixed-preference Hamiltonian; comparison then yields convergence to the embedded fixed-preference value.

For original-state strictness, suppose two augmented states (c, u_1) and (c, u_2) share the same original state c but have different optimal Hamiltonian maximizers or terminal valuations with positive payoff margin. An original-state fixed-preference Markov rule $\pi(t, c)$ must choose the same action at both states, so it loses value at one of them. For fixed-valuation strictness against Π_{AO}^{aug} , compare the Bellman/Hamiltonian operator with θ frozen to $\bar{\theta}$ against the full joint-control operator. A positive full-minus-frozen Hamiltonian or terminal margin on a reachable set, net of adjustment cost and stable under the transition kernel, gives a strict dynamic-programming margin by the one-step Bellman comparison and backward induction (or by the viscosity comparison analogue in continuous time). Example 4.10 gives both a pure omitted-state separation and a same-information fixed- θ separation; the exact-DP inventory certificate in Table 12 gives the packaged finite-state analogue. $\square$

C.2. Proof of Theorem 4.3 (Operator-Level Non-Collapse)

Proof. A pointwise supremum of affine functionals is affine only if the same affine functional is active everywhere, or all active affine functionals have identical slope and intercept. The stated switching condition rules this out: there is one continuation value under which θ_0 is strictly preferred and another under which θ_1 is strictly preferred. Hence $\Phi_{x,a}$ has a kink or, more generally, non-affine support-function geometry. It cannot equal a fixed payoff-plus-transition affine form for all V .

Because the contradiction arises at a reachable state-action pair, the Bellman recursion propagates it backward by monotonicity. Thus an action-only augmented MDP with fixed reward and fixed transition primitives cannot reproduce the NDU operator except in the payoff-irrelevant or globally dominated valuation-control cases. $\square$

C.3. Proof of Corollary 4.4 (Non-Equivalence to Common Reductions)

Proof. For exogenous adaptive preferences, fixing the physical action fixes the transition law of the augmented state, hence the Bellman contribution is affine in V . For a static-multiplier CMDP, the Lagrangian reward $r - \eta c$ is fixed once η is chosen, and the action-conditioned Bellman contribution is again affine in V . By Theorem 4.3, an active NDU valuation channel gives a non-affine support operator in V ; therefore these classes cannot be operator-equivalent.

For entropic risk-sensitive control, the action-conditioned continuation term has the form

$$\rho^{-1} \log \mathbb{E}[\exp(\rho V(X'))],$$

whose second derivative in value directions is proportional to a conditional variance under the tilted measure. It is smooth and strictly curved whenever the next-state value is non-degenerate. A finite valuation-control support function is piecewise affine and generally nonsmooth at switching boundaries. These two operators cannot coincide on an open set unless both reduce to the same affine operator. That is again the payoff-irrelevant valuation-control case.

For robust control with a minimizing nature, the action-conditioned robust term is an infimum of affine functionals of V , hence concave in V . The NDU term is a supremum of affine functionals, hence convex. A function that is both convex and concave on a convex set is affine. Thus robust-control equivalence on an open set requires affine collapse, contradicting the switching condition. If one instead lets the decision maker choose the ambiguity model so that the robust envelope reproduces $\sup_{\theta} L_{\theta}$, then θ has simply been added back to the action/evaluator, which is outside an evaluator-preserving fixed-preference reduction. $\square$

C.4. Proof of Theorem 4.5 (Inventory Service-Pressure Structure)

Proof. For fixed θ , the inventory part of (29) is the standard convex periodic-review inventory objective: increasing S trades ordering/holding cost against expected backlog or lost-sales cost. Under the convexity and increasing-differences hypotheses, the Bellman value preserves convexity and submodularity by backward induction. Therefore the optimal order quantity has an order-up-to form $q = (S^* - I)^+$.

The comparative statics follow from increasing differences. Higher backlog B or a stochastically larger demand regime z increases the marginal value of raising S , so Topkis monotonicity gives S_t^* nondecreasing in B and z . Since $p'(\theta) > 0$, a larger service-pressure level increases the marginal penalty of unmet demand and therefore also increases the optimal order-up-to level.

For the valuation selector, the one-step objective has increasing differences between θ and the shortage/backlog stress variables because $p(\theta)$ multiplies expected unmet demand. Hence the maximizer θ_t^* is nondecreasing in B and z . The term $-\lambda(\theta - \theta^-)^2$ gives strict concavity and pulls the maximizer toward θ^- ; larger λ increases that pull. Strong concavity in θ gives uniqueness and Lipschitz continuity by the standard strong-monotonicity argument for argmax maps. $\square$

C.5. Proof of Theorem 4.6 (Valuation-Control Reduction)

Proof. For fixed (x, a, p) , the θ -dependent part of the Hamiltonian is a strongly concave quadratic over a compact convex set. The unconstrained first-order condition is

$$\eta(x, a, p) - \lambda R(\theta - \bar{\theta}(x)) = 0,$$

which gives $\bar{\theta}(x) + \lambda^{-1}R^{-1}\eta(x, a, p)$. Adding the convex constraint Θ yields the weighted projection (32). This proves uniqueness and the analytic elimination formula (33).

Weighted projections are nonexpansive in the R -norm. Combining nonexpansiveness with the Lipschitz property of η and $R \succeq mI$ gives (34). Finally, strong concavity implies the standard quadratic-growth inequality

$$\mathcal{H}(\theta^*) - \mathcal{H}(\tilde{\theta}) \leq \frac{1}{2\lambda m} \|\nabla_{\theta} \mathcal{H}(\tilde{\theta})\|^2,$$

with the projected-gradient variant under constraints. This proves the greedy-gap and computational-reduction claims, including the replacement of a joint $N_A N_{\Theta}^{d_{\theta}}$ grid by N_A physical-action evaluations plus the closed-form projection. $\square$

C.6. Proof of Theorem 4.7 (Stable Physical-Cost Improvement)

Proof. The derivative of the physical cost is

$$C'_z(S) = hF_z(S) - p_z(1 - F_z(S)) = (h + p_z)F_z(S) - p_z.$$

Thus $C'_z(S_z^*) = 0$, and S_z^* is the unique minimizer on the relevant interval. The second derivative is

$$C''_z(S) = (h + p_z)f_z(S) \geq (h + p_z)f_{\min}.$$

Strong convexity therefore gives

$$C_z(S_z^{\bar{\theta}}) - C_z(S_z^*) \geq \frac{1}{2}(h + p_z)f_{\min}(S_z^{\bar{\theta}} - S_z^*)^2.$$

Summing over z gives (39). The inequality is strict because no fixed $\bar{\theta}$ matches all p_z , so at least one positive-probability regime has a different critical fractile and, by the density lower bound, a different quantile. Subtracting the adjustment-cost term gives (40).

For stability, the density lower bound implies local quantile Lipschitz continuity: if $\sup_s |\hat{F}_z(s) - F_z(s)| \leq \varepsilon$, then each relevant critical-fractile quantile moves by at most $\varepsilon/f_{\min}$ on the interval. Since the cost functions have bounded local slopes there, each physical-cost difference changes by at most a constant times $\varepsilon/f_{\min}$. Summing over finitely many regimes gives the stated $K\varepsilon/f_{\min}$ perturbation bound. A population gap larger than this bound remains positive for sufficiently accurate demand estimation. $\square$

C.7. Proof of Theorem 2.5 (NBO HJB-Solver Certificate)

Proof. Write the HJB operator as

$$\mathcal{F}[V](t, x) = -\partial_t V(t, x) - \sup_a H(x, a, \nabla V(t, x), \nabla^2 V(t, x)).$$

When the verified residual is R_{β}^* , the interior optimal-Hamiltonian defect of V_{β} is bounded directly by δ_H and the terminal defect is bounded by δ_T . Let κ be the discount/comparison constant on

the compact domain and set $B(t) = \delta_T + (T - t)\delta_H$ for the optimal-value comparison. The shifted functions $V_\beta + B$ and $V_\beta - B$ are respectively super- and subsolutions up to the sign convention of $\mathcal{F}$, with terminal data bracketing G . By the comparison principle and the standard doubling/barrier argument for bounded viscosity solutions on compact parabolic domains,

$$V_\beta - B_C \leq V^* \leq V_\beta + B_C,$$

where $B_C \leq C_D(\delta_T + T\delta_H)$ after absorbing the Hamiltonian Lipschitz and discount factors into C_D . This proves the value bound.

The boundary issue in this argument is only the parabolic boundary of the compact certificate domain. The theorem assumes comparison for that bounded viscosity class, so no global growth condition outside D is used. If the compact certificate is implemented with state constraints or reflected/numerical boundary faces, the same statement applies to the corresponding comparison problem; otherwise the certificate must include the boundary residual on the nonterminal faces as part of δ_H . This is why the experiments distinguish exact compact, uniform, finite-cover, and cover-refinement certificates from sampled large-domain neural rows.

For the policy bound, fix the actor A_η and let V^{A_η} be the value of the controlled diffusion with that Markov feedback. The actor-induced Hamiltonian differs from the optimal Hamiltonian by at most δ_A , so the actor-controlled linear HJB has interior defect $\delta_H + \delta_A$ and terminal defect δ_T . Comparing V^{A_η} with V_β and then with V^* yields

$$0 \leq V^*(t, x) - V^{A_\eta}(t, x) \leq C_D(\delta_T + T\delta_H + T\delta_A).$$

If the available residual is instead $R_{\beta, \eta}$, then

$$|R_\beta^*| \leq |R_{\beta, \eta}| + \sup_a H(x, a, \nabla V_\beta, \nabla^2 V_\beta) - H(x, A_\eta, \nabla V_\beta, \nabla^2 V_\beta),$$

so the same comparison uses $\tilde{\delta}_H + \delta_A$ in place of δ_H . If a sequence of NBO critics and actors has the relevant sup-norm residual, terminal gap, and greedy gap tending to zero, the preceding estimates imply local uniform convergence of critics to the HJB value and vanishing policy suboptimality. This proves the NBO HJB-solver certificate. The compact theory-covered audit satisfies the certificate directly; the large sampled-domain HBO/NBO tables report certificate-gap closure metrics—residual, terminal, greedy gap, and rollout value—on the benchmark domains. $\square$

C.8. Proof of Proposition 5.1 (Decision-Layer Certificate)

Proof. For class inclusion, the embedding assumption means that every fixed-preference policy $\pi_0 \in \Pi_0$ has an embedded valuation-control policy $\iota(\pi_0) \in \Pi_F$ with the same physical action law and the same evaluator value:

$$J(\iota(\pi_0)) = J(\pi_0).$$

Therefore

$$\sup_{\pi \in \Pi_F} J(\pi) \geq \sup_{\pi_0 \in \Pi_0} J(\iota(\pi_0)) = \sup_{\pi_0 \in \Pi_0} J(\pi_0).$$

This proves part (i).

Part (ii) is a direct attained-policy comparison under the fixed benchmark tuple. If $\hat{\Delta} = J(\hat{\pi}_F) - J(\hat{\pi}_0) > 0$, then the reported valuation-control policy has higher held-out value than the named comparator under the same rollout law and evaluator. Since neither policy need be globally optimal, this is a lower-bound or audit statement rather than a class-optimality theorem.

For part (iii), by definition of the margin half-width the population margin is at least $\hat{\Delta} - h$ on the stated confidence event. The baseline optimizer-gap certificate gives

$$\sup_{\pi \in \Pi_0} J(\pi) \leq J(\hat{\pi}_0) + g_0.$$

Also $\sup_{\pi \in \Pi_F} J(\pi) \geq J(\hat{\pi}_F)$. Combining these inequalities gives

$$\sup_{\pi \in \Pi_F} J(\pi) - \sup_{\pi \in \Pi_0} J(\pi) \geq J(\hat{\pi}_F) - J(\hat{\pi}_0) - h - g_0 = \hat{\Delta} - h - g_0.$$

Thus a strictly positive right-hand side yields class-level separation for the finite benchmark/evaluator. In an exact finite-state dynamic program the optimizer gap and sampling half-width for the solved discrete problem are zero. In the CEM rows of this paper no global optimizer-gap certificate is asserted, so the rows are read as attained-policy audits.

For part (iv), $\mathcal{T}\hat{\pi}_F$ is evaluated under the same exogenous rollout law and differs from $\hat{\pi}_F$ only by the valuation-channel intervention specified in the benchmark. A strict inequality $J(\hat{\pi}_F) > J(\mathcal{T}\hat{\pi}_F)$ means that freezing, permuting, or removing that channel changes the induced trajectory/evaluator distribution in a value-relevant direction. Hence the learned valuation channel is active for that benchmark/evaluator.

Part (v) follows because E is simply another fixed evaluator on the same realized trajectories. Agreement between J and E is additional physical-performance evidence; disagreement means that the valuation objective and the external physical diagnostic encode different tradeoffs. The certificate reads the HBO/NBO, CEM, and exact-DP rows at the benchmark/evaluator level. $\square$

C.9. Proof of Proposition 2.30 (Consistency and Robustness)

Proof. We write the regularized joint control as $\alpha_t^\varepsilon = (a_t^\varepsilon, \theta_t^\varepsilon)$ and decompose the shadow price as $P_t^\varepsilon = (P_{c,t}^\varepsilon, P_{u,t}^\varepsilon)$. The key point is that full NDU is blockwise: the overt-action component is attached to the external-state block c , whereas the delicate singular limit lives in the preference block u .

1. *Joint-control consistency.* Let V^ε denote the value function of the ε -regularized problem and assume the comparison and regularity hypotheses used in Proposition 3.11. Under the additional smoothness assumptions required for classical verification, define

$$P_t^\varepsilon = \nabla_x V^\varepsilon(t, X_t^\varepsilon).$$

When $\partial_a h \equiv 0$, or when the a -dependence of the Hamiltonian is separable from P_u , the action-sensitive part of the Hamiltonian is recovered by standard current-value stochastic-control arguments through the nondegenerate c -block. The preference-sensitive part is different: when the u -noise is small, the diffusion projection no longer reveals $P_{u,t}^\varepsilon$ stably from the preference block of Z_t^ε , so the preference-control component must be parameterized directly through the shadow price itself. If a also affects the preference drift, then the same shadow-price parameterization is needed for the whole joint argmax because the action first-order condition contains $P_u^\top \partial_a h$.

The feedback law is therefore computed as

$$\alpha_t^\varepsilon \in \arg \max_{\alpha \in \mathcal{A} \times \Theta} \{P_t^{\varepsilon\top} b(X_t^\varepsilon, \alpha) + f(X_t^\varepsilon, \alpha) - \lambda \mathcal{L}(X_t^\varepsilon, \alpha)\}.$$

Under Assumption 3.1(ii), the selector is globally Lipschitz in (x, p) , so the regularized feedback is well defined. Fix a sequence $\varepsilon_k \downarrow 0$ along which $X^{\varepsilon_k} \rightarrow X^0$ in the state topology used for the SDE stability estimates and $P^{\varepsilon_k} \rightarrow P^0$ strongly in $L^2([0, T] \times \Omega; \mathbb{R}^{n+q})$. Define the limiting feedback by $\alpha_t^0 := \alpha^*(X_t^0, P_t^0)$. Then the selector Lipschitz bound yields $\|\alpha^{\varepsilon_k} - \alpha^0\|_{L^2([0, T] \times \Omega)} \leq L_\alpha (\|X^{\varepsilon_k} - X^0\|_{L^2([0, T] \times \Omega)} + \|P^{\varepsilon_k} - P^0\|_{L^2([0, T] \times \Omega)}) \rightarrow 0$. Thus the consistency claim is to be read as a continuity statement under strong convergence hypotheses, not as a compactness theorem for arbitrary weak sequential limits through the nonlinear argmax map. In particular, the action component a_t^ε is treated by the standard nondegenerate block only in the separable action-preference cases just described; in the general joint-control case both a_t^ε and θ_t^ε must be computed from the full shadow price.

2. *Vanishing auxiliary martingale in the preference block.* Only the fictitious preference noise vanishes as $\varepsilon \rightarrow 0$, so the relevant martingale term is

$$\mathcal{I}_t^\varepsilon := \int_0^t \sqrt{2\varepsilon} (P_{u,s}^\varepsilon)^\top dW_s^u.$$

By Doob's maximal inequality and Itô isometry,

$$\mathbb{E} \left[\sup_{t \in [0, T]} |\mathcal{I}_t^\varepsilon|^2 \right] \leq 4\mathbb{E} [|\mathcal{I}_T^\varepsilon|^2] = 8\varepsilon \mathbb{E} \int_0^T |P_{u,s}^\varepsilon|^2 ds.$$

If

$$\sup_{\varepsilon \in (0, 1]} \mathbb{E} \int_0^T |P_{u,s}^\varepsilon|^2 ds < \infty,$$

then the right-hand side is $O(\varepsilon)$ and therefore converges to zero. Hence the auxiliary preference-block martingale vanishes in $L^2(\Omega)$ as $\varepsilon \rightarrow 0$.

3. *Numerical robustness of the shadow-price parameterization.* In a standard deep-BSDE solver, one first learns Z_t^ε and then infers the relevant gradients by diffusion inversion. In the present blockwise setting,

$$Z_t^\varepsilon = \Sigma_\varepsilon (X_t^\varepsilon)^\top P_t^\varepsilon = \begin{bmatrix} \sigma_c(c_t)^\top P_{c,t}^\varepsilon \\ \sqrt{2\varepsilon} P_{u,t}^\varepsilon \end{bmatrix},$$

so the instability is concentrated in the preference block: recovering $P_{u,t}^\varepsilon$ from the second block requires division by $\sqrt{2\varepsilon}$. The action-relevant block is not affected in the same way only when the overt action does not enter h or when the Hamiltonian dependence on a is separable from P_u .

The shadow-price parameterization avoids this singular inversion by learning P_t directly and only then constructing Z_t^ε forward from the learned shadow price. As a consequence, the preference-control component, and any overt-action component coupled through $\partial_a h$, is handled through direct access to the full shadow price. This removes the exploding inversion of the preference block and explains why the regularized full-NDU solver remains numerically meaningful even when the u -noise is sent toward zero. $\square$

C.10. Proof of Lemma 3.6 (Restricted Fixed-Epsilon Verification)

Proof. Fix $\varepsilon > 0$ and a terminal window $I = [T - \delta, T]$. The Picard input space is the projected Hilbert subspace $\mathcal{H}_{\Sigma_\varepsilon}^2(I) \subset \mathcal{H}^2(I; \mathbb{R}^{d+q})$; equivalently one may work in reduced Brownian coordinates with square diffusion $(\Sigma_\varepsilon \Sigma_\varepsilon^\top)^{1/2}$. For an input volatility z , set

$$\alpha_\varepsilon^*(x, z) = \alpha^*(x, A_\varepsilon^\dagger z), \quad A_\varepsilon^\dagger = (\Sigma_\varepsilon \Sigma_\varepsilon^\top)^{-1} \Sigma_\varepsilon.$$

The induced volatility-Lipschitz constant is

$$\Lambda_\varepsilon \leq L_b^\alpha L_\alpha \|A_\varepsilon^\dagger\|_{\text{op}} \leq L_b^\alpha L_\alpha (2\varepsilon)^{-1/2},$$

where L_b^α is the drift Lipschitz constant in the control variable and L_α is the selector Lipschitz constant in p . Given z , solve

$$dX_t^z = b_\varepsilon(t, X_t^z, z_t) dt + \Sigma_\varepsilon d\mathbf{W}_t, \quad X_{T-\delta}^z = x,$$

and then solve the decoupled BSDE

$$dY_t^z = -g_\varepsilon(t, X_t^z, Y_t^z, Z_t^z) dt + (Z_t^z)^\top d\mathbf{W}_t, \quad Y_T^z = \Psi(X_T^z).$$

This defines $\Phi(z) = Z^z$.

For two inputs z^1, z^2 , standard SDE estimates give, for a constant C_X depending only on $L_{b,x}$,

$$\mathbb{E} \sup_{t \in I} |\Delta X_t|^2 \leq C_X \delta e^{C_X \delta} \Lambda_\varepsilon^2 \|\Delta z\|_{\mathcal{H}^2(I)}^2.$$

The BSDE stability estimate with terminal data $\Psi(X_T^1) - \Psi(X_T^2)$ and driver difference through $X^1 - X^2$ gives

$$\|\Delta Y\|_{\mathcal{S}^2(I)}^2 + \|\Delta Z\|_{\mathcal{H}^2(I)}^2 \leq C_B e^{C_B \delta} (L_\Psi^2 + \delta L_{g,x}^2) \mathbb{E} \sup_{t \in I} |\Delta X_t|^2,$$

where C_B depends on ρ and $L_{g,z}(\varepsilon)$. Combining the two displays,

$$\|\Phi(z^1) - \Phi(z^2)\|_{\mathcal{H}^2(I)}^2 \leq C_B C_X \delta e^{(C_B + C_X)\delta} (L_\Psi^2 + \delta L_{g,x}^2) \Lambda_\varepsilon^2 \|z^1 - z^2\|_{\mathcal{H}^2(I)}^2.$$

Choose $\delta_\varepsilon > 0$ so that the multiplier is strictly smaller than one. Banach's fixed-point theorem gives a unique fixed point Z , and the associated (X, Y) solve the coupled regularized system. The dependence of δ_ε on ε enters through Λ_ε and $L_{g,z}(\varepsilon)$, so no uniform small-time claim is made as $\varepsilon \downarrow 0$. $\square$

C.11. Proof of Proposition 3.8 (Global Well-Posedness under Small Coupling)

Proof. Use the same canonical projected map Φ as in the local proof, now on $\mathcal{H}_{\Sigma_\varepsilon}^2([0, T])$. For inputs z^1, z^2 , the forward stability estimate on the full horizon is

$$\mathbb{E} \sup_{t \in [0, T]} |\Delta X_t|^2 \leq C_X(T, L_{b,x}) \Lambda_\varepsilon^2 \|\Delta z\|_{\mathcal{H}^2(0, T)}^2.$$

The terminal condition enters the BSDE estimate through

$$\mathbb{E} |\Psi(X_T^1) - \Psi(X_T^2)|^2 \leq L_\Psi^2 \mathbb{E} |\Delta X_T|^2 \leq L_\Psi^2 \mathbb{E} \sup_{t \in [0, T]} |\Delta X_t|^2,$$

and the driver contribution through

$$\mathbb{E} \int_0^T |g_\varepsilon(t, X_t^1, Y_t^1, Z_t^1) - g_\varepsilon(t, X_t^2, Y_t^1, Z_t^1)|^2 dt \leq T L_{g,x}^2 \mathbb{E} \sup_{t \in [0, T]} |\Delta X_t|^2.$$

The standard Lipschitz BSDE estimate then yields

$$\|\Phi(z^1) - \Phi(z^2)\|_{\mathcal{H}^2(0, T)}^2 \leq C_B(T, \rho, L_{g,z}(\varepsilon)) C_X(T, L_{b,x}) (L_\Psi^2 + T L_{g,x}^2) \Lambda_\varepsilon^2 \|z^1 - z^2\|_{\mathcal{H}^2(0, T)}^2.$$

Thus (17) makes Φ a strict contraction, and Banach's theorem gives existence and uniqueness on $[0, T]$. $\square$

C.12. Proof of Propositions 3.11–3.12 (Viscosity Stability and Auxiliary Noise)

Proof. For each fixed $\varepsilon > 0$, the regularized HJB is uniformly parabolic on the controlled compact action set, with locally Lipschitz coefficients of at most linear growth and a proper discounted Hamiltonian. We work in the polynomial-growth comparison class assumed in Proposition 3.11; under that comparison/local-boundedness regime, standard viscosity theory yields a unique continuous viscosity solution V^ε independently of whether the FBSDE contraction argument from Section 3 remains globally valid at that ε . Let H_ε be the corresponding regularized Hamiltonian. As $\varepsilon \rightarrow 0$, $H_\varepsilon \rightarrow H_0$ locally uniformly. By the standard half-relaxed-limit stability theorem for viscosity solutions,

together with the assumed local uniform boundedness of $\{V^\varepsilon\}_{\varepsilon \in (0,1]}$, the upper and lower relaxed limits are respectively a subsolution and a supersolution of the limit equation. If the limit equation $\partial_t V + H_0 = 0$ satisfies a comparison principle in the relevant growth class, these envelopes coincide, and therefore $V^\varepsilon \rightarrow V^0$ locally uniformly.

For the auxiliary-noise term, let

$$\mathcal{I}_t^\varepsilon := \int_0^t \sqrt{2\varepsilon} (P_{u,s}^\varepsilon)^\top dW_s^u.$$

Then $\mathcal{I}^\varepsilon$ is a square-integrable martingale, so Doob's maximal inequality and Itô isometry give

$$\mathbb{E} \left[\sup_{t \in [0,T]} |\mathcal{I}_t^\varepsilon|^2 \right] \leq 4\mathbb{E} [|\mathcal{I}_T^\varepsilon|^2] = 8\varepsilon \mathbb{E} \left[\int_0^T |P_{u,s}^\varepsilon|^2 ds \right].$$

Under the uniform $\mathcal{H}^2$ bound and compatibility assumptions in Proposition 3.12, the right-hand side converges to zero as $\varepsilon \rightarrow 0$, so in particular $\mathcal{I}_T^\varepsilon \rightarrow 0$ in $L^2(\Omega)$. In neural parameterizations, proving such a compatible selection and uniform bound rigorously remains an additional analytical challenge in the vanishing-viscosity limit, which is why the proposition states them explicitly as assumptions rather than as derived estimates. $\square$

C.13. Proof of Lemma 3.16 (Flat-Preference Sufficient Regime)

Proof. Under the stated subclass assumptions, the regularized HJB depends on the preference coordinate u only through the drift term $\nabla_u V^\varepsilon \cdot h(c, u; a, \theta)$. Consider the function $\tilde{V}^\varepsilon(t, c, u) := \bar{V}^\varepsilon(t, c)$, where $\bar{V}^\varepsilon$ solves the reduced regularized HJB obtained by suppressing the irrelevant preference coordinate. Because $\tilde{V}^\varepsilon$ is independent of u , one has $\nabla_u \tilde{V}^\varepsilon = 0$, so the preference-drift contribution drops out. The remaining coefficients, reward terms, and terminal condition are exactly those of the full regularized HJB under assumptions (i)–(iii). Hence $\tilde{V}^\varepsilon$ is a continuous viscosity solution of the full regularized HJB with the same terminal condition as V^ε . By uniqueness of the continuous viscosity solution for the uniformly parabolic regularized HJB, $V^\varepsilon \equiv \tilde{V}^\varepsilon$ and therefore is independent of u . Consequently $P_{u,t}^\varepsilon = \partial_u V^\varepsilon(t, X_t^\varepsilon) = 0$ for all t , which immediately gives the stated uniform $\mathcal{H}^2$ bound. $\square$

C.14. Proof of Proposition 4.9 (Frozen-Valuation Consistency)

Proof. Fix a compact $K \subset \mathbb{R}^{n+q} \times \mathbb{R}^{n+q} \times \mathbb{S}^{n+q}$ and write

$$F(x, p, M, a, \theta) = \frac{1}{2} \text{Tr}(\Sigma \Sigma^\top M) + p^\top b(x, a, \theta) + f(x; a, \theta),$$

$$H^\lambda = \sup_{a, \theta} \{F - \lambda \mathcal{L}\}, \quad H_{RL} = \sup_a F(x, p, M, a, \bar{\theta}).$$

The common discount term is left in the PDE. Since $\mathcal{L}(x, a, \bar{\theta}) = 0$, $H^\lambda \geq H_{RL}$ on K . For the reverse bound, fix $\eta > 0$. Uniform continuity of F on $K \times \mathcal{A} \times \Theta$ gives $\delta > 0$ such that $|F(x, p, M, a, \theta) -$

$|F(x, p, M, a, \bar{\theta})| < \eta$ whenever $\|\theta - \bar{\theta}\| \leq \delta$. On $K_\delta = \{(a, \theta) : \|\theta - \bar{\theta}\| \geq \delta\}$, compactness and Assumption 4.8 give $m_\delta(K) = \inf_{\pi_x(K) \times K_\delta} \mathcal{L} > 0$, while $C_K = \sup_{K \times \mathcal{A} \times \Theta} |F| < \infty$. If $\lambda > 2C_K/m_\delta(K)$, every point in K_δ has payoff at most $C_K - \lambda m_\delta(K) < -C_K \leq H^\lambda$ and cannot maximize. Hence all maximizing selectors satisfy $\|\theta_\lambda^* - \bar{\theta}\| < \delta$ and therefore

$$0 \leq H^\lambda - H_{RL} \leq \sup_{a, \|\theta - \bar{\theta}\| \leq \delta} |F(x, p, M, a, \theta) - F(x, p, M, a, \bar{\theta})| < \eta$$

uniformly on K . Thus $H^\lambda \rightarrow H_{RL}$ locally uniformly on compact state/gradient/Hessian sets. The value bounds $V_{RL} \leq V^\lambda \leq V_{\text{free}}$ follow by embedding every RL policy with $\theta \equiv \bar{\theta}$ and by removing the nonnegative valuation-adjustment penalty. Local boundedness, half-relaxed-limit stability, and comparison for the limiting fixed-preference HJB then yield $V^\lambda \rightarrow V_{RL}$ locally uniformly on $[0, T] \times \mathbb{R}^{n+q}$. $\square$