

Limited-Memory Renewal Contracts: Participation, Quantization, and Exact Design

Anonymous manuscript for Operations Research

Abstract. A renewal provider must allocate accepted obligations through a limited execution alphabet. We solve the resulting joint design problem with endogenous branch promises, continuously chosen terminal levels, bounded realization exposure, and charges for the actual enabled levels. A branch-specific eligibility prefix gives the exact response to any promise. For quadratic operating costs and nonnegative quadratic level charges, eliminating lottery probabilities yields closed quadratic cells and a finite exact global algorithm. A second theorem rounds codewords and repairs their probabilities while preserving every branch promise and risk ceiling exactly. Its uniform error bound gives certified continuous-value intervals and an additive approximation scheme for a fixed alphabet budget. We retain the sharper saturated structure: cap-anchored nonhighest levels, a continuously optimized off-cap highest level, and contractual tail costs that remain Monge after minimization, permitting classical matrix search. Exact-alphabet thresholds, institutional-gap stability, homogeneous repeated-cap aggregation, and a charged interior pathwise-randomization example clarify the decision implications. Synthetic parameter studies compare continuous and catalog optima and separately measure the new algorithms. Established quantization, facility-opening, and quadratic-optimization principles are distinguished from the contractual reductions.

Keywords: renewal contracts; quantization; dynamic programming.

Subject classifications: Dynamic programming: service design; Programming: resource allocation.

Area of review: Optimization.

Submission category: Lengthy manuscript.

September 24, 2026

1 Introduction

An upstream renewal desk can know an accepted obligation that a downstream execution gateway does not retain. When only a small service symbol crosses that boundary, the provider must choose both its terminal instruction alphabet and the branch-dependent encoder. This is a quantizer-design problem, but its objective and feasible randomization are generated by a contract: an aggregate continuation promise, local participation limits, and the cost of intermediate rather than terminal service. We ask how to choose the alphabet and allocate the promise jointly when individual lottery realizations have an overrun limit and enabling an additional execution level has a cost.

The main result gives a single design formulation for every feasible aggregate promise, from zero to saturation, with continuously chosen execution levels, bounded realization exposure, and charges on the actual selected levels. A branch-specific eligibility prefix and a canonical adjacent lottery reduce each fixed-book problem to concave resource allocation. In the quadratic model, eliminating lottery probabilities then produces finitely many closed quadratic optimization cells. Exact stationary-face enumeration solves the continuous outer problem, including collisions and nonsaturated pathwise participation. This global reference algorithm is exponential, and we state that cost explicitly.

A second result supplies a computational alternative with a uniform guarantee. We round codewords downward, retain every branch promise, and reconstruct the lottery rather than retaining its probabilities. The resulting catalog policy preserves the promise and the risk bound exactly, even at a binding pathwise constraint. Its net-payoff loss has an explicit bound proportional to the catalog spacing, uniformly over promises and risk tolerances. Exact catalog allocation therefore gives certified lower and upper values for the continuous problem, and an additive approximation scheme when the symbol budget is fixed. This connection is what makes a numerical catalog a controlled approximation rather than a change of model.

The saturated contract remains a particularly tractable and informative member of this theory. Deterministic execution yields an ordered one-sided quantizer. Under expected participation, a quadratic optimum has cap-anchored nonhighest levels and one continuously optimized highest level, which can lie strictly between caps because intermediate service meets the shortfall. A crossing-difference identity preserves Monge order after continuous tail minimization. Classical matrix search then gives linear work per dynamic-programming layer. We retain these results, their complete proofs, the all-promise exact-alphabet phase diagram and stability bounds, and the faster saturated charged-catalog recurrence. The general algorithm complements rather than replaces these sharper special-case results.

Ordered quantizer dynamic programming, optimized reproduction points, unbiased random splitting, and matrix search are established ideas (Wu, 1991; Pagès and Wilbertz, 2012; Croci et al., 2022; Jourdain and Pagès, 2021). Likewise, a cost per enabled representative and quadratic active-set analysis are not generic methodological inventions of this paper. The contractual contribution is the interaction of endogenous promises, heterogeneous shortfall service, branch-specific realization ceilings, and a shared selected alphabet. That interaction yields the eligibility-prefix reduction, the exact closed-cell formulation, and a one-sided repair that needs no promise or risk slack. The heterogeneous tail and its optimized Monge cancellation remain a separate structural advantage at saturation.

The operational interpretation is a risk-neutral or actuarial acceptance rule followed by a restricted ex-

ecution interface; hard realization limits represent a different acceptance institution. We spell out those assumptions and distinguish documented interface-security analogies from evidence of an actual deployed contract. Exact computations map institutional gaps and joint promise–risk–price choices, compare continuous and catalog optima, and time the new algorithms separately from inherited benchmarks. These are synthetic design studies and implementation checks, not customer data or substitutes for mathematical proofs.

2 Quantization, Contracting, and the Contribution Boundary

Ordered scalar quantization. Wu (1991) accelerates the discrete optimal mean-square K -level quantization dynamic program from $O(KN^2)$ to $O(KN)$ by matrix searching. Later one-dimensional clustering treatments extend the cost structures and discuss reconstruction and space reductions (Grønlund et al., 2017). Thus contiguous-cell dynamic programming, continuous reproduction-point optimization, and an $O(mk)$ matrix-search consequence are not new as generic claims. In the present deterministic problem, participation forces a cell’s reconstruction to its smallest cap, not its least-squares centroid. Its loss includes branch-specific intermediate-service costs. The difference is a particular feasible distortion, not a new ordered-partition paradigm.

Unbiased randomization and dual quantization. The adjacent lottery is exactly the classical unbiased two-point quantizer. For distinct adjacent levels $u < v$ and $u \leq b \leq v$,

$$\Pr(C = v) = \frac{b - u}{v - u}, \quad \Pr(C = u) = \frac{v - b}{v - u}, \quad \mathbb{E}C = b. \quad (1)$$

Stochastic rounding uses this same rule (Crocì et al., 2022). Dual quantization is an even closer optimization predecessor: Pagès and Wilbertz (2012) optimize mean-preserving random splitting and relate it to piecewise-affine interpolation. Their extended formulation also uses nearest-neighbor projection outside the grid’s convex hull. We therefore do not claim that random splitting, optimized grids, interpolation, or treatment of out-of-hull points first arise here.

For the quadratic terminal reward, our in-hull loss is $q(b - u)(v - b)/2$, the corresponding quadratic dual-quantization error multiplied by $q/2$. Above the highest level z , however, the contractual loss is

$$f(b) - f(z) + h_j(b - z) = (r - qz)(b - z) + \frac{\gamma_j - q}{2}(b - z)^2.$$

It is generally not a common squared-distance projection penalty: it carries a level-dependent linear term and a heterogeneous intermediate-service curvature. The paper’s combined saturated structural result concerns this one-sided contractual tail, its continuously chosen endpoint, and the cancellation that preserves the terminal Monge order. The affine movement of an interior codeword follows the quadratic interpolation structure and is not advertised as an independent general quantization principle. Its sufficient conditions are explicit: saturated branch targets, quadratic terminal reward, uncharged codewords, and intermediate costs arising only above the highest level. Section 8 treats endogenous targets, and Section 11 treats actual level

charges without assuming the same anchoring lemma.

Optimized dual grids and enabled-level penalties. Jourdain and Pagès (2021) prove uniqueness of optimal one-dimensional dual grids for compactly supported log-concave densities and develop a Lloyd-like algorithm in the quadratic case. That result directly concerns optimized grids, not merely evaluation of a prescribed rounding rule. Here branch targets are themselves constrained decisions, intermediate shortfall costs can differ by branch, and each branch admits a different eligible prefix. Our exact continuous reduction does not assume a log-concave source density or uniqueness of a stationary grid; the grid approximation also preserves the contractual constraints at their boundary.

A charge for enabling a representative is a fixed-opening-cost idea, not a new type of resource pricing. For example, Aardal et al. (2014) distinguish facility-opening costs from demand-dependent service costs under a cardinality restriction. In our model, a selected level can implement adjacent random splitting and branch-specific intermediate shortfalls, rather than serving a client by ordinary distance to one center. The opening charge is not entropy, expected message length, or a cost per token use. The new issue is its joint optimization with exact promises and realization eligibility. The finite quadratic search in Theorem 9.3 invokes a classical optimization principle; its contribution is the contractual reduction to those closed cells. Theorem 10.1 gives the separate, constraint-preserving approximation guarantee.

Finite-rate control and contracts. An informed encoder and a downstream finite-alphabet decoder are standard in communication-constrained control (Fu, 2024). Our result is not a stabilization or channel-capacity theorem. There is no dynamic plant whose state is estimated through noisy feedback, no hidden type report, and no incentive-compatibility claim. What the contract adds is the allocation of accepted continuation payments and the distinction between conditional expected and realization-wise participation. Continuation promises themselves have a long contracting history (Spear and Srivastava, 1987; Thomas and Worrall, 1988), including finite policy-graph implementations (Zhang, 2012). Below the exact alphabet threshold, participation determines the allowable quantizer rather than communication alone determining the objective.

Classical matrix search, model-specific cost structure. Quadrangle-inequality dynamic programming (Yao, 1980) and full totally monotone matrix search (Aggarwal et al., 1987) supply the algorithmic framework. Generalized and partial Monge matrix searching require their own domain analysis (Klawe and Kleitman, 1990; Chan, 2021; Kaplan et al., 2017). Our suffix and prefix domains have the favorable completion orientation; this is not an improvement on the general partial-matrix bounds. We retain complete proofs of the finite tie rule, wholly padded selected rows, column offsets, and reachable predecessor intervals so that the implementation’s use of the classical result is checkable. The essential contractual input is the inequality after continuous tail minimization, not merely convexity of a scalar reproduction-point objective. We do not claim that minimizing a reproduction point preserves Monge structure for arbitrary quantizer losses.

Table 1 states the predecessor principle and contractual increment for each result.

3 The Contract and Its Execution Boundary

3.1 A saturated renewal agreement

At the renewal date, a public branch $j \in \{1, \dots, k\}$ is observed with probability $\pi_j > 0$, where $\sum_j \pi_j = 1$. Its continuation cap is $0 < b_1 < \dots < b_k < 1$. The provider chooses an intermediate tier $Y_j \in [0, 1]$ and a terminal tier $C_j \in [0, 1]$. All branches subsequently reach the same terminal public vertex. Unit payment coefficients and unit discounting make total remaining payment $Y_j + C_j$. A fixed service-benefit continuation minus the customer's outside continuation gives the cap. At the root the provider has committed to

$$\sum_j \pi_j \mathbb{E}[Y_j + C_j] = \bar{b}, \quad \bar{b} := \sum_j \pi_j b_j. \quad (2)$$

Under expected participation, $\mathbb{E}[Y_j + C_j] \leq b_j$ for every j , with expectation conditional on that observed branch and taken before the internal draw. Positivity of the probabilities and (2) imply equality separately on every branch. Under pathwise participation, the stronger constraint is $Y_j + C_j \leq b_j$ almost surely. The same argument then implies equality almost surely, not merely in expectation.

The operating payoff is $f(C_j) - h_j(Y_j)$, with increasing concave f and convex nondecreasing h_j satisfying $h_j(0) = 0$. General shape assumptions suffice for the deterministic and pathwise results. The global continuous algorithm specializes to the quadratic primitives in (8). The negative intermediate contribution represents the net cost of an advance or expedited tier, while the terminal tier contributes net operating return. It does not mean that the customer has negative service benefit: that benefit has already entered the continuation cap. There are no endogenous public transitions, coupled commitments, or switching costs in the renewal theorem.

3.2 What a writable code means

At the observed branch, an encoder selects a symbol $Z \in \{1, \dots, m\}$. At the terminal vertex the decoder observes only the current public vertex and Z . Its stored program can contain arbitrary read-only model constants but cannot recover the realized branch from those constants. An inherited intermediate tier, a customer identifier with a branch lookup, a retained shared random seed, and a transmitted numerical payment are all additional history channels; none is free in this interface. An m -symbol alphabet requires $\lceil \log_2 m \rceil$ fixed-width writable bits, not m bits. Public-state storage and read-only program storage are counted separately.

Private encoder randomization is allowed under expected participation. Fresh decoder randomness is independent of earlier private choices conditional on the symbol. Any seed that correlates those choices across dates must itself be encoded in the charged symbol. Concavity permits a deterministic terminal decoder at the expected-participation optimum, so a codebook is a collection of terminal levels. The lower bound for exact implementation continues to allow randomized encoders and decoders under this accounting. Distinct full-information terminal actions and strict concavity remain visible assumptions; Proposition 8.2 handles repeated caps; uniqueness is required for the exact-action lower bound.

3.3 Acceptance timing

The renewal desk observes the branch and offers the intermediate tier and terminal lottery. Under expected participation, the customer accepts before the ticket draw; after acceptance only the ticket reaches the common gateway. The realized ticket may be disclosed, but no new costless exit option is supplied after the draw within this episode. Under pathwise participation, each possible realized obligation must instead satisfy the branch cap, and a draw-specific intermediate tier is allowed. A seed known before acceptance changes that conditioning information. Repeated-customer effects require an enlarged state and are not part of the three-date theorem. The full stipulated mechanism and its qualifications are given in the electronic companion; no empirical or legal claim about an actual provider is made. Section 8 allows a nonsaturated root promise, and Section 11 specifies the intermediate risk protocol.

3.4 Preferences and a restricted execution interface

A concrete preference interpretation of expected participation is quasilinear, risk-neutral continuation utility $A_j - \mathbb{E}(Y + C)$ with outside value O_j , so that $b_j = A_j - O_j$ after normalization. Alternatively the expectation bound can be an explicitly agreed actuarial budget. Neither interpretation follows from a finite communication alphabet. A customer with a hard realization-level liquidity limit instead requires $Y + C \leq b_j$. The intermediate institution retains the actuarial cap and separately authorizes at most δ realization overrun. Its intermediate tier is committed before the draw so that the disclosed lottery and its ceiling can be checked together. Allowing a fresh, costless rejection after a high draw changes the feasible problem; the expected-participation comparison assumes that option is absent, not that every real renewal contract can eliminate it.

The downstream restriction is an interface restriction, not an inability to store a model. An upstream desk keeps the branch record; a downstream gateway has a read-only map from a small session token to tested execution levels and no access to the upstream record or a branch-indexed customer identifier. A relevant documented component analogy is operational-technology allowlisting and least functionality: NIST SP 800-82 Rev. 3 discusses restricting executable applications, testing policy configurations, disabling nonessential interfaces, and restricting writable set points (Stouffer et al., 2023, Sections 5.2.5.1, 6.2.4.1 and control CM-7). These are precedents for a constrained execution boundary, not evidence that a particular provider uses our entire renewal or lottery protocol.

Under this design, opening an additional level can incur a one-time cost of engineering, test fixtures, authorization, and configuration assurance. This is the modeled source of $\rho(c)$: it is paid if the level is enabled, independently of the frequency with which its token is used. Per-use charges and entropy-based communication charges would be different objective terms. An exogenous catalog represents levels already eligible for certification; a numerical catalog chosen by the operator instead calls for the approximation theorem in Section 10. The model does not assert that every read-only contract must be inaccessible downstream; it quantifies the consequence when that access is deliberately excluded. No empirical adoption, calibration, or enforceability conclusion is inferred from the analogy.

4 The Deterministic Service-Code Frontier

Consider three dates. The root tier is fixed at zero. At the intermediate date a branch j is observed with probability $\pi_j > 0$, where $\sum_j \pi_j = 1$. All branches then enter the same terminal public vertex. Payments and discounting equal one. Intermediate and terminal tiers are in $[0, 1]$. Branch j has cap $0 < b_1 < \dots < b_k < 1$, and the root promise is $\bar{b} = \sum_j \pi_j b_j$. At the intermediate date the reward is $-h_j(y)$, where h_j is convex, nondecreasing, and $h_j(0) = 0$. The common terminal reward $f(c)$ is continuous, concave, and strictly increasing on $[0, 1]$. The terminal cap is one. The terminal decoder observes only its public vertex and its memory symbol, not the inherited intermediate tier.

Theorem 4.1 (Heterogeneous renewal frontier). *The unique full-information policy is $y_j = 0$, $c_j = b_j$, with value $J = \sum_j \pi_j f(b_j)$. Exact implementation requires k symbols under strict cap ordering; repeated caps are covered by Proposition 8.2. With at most m deterministic terminal symbols, an optimal encoding has contiguous nonempty cells in cap order. For a cell $i, \dots, j$, its terminal tier is b_i and its loss is*

$$C(i, j) = \sum_{h=i}^j \pi_h [f(b_h) - f(b_i) + h_h(b_h - b_i)]. \quad (3)$$

The exact minimum loss is $D_m = D(\min(m, k), k)$, where

$$D(s, j) = \min_{s-1 \leq i < j} \{D(s-1, i) + C(i+1, j)\}, \quad D(0, 0) = 0, \quad (4)$$

and all other infeasible entries equal $+\infty$. With a precomputed cell-cost table this requires $O(mk^2)$ arithmetic operations. For quadratic f and quadratic h_j , prefix sums of the probability-weighted coefficients compute each cell cost in constant arithmetic time.

The proof is in Electronic Companion Section EC.1.1.

This strengthens and includes the quadratic equal-probability frontier. With $f(c) = 2c - c^2/2$, $h_j(y) = y^2/2$, and $\pi_j = 1/k$, (3) becomes

$$C(i, j) = \frac{2 - b_i}{k} \sum_{h=i}^j (b_h - b_i). \quad (5)$$

For $k = 8$ and $b_j = j/9$, one, two, and three symbols lose $119/162$, $5/18$, and $49/324$, respectively. Two writable bits permit four symbols, not three. The segmentation recurrence is standard; the contract-specific cell loss, ordered-cell sufficiency, and exact memory interpretation are the substantive statements.

5 When Randomization Improves a Restricted Memory Budget

The exact-optimum lower bound in Theorem 8.1 does not imply that randomization is useless below that alphabet size. The renewal family permits an explicit answer under the declared conditional-expectation participation institution. At an intermediate public branch, the controller may randomize its terminal symbol;

it carries no uncharged seed or branch record. Participation is evaluated before that private draw. Terminal actions may depend only on the terminal public vertex and the retained symbol.

For terminal levels $c_1 < \dots < c_s$, let F_c be the linear interpolant through $(c_i, f(c_i))$. Since f is concave and strictly increasing, this interpolant is concave and strictly increasing on $[c_1, c_s]$; for $s = 1$ it is defined at its sole point. Retain the caps, probabilities, and convex intermediate costs of Theorem 4.1.

Theorem 5.1 (Optimal lotteries for a fixed terminal alphabet). *For a fixed set of terminal levels, a feasible randomized controller exists if and only if $c_1 \leq b_1$. Its optimum at branch j has deterministic intermediate tier and terminal mean*

$$\mu_j = \min\{b_j, c_s\}, \quad y_j = b_j - \mu_j.$$

It randomizes over at most the two adjacent terminal levels bracketing μ_j . The optimal loss with at most m writable terminal symbols is consequently

$$D_m^{\text{ran}} = \min_{\substack{1 \leq s \leq m, \\ 0 \leq c_1 < \dots < c_s \leq 1 \\ c_1 \leq b_1}} \sum_j \pi_j \{f(b_j) - F_c(\min\{b_j, c_s\}) + h_j((b_j - c_s)_+)\}. \quad (6)$$

The minimum allows coincident levels by deleting unused duplicates. It satisfies $D_m^{\text{ran}} \leq D_m$. Equality with the full-information optimum still requires k symbols when the full-information actions are distinct and rewards are strictly concave. Formula (6) optimizes the encoding for each continuous codebook; it is not a claim that the joint codebook problem is solved by the deterministic segmentation recurrence.

The proof is in Electronic Companion Section EC.1.2.

Proposition 5.2 (A strict, exactly solved two-symbol separation). *In the quadratic renewal model, take three equiprobable branches with caps $1/4, 1/2, 3/4$, $f(c) = 2c - c^2/2$, and $h_j(y) = y^2/2$. Then*

$$D_2 = \frac{1}{8}, \quad D_2^{\text{ran}} = \frac{1}{96}, \quad D_3 = D_3^{\text{ran}} = 0. \quad (7)$$

Two symbols thus reduce the randomized loss to one twelfth of the best deterministic loss, without reducing the three-symbol requirement for exact implementation.

The proof is in Electronic Companion Section EC.1.3.

The institutional qualification is substantive. On the middle branch, the high private draw delivers terminal payment $3/4$ against an expected-continuation cap $1/2$. The lottery is feasible when participation is checked before the private draw, as stipulated here, but not when that branch cap must hold for every private draw. We do not transfer the randomized frontier to the pathwise institution. The deterministic frontier and the exact-memory lower bound retain their original domains. This separation explains why a statement about the minimum alphabet at the exact optimum cannot settle the economic value of randomization under a smaller alphabet.

6 The Global Continuous Lottery Codebook

The fixed-codebook formula in Theorem 5.1 leaves a continuous design problem: both the terminal service levels and their assignment probabilities must be chosen. We now solve that problem globally for quadratic terminal rewards and heterogeneous quadratic intermediate costs. The solution also determines when imposing participation after, rather than before, the private draw eliminates the value of randomization. These are restricted-alphabet design results, distinct from the exact-implementation alphabet in Theorem 8.1.

Retain the three-date renewal architecture of Theorem 4.1, including the saturated root promise $\bar{b} = \sum_{j=1}^k \pi_j b_j$, strictly ordered caps $0 < b_1 < \dots < b_k < 1$, and positive probabilities summing to one. Throughout this section,

$$f(c) = rc - \frac{q}{2}c^2, \quad h_j(y) = \frac{\gamma_j}{2}y^2, \quad r \geq q > 0, \quad \gamma_j \geq 0. \quad (8)$$

The parameters and probabilities can differ across branches as indicated. Terminal reward is strictly increasing on the feasible cap range, including when $r = q$. All private information retained until the common terminal vertex must be encoded in the charged symbol. The decoder and public observation convention are unchanged.

6.1 Cap anchors and one continuous terminal level

For $u \leq b \leq v$, quadratic concavity gives the interpolation loss

$$f(b) - \frac{v-b}{v-u}f(u) - \frac{b-u}{v-u}f(v) = \frac{q}{2}(b-u)(v-b), \quad (9)$$

with the endpoint value understood when levels coincide. Thus an interior codeword enters the loss affinely until it crosses a branch cap. The highest codeword is different: reducing it below a branch cap creates an intermediate payment and its associated cost.

Lemma 6.1 (Cap-anchored representation). *For every alphabet budget $m \geq 1$, an optimum of (6) exists with the following representation. Its lowest codeword is b_1 ; every codeword except the highest belongs to $\{b_1, \dots, b_k\}$; and the highest belongs to $[b_1, b_k]$. A one-symbol optimum is $\{b_1\}$. The highest codeword need not be a branch cap.*

The proof is in Electronic Companion Section EC.1.4.

The lemma does not discretize the highest level. A cap-only restriction can be strictly suboptimal, as Proposition 6.3 shows. It also does not assert that every optimal codebook is anchored; it supplies an optimum with this useful form.

6.2 A finite exact algorithm

For adjacent anchored levels $b_i < b_j$, define

$$E(i, j) = \frac{q}{2} \sum_{h=i}^j \pi_h (b_h - b_i)(b_j - b_h). \quad (10)$$

Let $P(s, j)$ be the minimum interpolation loss on caps through b_j using exactly s anchored levels, starting at b_1 and ending at b_j . Set $P(1, 1) = 0$, $P(1, j) = +\infty$ for $j > 1$, and use the ordered recurrence

$$P(s, j) = \min_{i < j} \{P(s-1, i) + E(i, j)\}, \quad 2 \leq s \leq j. \quad (11)$$

This is a standard shortest-path or segmentation recurrence. The substantive reduction is Lemma 6.1, which makes it applicable to a jointly optimized continuous lottery codebook.

Suppose that b_i is the last anchored level and z is the final level. For $\ell \geq i$ and $b_\ell \leq z \leq b_{\ell+1}$, put $b_{k+1} = b_k$ and define

$$\begin{aligned} \Psi_{i\ell}(z) = & \frac{q}{2} \sum_{h=i+1}^{\ell} \pi_h (b_h - b_i)(z - b_h) \\ & + \sum_{h=\ell+1}^k \pi_h \left\{ f(b_h) - f(z) + \frac{\gamma_h}{2} (b_h - z)^2 \right\}. \end{aligned} \quad (12)$$

The first sum is the lottery loss below z ; the second is the terminal-reward and intermediate-payment loss above it. Empty sums are zero. For $\ell < k$ the function is strictly convex in z and has the constrained minimizer

$$z_{i\ell} = \left[\frac{r \sum_{h>\ell} \pi_h + \sum_{h>\ell} \pi_h \gamma_h b_h - \frac{q}{2} \sum_{h=i+1}^{\ell} \pi_h (b_h - b_i)}{q \sum_{h>\ell} \pi_h + \sum_{h>\ell} \pi_h \gamma_h} \right]_{[b_\ell, b_{\ell+1}]}. \quad (13)$$

For $\ell = k$, take $z_{ik} = b_k$. Define the terminal cost

$$T_i = \min_{i \leq \ell \leq k} \Psi_{i\ell}(z_{i\ell}). \quad (14)$$

Theorem 6.2 (Global frontier and exact complexity). *Under (8), the optimum with at most m symbols is*

$$D_m^{\text{ran}} = \min \left\{ D_1, \min_{\substack{1 \leq s \leq \min(m-1, k) \\ 1 \leq j \leq k}} \{P(s, j) + T_j\} \right\}, \quad (15)$$

where the inner minimum is empty for $m = 1$. Backtracking produces a globally optimal codebook and the adjacent lotteries implementing it. For ordered inputs and $1 \leq m \leq k$, all frontiers through budget m can be computed in $O(mk^2)$ exact rational operations and $O(mk + k)$ stored rational numbers and indices, including backpointers. No codeword grid is required.

For rational input of total binary length L , a reduced-rational implementation has polynomial Turing complexity. A conservative bound is $O(mk^2 S^3)$ bit operations with $S = O(k[L + \log(k+2)])$. This includes exact cost comparisons, clamping, and fraction reduction. For $m \geq k$ the optimal loss is zero; listing additional budget entries requires only their output cost. These guarantees concern the stated algorithm, not an arbitrary nonlinear optimizer for (6).

The proof is in Electronic Companion Section EC.1.5.

6.3 A continuous optimal highest level

Proposition 6.3 (An off-cap terminal level is optimal). *Let $(b_1, b_2, b_3) = (1/4, 1/2, 3/4)$, $(\pi_1, \pi_2, \pi_3) = (7/20, 3/5, 1/20)$, $r = 2$, $q = 1$, and $\gamma_j = 1$. The globally optimal two-symbol expected-participation codebook is $(1/4, 5/8)$, and*

$$D_2^{\text{ran}} = \frac{23}{1280} < \frac{3}{160} = \min_{\substack{|c| \leq 2 \\ c \subseteq \{b_1, b_2, b_3\}}} D(c) = D_2, \quad (16)$$

where the cap-only minimum is over feasible codebooks and $D(c)$ denotes their expected-participation loss. Thus even optimizing all cap-only lotteries misses the unrestricted optimum.

The proof is in Electronic Companion Section EC.1.6.

6.4 Participation after the private draw

The preceding frontier is appropriate when a renewal review approves an expected continuation before the protocol draws its private symbol. Consider instead an agreement that requires the same cap to hold for every private realization after that review. A draw-specific intermediate tier is allowed; only the terminal public state and the charged symbol reach the terminal decoder. Write $D_m^{\text{pw,ran}}$ for the best randomized loss under these pathwise constraints and the same saturated root promise.

Theorem 6.4 (Pathwise randomization has no frontier advantage). *For the general renewal family of Theorem 4.1, with strictly increasing concave f and convex nondecreasing h_j , and the observation architecture just specified,*

$$D_m^{\text{pw,ran}} = D_m \quad \text{for every } m \geq 1. \quad (17)$$

Consequently (15) and the deterministic frontier give an exact comparison of the two participation institutions under (8). Their exact-optimum alphabet thresholds are both k when strict concavity makes the k distinct full-information actions unique.

The proof is in Electronic Companion Section EC.1.7.

This institutional result is not an assertion that randomization is generally ineffective in dynamic contracts. It uses the saturated root promise, an exogenous recombination, and the charged-information interface. Under expected participation, Theorem 6.2 chooses a less expensive lottery protocol; under pathwise participation, the same service model calls for the deterministic frontier. In either institution an operator can price a symbol budget, add its implementation charge to the computed loss, and optimize the budget without confusing read-only program size with retained private information.

7 Monge Costs and Classical Linear Matrix Search

The finite dynamic program still optimizes a continuous highest service level. We first prove the contractual cost inequalities and then verify the classical easy-direction completion with the exact order used by the code. A search routine cannot establish the economic reduction that supplies its entries.

7.1 The inequality survives continuous terminal minimization

Write $I_\ell = [b_\ell, b_{\ell+1}]$, including $I_k = \{b_k\}$, and $Q(i, \ell) = \min_{z \in I_\ell} \Psi_{i\ell}(z)$ for $i \leq \ell$. A partial array is Monge here when

$$A(r, c) + A(s, d) \leq A(r, d) + A(s, c), \quad r < s, c < d, \quad (18)$$

whenever all four entries are defined. This fixes the inequality convention before any orientation is named.

Lemma 7.1 (Contractual finite-cost inequalities). *The anchored edge costs E and deterministic cell costs C satisfy $A(a, c) + A(b, d) \leq A(a, d) + A(b, c)$ for $a \leq b \leq c \leq d$. The continuously minimized terminal array Q is Monge on its feasible domain. These statements permit arbitrary positive probabilities and heterogeneous nonnegative intermediate curvatures.*

The proof is in Electronic Companion Section EC.1.8.

Lemma 7.2 (Previous-layer values are column offsets). *Adding any finite column term v_c preserves the Monge inequality. For fixed finite predecessor values, the arrays $A(j, i) = P(s, i) + E(i, j)$ and $B(j, i) = D(s - 1, i) + C(i + 1, j)$ are therefore Monge wherever all four entries are feasible.*

The proof is in Electronic Companion Section EC.1.9.

7.2 Completion with the implemented strict and tie orders

Give a feasible entry the lexicographic key

$$K(r, c) = (0, A(r, c), c). \quad (19)$$

Finite ties select the smaller column, and distinct columns have distinct keys. The exact total-monotonicity condition used by the implementation is

$$K(r, d) < K(r, c) \implies K(s, d) < K(s, c), \quad r < s, c < d. \quad (20)$$

A strict preference for the right column cannot reverse in a later row. This applies to arbitrary selected rows and columns. Because keys are distinct, a decrease of row-argmin indices in any submatrix would contradict (20).

Lemma 7.3 (Suffix-feasible completion). *Let row r have feasible columns $c \geq \ell_r$, with nondecreasing cutoffs ℓ_r . Suppose the finite entries are Monge. Complete every $c < \ell_r$ with $K(r, c) = (1, 0, -c)$. Then (20) holds. A completely infeasible row of a selected submatrix selects its largest selected column.*

The proof is in Electronic Companion Section EC.1.10.

Lemma 7.4 (Prefix-feasible completion). *Let row r have feasible columns $c \leq u_r$, with nondecreasing cutoffs u_r . Suppose the finite entries are Monge. Complete every $c > u_r$ with $K(r, c) = (1, 0, c)$. Then (20) holds. A completely infeasible selected row selects its smallest selected column.*

The proof is in Electronic Companion Section EC.1.11.

These lemmas specialize the known favorable completion setting in Section 2; they are not general partial-matrix search theorems. A helper accepting two bounds does not prove total monotonicity for arbitrary interval rows. For example, two zero-cost rows feasible only at columns 2 and 1 have a vacuously Monge finite part but decreasing minima. Uniform infinite padding with ordinary left ties also fails on the required suffix domain: rows with cutoffs 2, 3 restricted to columns 1, 2 select 2, 1. The opposite internal padding orders above are deliberate.

7.3 Reachability and three search interfaces

Lemma 7.5 (Contiguous reachable predecessor sets). *For the anchored program, $P(1, j)$ is finite exactly at $j = 1$, and for $s \geq 2$, $P(s, j)$ is finite exactly when $s \leq j \leq k$. For the deterministic program, $D(0, j)$ is finite only at $j = 0$, and $D(s, j)$ is finite exactly when $s \leq j \leq k$ for $s \geq 1$. Each globally reachable predecessor set that can extend to a further endpoint is a contiguous interval or the initial singleton.*

The proof is in Electronic Companion Section EC.1.12.

The following pseudocode uses one-based cap indices and zero for the empty deterministic prefix. “Search” means ordinary full-matrix SMAWK on the supplied key. All prior-layer values are known before a search.

Terminal interface. Rows and columns are $1, \dots, k$. At (i, ℓ) return $(1, 0, -\ell)$ if $\ell < i$; otherwise minimize (12) by (13) and return $(0, Q(i, \ell), \ell)$. Search all rows. Reevaluate each winning cell once to store T_i and its continuous optimizer.

Anchored-prefix interface. To form layer $s+1$, use rows $j = s+1, \dots, k$. Columns are $\{1\}$ for $s = 1$, and $s, \dots, k-1$ otherwise. At (j, i) return $(1, 0, i)$ if $i \geq j$, and $(0, P(s, i) + E(i, j), i)$ otherwise. Search and store values and predecessors. Scan the current reachable endpoints with T_i to update the at-most-budget incumbent; backtrack only its winning codebook.

Deterministic-prefix interface. To form layer s , use rows $j = s, \dots, k$. Columns are $\{0\}$ for $s = 1$, and $s-1, \dots, k-1$ otherwise. Return $(1, 0, i)$ if $i \geq j$, and $(0, D(s-1, i) + C(i+1, j), i)$ otherwise. Search, store predecessors, and backtrack the solution ending at k .

Lemma 7.5 justifies every discarded column. The terminal interface has shrinking feasible suffixes; the program interfaces have expanding feasible prefixes. The completion lemmas include infeasible rows in submatrices created internally by SMAWK, although every original searched row has a feasible entry. Padded comparisons do not evaluate an undefined economic cost.

Theorem 7.6 (Contractual corollary of classical matrix search). *For ordered rational inputs satisfying (8) and $1 \leq m \leq k$, both expected-participation and deterministic/pathwise frontiers through budget m , including an optimal codebook for every budget, require $O(mk)$ rational operations and comparisons and $O(mk + k)$ stored rational numbers and indices. The highest randomized level remains continuous. With total input length L and $S = O(k[L + \log(k+2)])$, a conservative Turing bound is $O(mkS^3)$. Sorting unordered distinct caps adds $O(k \log k)$ exact comparisons.*

The proof is in Electronic Companion Section EC.1.13.

The retained divide-and-conquer implementation is a valid $O(mk \log(k+1))$ alternative using ordinary row-minimum monotonicity. It is an implementation comparator rather than another headline theorem. $O(k)$ exact-rational oracle/comparison work per layer does not imply uniform wall-time superiority for a particular Python implementation.

8 Promise Allocation Away from Saturation

The saturated promise is an endpoint of a larger allocation problem, not a normalization. Fix the branch caps and replace the root promise by $B \in [0, \bar{b}]$, where $\bar{b} = \sum_j \pi_j b_j$. Write $V_m^E(B)$, $V_m^P(B)$, and $V_m^D(B)$ for optimal value with at most m symbols under expected participation, pathwise participation allowing draw-specific intermediate tiers, and deterministic execution, respectively. The information interface remains unchanged. There is no codeword-dependent charge in this section, and terminal levels may range continuously over $[0, 1]$. The full-information value is $V_\infty(B)$.

8.1 An exact threshold for every root promise

Theorem 8.1 (Clipped service levels and exact alphabet). *Suppose f is strictly increasing and strictly concave, and each h_j is convex, nondecreasing, with $h_j(0) = 0$. For every $B \in [0, \bar{b}]$, full information uniquely implements*

$$Y_j = 0, \quad C_j = \min\{b_j, \tau\}, \quad \sum_j \pi_j \min\{b_j, \tau\} = B. \quad (21)$$

At $B = 0$ take $\tau = 0$, and at $B = \bar{b}$ take $\tau = \max_j b_j$. The minimum alphabet attaining this value under any of the three institutions is

$$n(B) = |\{\min(b_j, \tau) : 1 \leq j \leq k\}|.$$

In particular, for $0 \leq B \leq \min_j b_j$, one symbol attains $V_m^I(B) = f(B)$ for every $m \geq 1$ and $I \in \{E, P, D\}$.

Proof. For a full-information feasible policy, let $t_j = \mathbb{E}(Y_j + C_j) \leq b_j$. Replacing it on branch j by $Y_j = 0, C_j = t_j$ preserves the root promise and satisfies even pathwise participation. Jensen's inequality, strict increase of f , and $h_j \geq 0$ show that the replacement weakly improves value, strictly if $\mathbb{E}Y_j > 0$ or C_j is nonconstant. Hence it suffices to maximize $\sum_j \pi_j f(t_j)$ with $0 \leq t_j \leq b_j$ and $\sum_j \pi_j t_j = B$. Strict concavity implies uniqueness. If two uncapped coordinates differ, transferring equal probability-weighted amounts from the larger to the smaller strictly improves this sum; the same transfer shows that a capped coordinate below the common free level must be filled. Thus the unique optimizer is (21), with the endpoint cases interpreted as stated. Continuity of the clipped mean gives existence of τ .

Equality in these improvements requires deterministic terminal action $\min(b_j, \tau)$ on each branch. A symbol's terminal decoder cannot implement two different deterministic actions at the same public vertex. Its conditional output distribution is independent of the originating branch, including when fresh decoder randomization is permitted. Consequently distinct optimal actions require distinct symbols. Assigning one

symbol per distinct action attains the lower bound under all institutions. When $B \leq \min b_j$, all optimal actions equal B . $\square$

For strictly ordered caps, the thresholds at which $n(B)$ changes are explicit: put $B_\ell = \sum_j \pi_j \min(b_j, b_\ell)$. The exact alphabet is $\ell + 1$ for $B_\ell < B \leq B_{\ell+1}$, $1 \leq \ell < k$, and it is one on $[0, B_1]$. At $B = B_\ell$ it is ℓ . Thus exact implementation has a promise-dependent phase diagram, rather than requiring k symbols at every promise.

Proposition 8.2 (Repeated caps at the saturated endpoint). *For $B = \bar{b}$, branches with a common cap b can be aggregated in every deterministic, expected, and pathwise frontier. Their aggregate probability is $p_b = \sum_{j:b_j=b} \pi_j$ and their aggregate intermediate cost is $h_b(y) = p_b^{-1} \sum_{j:b_j=b} \pi_j h_j(y)$. Quadratic costs therefore aggregate by their probability-weighted curvature. The exact alphabet threshold is the number of distinct caps. For arbitrary B , the full-information threshold remains (21), but aggregation of the restricted-alphabet problem with heterogeneous intermediate costs is not asserted.*

Proof. At saturation, the fixed-book expected rule chooses the same terminal mean and intermediate tier for all branches with the same cap. The deterministic and pathwise rules choose the same highest eligible terminal level. Summing their rewards gives precisely the aggregate cost. Optimize over the common books and apply Theorem 8.1 for the threshold. Away from saturation the branch promises themselves can differ, so the saturated aggregation argument does not apply. $\square$

8.2 Uniform transport between any two root promises

Assume in addition that f is Lipschitz with constant L_f on $[0, 1]$ and each h_j is Lipschitz with constant H_j on $[0, b_j]$; put $H = \max_j H_j$. These constants are independent of the alphabet budget. For quadratic primitives one may take $L_f = r$ and $H = \max_j \gamma_j b_j$. The smaller intermediate domain is sufficient because the mean and pathwise replacements used below have tiers at most b_j .

Theorem 8.3 (All-promise value and institutional-gap stability). *For any $0 \leq B_1 \leq B_2 \leq \bar{b}$, $\Delta = B_2 - B_1$, every $m \geq 1$, and $I \in \{E, P, D\}$,*

$$V_m^I(B_2) - L_f \Delta \leq V_m^I(B_1) \leq V_m^I(B_2) + H \Delta. \quad (22)$$

If $G_m(B) = V_m^E(B) - V_m^P(B)$, then

$$|G_m(B_2) - G_m(B_1)| \leq (L_f + H) \Delta. \quad (23)$$

The same bound holds for changes in the loss $V_\infty(B) - V_m^I(B)$. In particular, a certified strict advantage $G_m(B_0) > 0$ persists for every feasible B with $|B - B_0| < G_m(B_0)/(L_f + H)$.

Proof. First contract any policy at $B_2 > 0$ by multiplying all intermediate and terminal actions by $\alpha = B_1/B_2$. The decoder uses the same symbols, and both participation conventions remain feasible. Intermediate cost weakly decreases. The decrease in terminal reward is at most $L_f(1 - \alpha)\mathbb{E}C \leq L_f \Delta$, where expectations include branch weights. The case $B_2 = 0$ is immediate. Taking suprema gives the left inequality.

For the other direction, start from a policy at B_1 . Under expected participation, replace the terminal decoder given a symbol by its mean and the intermediate tier on each branch by its mean. These changes weakly improve reward, preserve the alphabet, and preserve expected obligations. Write $\mu_j = \mathbb{E}C_j$, $y_j = \mathbb{E}Y_j$, and $s_j = b_j - y_j - \mu_j \geq 0$. If $B_1 < \bar{b}$, put $\theta = \Delta/(\bar{b} - B_1)$ and replace y_j by $y_j + \theta s_j$. The total added obligation is $\theta \sum_j \pi_j s_j = \Delta$. Its cost is at most $H\Delta$, and the book and its encoding probabilities are unchanged.

For pathwise participation, replace fresh terminal decoder randomness conditional on symbol z by its mean c_z . Conditional independence from the earlier tier ensures $Y_j + c_z \leq b_j$ almost surely after this replacement; terminal concavity weakly improves reward. Use the draw-specific tier

$$Y'_j = Y_j + \theta(b_j - Y_j - c_z).$$

This is admissible at the stipulated intermediate timing, remains between Y_j and $b_j - c_z$, and adds expected obligation Δ with cost at most $H\Delta$. It requires no extra information at the terminal gateway. The deterministic case is the same construction without random draws. If $B_1 = \bar{b}$, necessarily $\Delta = 0$. Taking suprema, or limits of near-optimal policies, proves the right inequality without requiring a maximizer.

For each institution the value increment lies in $[-H\Delta, L_f\Delta]$. Subtracting two such increments proves the institutional-gap bound. The same constructions apply with unrestricted full information, yielding the loss bound. Every transformation preserves the original number of writable symbols; no concavity of the optimized codebook envelope is assumed. $\square$

In the equal-probability three-branch example of Proposition 5.2, $\bar{b} = 1/2$, $G_2(\bar{b}) = 11/96$, and $L_f + H = 2 + 3/4 = 11/4$. Therefore

$$G_2(B) \geq \frac{11}{96} - \frac{11}{4} \left(\frac{1}{2} - B \right) > 0 \quad \text{for } \frac{11}{24} < B \leq \frac{1}{2}. \quad (24)$$

This is a sufficient certified interval of strict advantage, not a maximal interval or an exact nonsaturated loss formula. The global cell algorithm below computes exact values outside this conservative interval as well. At the other endpoint, $B \leq 1/4$ gives zero loss with one symbol. In between, cap anchoring at the original b_j and pathwise-deterministic equality do not follow from the saturated proofs. Indeed the one-symbol full optimum $C = B < \min b_j$ already shows why original-cap anchoring cannot hold for all promises. The decision insight persists without implying that the saturated algorithm solves every interior promise.

8.3 The nonsaturated allocation problem

The larger problem has an exact reduction even when it does not have the saturated finite-anchor reduction. For a fixed book $c = (c_1, \dots, c_s)$ let F_c be its concave interpolant and define

$$v_{j,c}(t) = F_c(\min\{t, c_s\}) - h_j((t - c_s)_+), \quad c_1 \leq t \leq b_j. \quad (25)$$

Proposition 8.4 (Joint promise allocation and codebook design). *The expected-participation value at any B*

is exactly

$$V_m^E(B) = \max_{\substack{|c| \leq m, 0 \leq c_1 < \dots < c_s \leq 1 \\ c_1 \leq t_j \leq b_j, \sum_j \pi_j t_j = B}} \sum_j \pi_j v_{j,c}(t_j). \quad (26)$$

For a fixed book this is a separable concave resource-allocation problem. Branch targets are the chosen t_j , not mechanically the original caps. Saturation fixes $t_j = b_j$ and recovers the earlier frontier.

Proof. Set $t_j = \mathbb{E}(Y_j + C_j)$ for an arbitrary feasible policy. The mean replacements in Theorem 5.1 are still valid. For a given t_j , maximizing $F_c(\mu) - h_j(t_j - \mu)$ over $c_1 \leq \mu \leq \min(t_j, c_s)$ selects its largest feasible mean by monotonicity. Conversely that mean and its adjacent lottery implement (25) with intermediate tier $(t_j - c_s)_+$. This proves both directions of (26). The response is concave: to the left of c_s it is F_c , to the right it is $f(c_s) - h_j(t - c_s)$, and the one-sided slopes drop from a nonnegative chord slope to a nonpositive cost slope. This also covers a singleton book. Allocation with a common supporting multiplier, allowing subgradient intervals at kinks, solves the fixed-book problem. No interchange of codebook optimization and a dual minimization is used. $\square$

Electronic Companion Proposition EC.6.1 gives an exact supporting-price algorithm and a separately checkable certificate for every fixed-book allocation. Exhaustive certified catalog search extends that computation to all promises on small catalogs; its exponential cost is stated explicitly.

Theorem 9.1 extends the fixed-book response to every realization ceiling. Theorems 9.3 and 10.1 solve the continuous outer design and certify its catalog approximation; thus the inner allocation is no longer the endpoint of the nonsaturated analysis.

9 Joint Continuous Design at Every Promise

We now optimize the continuous alphabet, branch promises, bounded realization exposure, and actual selected-level charges in one problem. Fix $B \in [0, \bar{b}]$ and $\delta \geq 0$. On branch j , an intermediate tier y_j is committed before a terminal lottery is drawn, and

$$t_j = y_j + \mathbb{E}C_j \leq b_j, \quad y_j + C_j \leq b_j + \delta \text{ almost surely}, \quad \sum_j \pi_j t_j = B. \quad (27)$$

A book c has at most m distinct levels in $[0, 1]$ and incurs $\sum_{z \in c} \rho(z)$, once per selected level. Let $\mathcal{V}_m^\delta(B; \rho)$ denote maximum operating payoff less this charge. A finite catalog restriction is denoted $\mathcal{V}_{m, \mathcal{A}}^\delta(B; \rho)$. The expectation cap is not replaced by the weaker realization ceiling. Taking $\delta \geq 1$ makes the realization restriction redundant at the canonical optimum below. Taking $\delta = 0$ gives the full pathwise value, even though that institution initially permits draw-specific intermediate service. Thus these are endpoints of one design problem, not three unrelated extensions.

9.1 Eligibility prefixes replace saturated targets

For a strictly ordered book $c = (c_1, \dots, c_s)$ and branch j , define

$$c^{j,\delta} = \{z \in c : z \leq b_j + \delta\}, \quad d_j = \max c^{j,\delta},$$

and let $F_{j,c}^\delta$ be the piecewise-affine interpolant of f on this prefix, including the singleton interpretation. The prefix depends on the cap and risk ceiling, not on the endogenous target.

Theorem 9.1 (Canonical allocation under a realization ceiling). *Suppose f is strictly increasing and concave and h_j is convex, nondecreasing, and zero at zero. A fixed book is feasible exactly when $c_1 \leq \min_j b_j$ and $c_1 \leq B \leq \bar{b}$. Its optimal branch response at target $c_1 \leq t \leq b_j$ is*

$$v_{j,c}^\delta(t) = F_{j,c}^\delta(\min\{t, d_j\}) - h_j((t - d_j)_+). \quad (28)$$

It is attained by $y = (t - d_j)_+$ and an adjacent lottery of mean $\min(t, d_j)$ in the eligible prefix. Consequently

$$\mathcal{V}_m^\delta(B; \rho) = \max_{\substack{|c| \leq m, \ c_1 \leq t_j \leq b_j \\ \sum_j \pi_j t_j = B}} \left\{ \sum_j \pi_j v_{j,c}^\delta(t_j) - \sum_{z \in c} \rho(z) \right\}. \quad (29)$$

The fixed-book inner problem is concave. With zero charges, $\mathcal{V}_m^0(B; 0) = V_m^P(B)$ and $\mathcal{V}_m^\delta(B; 0) = V_m^E(B)$ for $\delta \geq 1$, at every feasible B .

Proof. Every used terminal level satisfies $C \leq b_j + \delta$ because $y \geq 0$. Writing $\mu = \mathbb{E}C$, the payoff is at most $F_{j,c}^\delta(\mu) - h_j(t - \mu)$. Both terms are nondecreasing as μ increases on $[c_1, \min(t, d_j)]$, so this upper bound is maximized at its right endpoint. If $t \leq d_j$, the adjacent lottery has $y = 0$ and all its support lies below $b_j + \delta$. If $t > d_j$, it is the singleton d_j and $y + d_j = t \leq b_j$. Thus the upper bound is attained without relaxing either constraint. The response is concave: its nonnegative interpolation slopes decrease, and its tail has the nonpositive slopes of $-h_j(t - d_j)$. The feasibility conditions follow by using the lowest level and allocating intermediate service; weighted targets fill the entire interval $[c_1, \bar{b}]$.

For pathwise participation with possibly draw-specific Y , all used C satisfy $C \leq b_j$. At fixed $t = \mathbb{E}(Y + C)$, concavity of f on the discrete support and Jensen's inequality for h_j give the same upper bound with $\delta = 0$ and $\mathbb{E}Y = t - \mu$. The canonical policy attains it and is pathwise feasible. Fresh decoder randomization can first be replaced by its conditional mean for each symbol: this increases terminal reward and preserves all the ceilings under the stipulated absence of an uncharged shared seed. Under expected participation the same mean argument has no eligibility truncation; $\delta \geq 1$ includes every level in $[0, 1]$. These observations establish the two endpoint identities and both directions of (29). $\square$

For quadratic primitives (8), chord $[u, v]$ has slope $r - q(u + v)/2$. On branch j its target capacity is

$$\left[\min\{b_j, v\} - u \right]_+ \mathbf{1}\{v \leq b_j + \delta\}.$$

Starting at $t_j = c_1$, fill these positive slopes in their common decreasing order, followed by zero-cost

intermediate capacity and the quadratic tails. A water-level sweep over events $\gamma_j(b_j - d_j)_+$ fills the remaining targets. It uses $O(k s + k \log(k + 1))$ rational arithmetic operations, including the certificate

$$\sum_j \pi_j v_{j,c}^\delta(t_j) = \lambda B + \sum_j \pi_j \max_{c_1 \leq t \leq b_j} \{v_{j,c}^\delta(t) - \lambda t\}. \quad (30)$$

The right-hand branch maxima are independently checked at eligible knots, the cap, and the clipped tail stationary point. This multiplier certifies the allocation for that book; it is not by itself a global codebook certificate. The outer problem is solved next.

9.2 A finite exact reduction of the continuous outer problem

Fix $s \leq m$ and temporarily allow $0 \leq c_1 \leq \dots \leq c_s \leq 1$. For each branch select one of $2s - 1$ modes. A singleton mode S_ℓ and an adjacent-lottery mode L_ℓ have the following closed descriptions:

$$S_\ell : \quad c_\ell \leq t_j \leq b_j, \quad g_j^S = r c_\ell - \frac{q}{2} c_\ell^2 - \frac{\gamma_j}{2} (t_j - c_\ell)^2; \quad (31)$$

$$L_\ell : \quad c_\ell \leq t_j \leq c_{\ell+1}, \quad t_j \leq b_j, \quad c_{\ell+1} \leq b_j + \delta, \quad g_j^L = r t_j - \frac{q}{2} [(c_\ell + c_{\ell+1}) t_j - c_\ell c_{\ell+1}]. \quad (32)$$

For S_ℓ , use $y_j = t_j - c_\ell$ and terminal level c_ℓ . For L_ℓ , use $y_j = 0$ and the adjacent unbiased lottery. At $c_\ell = c_{\ell+1}$ the constraints force $t_j = c_\ell$, and the formula is the singleton reward. Therefore collisions are not discarded as singular probability formulas. The quadratic interpolation identity eliminates the ratio in the lottery probabilities before codewords become optimization variables.

Each joint mode assignment, the ordered-book inequalities, and $\sum_j \pi_j t_j = B$ define a compact polytope. If $\rho(c) = \alpha c^2 + \beta c + \eta$ is nonnegative on $[0, 1]$, its objective

$$Q(c, t) = \sum_j \pi_j g_j(c, t) - \sum_{\ell=1}^s \rho(c_\ell) \quad (33)$$

is a quadratic polynomial, not necessarily concave. Nonnegative charges, rather than signs of their coefficients, are the assumption. For example, a decreasing nonnegative certification charge is allowed.

Lemma 9.2 (Finite stationary-face search). *A quadratic function on a nonempty compact polytope attains its maximum at a point that is either a vertex or the unique stationary point on the affine hull of some face with negative-definite restricted Hessian. With rational input data, some maximizing point and its value are rational and can be found by finitely many rational linear solves and feasibility tests.*

Proof. Among all faces whose relative interiors contain a global maximizer, choose one of least dimension and a maximizing point x in its relative interior. The derivative along the face's tangent space vanishes and the restricted Hessian is negative semidefinite. If this restriction has a nonzero null direction z , the quadratic is constant along $x + az$ in the affine hull: both its linear and quadratic terms in a vanish. Compactness makes this line reach a proper face while remaining feasible, contradicting the choice of dimension. The restriction is therefore negative definite unless the tangent space is zero dimensional.

Enumerate independent subsets of the defining inequalities, together with the equality constraints, and solve the stationarity equations on each associated affine space. On the face just identified the bordered stationarity matrix is nonsingular; at a vertex, the tangent space is empty and the same conclusion holds. Test feasibility in every original inequality and retain the largest objective among all feasible candidates. Each retained candidate is feasible, and at least one is a global maximizer. Rational elimination gives the rationality claim. Singular bordered systems may be skipped because the argument guarantees a nonsingular maximizing face of smaller dimension; it does not assume nondegeneracy of the original polytope. $\square$

Theorem 9.3 (Exact joint continuous design). *For rational quadratic primitives, rational caps and probabilities, rational $B \in [0, \bar{b}]$ and $\delta \geq 0$, and a rational quadratic charge nonnegative on $[0, 1]$, the maximum of (33) over the closed cells (31)–(32), for every $1 \leq s \leq m$, equals $\mathcal{V}_m^\delta(B; \rho)$. The optimum is attained by a rational book, rational targets, and rational lottery probabilities. Deleting the lottery modes gives the exact deterministic frontier at every promise, with the same charges.*

An exact algorithm enumerates at most $(2s - 1)^k$ cells at size s . After eliminating one target, each cell has dimension at most $s + k - 1$ and at most $s + 1 + 3k$ inequalities. Enumerating its independent active sets and applying Lemma 9.2 gives a finite global algorithm with at most

$$\sum_{s=1}^m (2s - 1)^k 2^{s+1+3k} \quad (34)$$

linear systems, each of polynomial dimension and rational encoding length in the input size. This is an exponential reference algorithm, not a polynomial-time claim or an extension of the saturated Monge complexity.

Proof. Theorem 9.1 turns every optimal book and target allocation into singleton or adjacent-lottery branch modes represented in these cells. Conversely every cell point implements a feasible policy by the displayed reconstruction; all constraints, including the exact root promise, hold. If ordered slots collide or some slots are unused, deleting redundant slots cannot increase a nonnegative charge. The resulting smaller book and its canonical branch rules appear among the enumerated sizes. Hence overcounting charges at duplicate slots cannot change the maximum across all sizes. These two directions prove equality with the original optimization problem, not merely a relaxation or a fixed-book upper bound.

There are finitely many compact cells and each has a continuous quadratic objective. Lemma 9.2 supplies a rational maximizing candidate on each nonempty cell. Adjacent distinct rational levels reconstruct rational probabilities; coincident levels reconstruct a singleton. The mode count, inequality count, and all-subsets upper bound give (34). Standard determinant bounds for rational linear systems give polynomial encoding length for each candidate. Singleton-only modes exactly represent deterministic terminal assignments and intermediate tiers, proving the deterministic statement. $\square$

The reduction solves the unrestricted continuous outer problem at an interior promise, rather than relating fixed-book allocation as global design. Its contract-specific content is the eligibility-prefix reduction and the closed quadratic elimination of lottery probabilities with no loss at collisions. Rational quadratic optimization and active-set reasoning are classical; polynomial-size witnesses are not a new general complexity

principle (Del Pia et al., 2017). The fast saturated Monge theorem remains valuable precisely because the general finite search is exponential.

9.3 Pathwise comparison, repeated caps, and small charges

Theorem 9.3 computes $V_m^E(B)$, $V_m^P(B)$, and $V_m^D(B)$ for every rational promise, using respectively all lottery modes at $\delta = 1$, all modes at $\delta = 0$, and singleton modes. No claim of pathwise–deterministic equality away from saturation is needed for this characterization. Section 14 compares their exact values along an entire declared parameter path. The zero-memory-loss interval $B \leq \min b_j$ in Theorem 8.1 remains important: a positive institution gap is not expected throughout the feasible promise range.

Proposition 9.4 (Pathwise randomization can matter in the interior). *For one branch with $b = 9/10$, $B = 9/20$, $r = 2$, $q = \gamma = 1$, two available symbols, and the nonnegative charge $\rho(c) = 3c(1 - c)$,*

$$\mathcal{V}_2^0(B; \rho) = \frac{171}{400} > \mathcal{V}_2^D(B; \rho) = \frac{9}{160}.$$

The pathwise optimum uses book $(0, 9/10)$, equal draw probabilities, and no intermediate tier. The deterministic optimum uses the single level $9/20$.

Proof. A deterministic policy uses only one level $0 \leq c \leq B$; deleting unused levels weakly lowers its charge. Its net payoff is $-81/800 - (11/20)c + 2c^2$, whose maximum over $[0, B]$ is $9/160$ at B . Every nonsingleton canonical pathwise mode has $0 \leq u \leq B \leq v \leq b$, and net payoff $2B - (u+v)B/2 + uv/2 - 3u(1-u) - 3v(1-v)$. This is convex in (u, v) , so its maximum on that rectangle is at a corner. Checking its four corners, together with the singleton case, gives $(u, v) = (0, 9/10)$ and value $171/400$. Both realizations meet the hard cap and their mean is exactly B . $\square$

This example uses a specified nonnegative, nonmonotone level charge; it is not a counterexample asserted for the uncharged problem. It shows why the saturated pathwise collapse cannot simply be imported into the unified interior design model.

Proposition 9.5 (Homogeneous repeated-cap aggregation). *For the randomized institutions in (27), branches with identical caps and identical intermediate-cost functions can be aggregated by adding their probabilities, at every B , δ , and codeword charge. This holds for each fixed book and after global book optimization.*

Proof. Such branches have identical eligible prefixes and the same concave response (28). Replacing their targets by their probability-weighted mean preserves the root promise and does not lower total reward by Jensen’s inequality. Replicating an aggregate target gives the converse inequality. Charges depend only on the common book. This argument does not aggregate heterogeneous costs or establish a corresponding deterministic restriction. $\square$

The strict off-cap conclusion is also stable under genuinely selected-level charges. In Proposition 6.3, the uncharged two-level book $(1/4, 5/8)$ beats every book of at most two cap levels by $1/1280$. If $0 \leq \rho(c) \leq \eta$ on $[0, 1]$ and $\eta < 1/2560$, evaluating this same off-cap book incurs less than that gap, while a cap-only book

cannot improve its uncharged payoff after a nonnegative charge. Thus no charged optimum is cap-only. This statement preserves the off-cap property, not necessarily the numerical location 5/8.

10 Catalog Approximation Without Promise or Risk Slack

Finite catalogs have two different interpretations. An externally certified menu is a restriction of the feasible technology, for which an exact catalog optimum is the relevant answer. Alternatively the operator may choose a numerical catalog to approximate continuous design. The latter interpretation needs a quantitative comparison to the unrestricted optimum; exactness on the catalog alone does not supply it.

Let $\mathcal{A} \subset [0, 1]$ contain zero and one, and let h be its largest consecutive spacing. Assume f is increasing, concave, and L_f -Lipschitz, and each h_j is convex, nondecreasing, and H_j -Lipschitz on $[0, b_j]$. Let the nonnegative charge ρ be L_ρ -Lipschitz on $[0, 1]$. Define

$$K_m = L_f + \sum_j \pi_j H_j + m L_\rho. \quad (35)$$

For the quadratic model, one can take $L_f = r$ and $H_j = \gamma_j b_j$. Neither the root promise nor the risk tolerance is discretized.

Theorem 10.1 (Uniform exactly feasible catalog certificate). *For every $B \in [0, \bar{b}]$, every $\delta \geq 0$, and every $m \geq 1$,*

$$0 \leq \mathcal{V}_m^\delta(B; \rho) - \mathcal{V}_{m, \mathcal{A}}^\delta(B; \rho) \leq K_m h. \quad (36)$$

More strongly, every canonical continuous-book policy has a catalog policy with no more symbols that preserves each branch target exactly, satisfies the same realization ceiling exactly, and loses at most $K_m h$ in net payoff. The conclusion includes the pathwise endpoint $\delta = 0$ and requires no strict feasibility margin.

Proof. Round each original level c downward to $a(c) = \max\{a \in \mathcal{A} : a \leq c\}$ and merge collisions. Fix a branch target t and let d be its old eligible top. Every old eligible level remains eligible after rounding. In particular the new eligible top d' satisfies $d' \geq a(d) \geq d - h$. The old canonical terminal mean is $\mu = \min(t, d)$; coupling its draws to their rounded levels gives a feasible-support lottery of mean $\tilde{\mu} \in [\mu - h, \mu]$ and expected terminal reward at least the old reward minus $L_f h$.

Do not retain those coupled probabilities. Recompute the canonical lottery on the rounded eligible prefix at the *same* target t . Its terminal mean $\mu' = \min(t, d')$ is at least $\tilde{\mu}$: every rounded old draw is at most d' and its mean is at most t . The new interpolant is increasing, and at mean $\tilde{\mu}$ it dominates the coupled lottery reward. Therefore its reward at μ' also dominates that reward. Moreover $\mu' \geq \mu - h$, so its intermediate tier $t - \mu'$ exceeds the old tier by at most h . The branch operating-payoff loss is at most $(L_f + H_j)h$. Theorem 9.1 ensures exact mean and risk feasibility for this reoptimized lottery, including when the risk bound is binding. All t_j are unchanged, so the root promise remains exactly B .

For each original level, $\rho(a(c)) \leq \rho(c) + L_\rho h$. After merging, the total charge is no larger than the sum of these rounded-slot charges because charges are nonnegative. At most m slots were present. Summing the branch bounds and charge bound gives $K_m h$. Apply the construction to an optimizer, or to

a sequence approaching the supremum, and use inclusion of catalog policies in continuous policies for the other inequality. $\square$

One-sided rounding is necessary but not sufficient without the probability repair. For example, let $b = t = 1/2$, $\delta = 1/8$, and old support $(1/4, 5/8)$ with upper probability $2/3$. Rounding the lower level to zero while retaining the probabilities reduces the mean to $5/12$. Compensating by $y = 1/12$ produces a maximum obligation $17/24 > 5/8$, violating the risk limit. Instead, the new canonical lottery on $(0, 5/8)$ uses upper probability $4/5$ and $y = 0$. It preserves the target and the ceiling exactly. This is the distinction between unconstrained quantizer rounding and the contractual transport in Theorem 10.1.

Corollary 10.2 (Certified approximation for a fixed alphabet budget). *For rational quadratic operating primitives and a nonnegative charge whose rational catalog evaluations and Lipschitz bound are supplied, enumerate every feasible catalog book of size at most m and solve its allocation by (30). With $N = |\mathcal{A}|$, this gives an exact catalog optimum in*

$$O\left(\left[\sum_{s=1}^{\min(m, N)} \binom{N}{s}\right] (km + k \log(k + 1))\right) \quad (37)$$

rational arithmetic/evaluation operations. A uniform mesh with denominator $\lceil K_m/\varepsilon \rceil$ (or one when that ceiling is zero) gives additive error at most ε . The algorithm is polynomial in k , K_m/ε , and the bit complexity of the sampled data for fixed m ; it is not claimed to be polynomial in a variable m , nor is this a multiplicative approximation ratio.

Proof. Increasing subsets exhaust the allowed books and the supporting-price algorithm optimizes every inner problem exactly. Take the largest charge-adjusted value and apply Theorem 10.1. The mesh includes zero, so every feasible promise is implementable even when it is smaller than the first positive grid point. The bound counts evaluation work; a black-box charge oracle is not assumed to have unit bit complexity unless that is explicitly part of the input model. $\square$

The guarantee is uniform in the two contractual constraints. For any institution pair, subtracting their certified value intervals gives a certified interval for the institutional gap. This permits a global accuracy statement away from saturation without presuming that original cap anchors or Monge order survive. When $B = \bar{b}$, the faster charged path recurrence in Section 11 replaces subset enumeration. When the saturated expected problem is uncharged, the continuous Monge algorithm avoids a mesh entirely. Thus the specialized algorithms are faster members of one feasible-policy hierarchy, while (36) connects numerical catalogs to continuous risk-and-charge design.

11 Joint Design of a Risk-Limited Certified Alphabet

A finite certification catalog gives a general intermediate institution and a nontrivial codebook-dependent implementation problem. Let $\mathcal{A} = \{a_1 < \dots < a_N\} \subset [0, 1]$ be the allowable terminal levels, with $a_1 \leq b_1$. Selecting a level a incurs an arbitrary nonnegative charge $\rho(a)$; for example, it may combine a declared

precision/storage charge and a per-level certification charge in the same units as service loss. These are stipulated prices, not estimated operational costs. A book is any subset of $\mathcal{A}$ of size at most m . The objective includes the charges of its actual selected levels, not merely its declared capacity.

The root promise is saturated in this section. Intermediate service y_j is chosen before the symbol draw, with

$$y_j + \mathbb{E}C_j = b_j, \quad y_j + C_j \leq b_j + \delta \quad \text{almost surely.} \quad (38)$$

This explicitly defined institution has the two familiar endpoints: $\delta = 0$ attains the pathwise/deterministic catalog optimum, while $\delta \geq 1$ removes the overrun restriction. The pre-draw restriction is not asserted without loss for arbitrary intermediate risk protocols.

For adjacent selected levels $u < v$, assign to this edge the branches $u \leq b_j < v$ and define

$$K_\delta(u, v) = \sum_{j: u \leq b_j < v} \pi_j \ell_j^\delta(u, v), \quad (39)$$

where, on that interval,

$$\ell_j^\delta(u, v) = \begin{cases} f(b_j) - \frac{v - b_j}{v - u} f(u) - \frac{b_j - u}{v - u} f(v), & v - b_j \leq \delta, \\ f(b_j) - f(u) + h_j(b_j - u), & v - b_j > \delta. \end{cases} \quad (40)$$

At a selected cap the branch loss is zero, including the boundary $b_j = u$. The last selected level u carries tail cost

$$T(u) = \sum_{j: b_j \geq u} \pi_j \{f(b_j) - f(u) + h_j(b_j - u)\}. \quad (41)$$

No original-cap anchoring is imposed on catalog entries.

Theorem 11.1 (Global catalog frontier with level-dependent charges). *For strictly increasing concave f and convex nondecreasing h_j , the following dynamic program globally minimizes service loss plus selected-level charges under the bounded-overrun institution:*

$$\begin{aligned} H_1(v) &= \begin{cases} \rho(v), & v \leq b_1, \\ +\infty, & v > b_1, \end{cases} \\ H_s(v) &= \rho(v) + \min_{u \in \mathcal{A}: u < v} \{H_{s-1}(u) + K_\delta(u, v)\}, \quad s \geq 2, \\ W_m^\delta &= \min_{1 \leq s \leq \min(m, N), v \in \mathcal{A}} \{H_s(v) + T(v)\}. \end{aligned} \quad (42)$$

Backtracking recovers the book, intermediate tiers, and lottery probabilities. All budgets through m require $O(kN^2 + mN^2)$ arithmetic/evaluation work with a direct branch-sum precomputation, and $O(N^2 + mN)$ storage. For rational quadratic data and rational catalog charges the calculation is exact in rational arithmetic. An arbitrary additional charge depending only on s can be added at the final minimization, but level charges must enter the recurrence itself.

Proof. For a fixed book, Proposition EC.5.1 gives the optimal branch rule. On $u \leq b_j < v$, if $v - b_j \leq \delta$ use the adjacent unbiased lottery of mean b_j and set $y_j = 0$; otherwise use u deterministically and set $y_j = b_j - u$. Branches above the last level also use that last level deterministically. Consequently the book's exact loss is the sum of its adjacent-edge costs and its terminal tail. Every branch is counted exactly once; cap endpoints contribute zero regardless of the interval convention.

Increasing paths through the catalog enumerate every possible sorted book. A path can begin exactly at levels no larger than b_1 , which is necessary and sufficient for feasibility. Charging ρ at each visited vertex adds the book's actual implementation cost. Induction on path length gives (42); appending the tail and minimizing over feasible lengths proves both inequalities needed for global optimality. If a selected level is unused, deleting it weakly lowers its nonnegative charge and does not worsen service. Thus the at-most-budget optimization also covers redundant symbols. Store predecessor indices and replay the fixed-book rule to reconstruct a controller.

There are $O(N^2)$ edges, each computable by k branch evaluations, and m layers of $O(N^2)$ comparisons/additions. Storing edges and backpointers gives the stated bounds. In the quadratic rational model every interpolation, difference, and cost is rational. There is no continuous optimizer or rounding tolerance in the recurrence. At $\delta = 0$ every strictly bracketing upper level is excluded, so the branch rule is the deterministic one; at $\delta \geq 1$ every possible upper level is admissible. The endpoint claims follow after optimization over the same catalog. $\square$

This theorem solves the saturated catalog specialization of the unified design problem for arbitrary k, N, m, δ and nonuniform ρ . An exact catalog optimum is not itself an optimum over all real codewords; Theorem 10.1 supplies the missing quantitative comparison. This recurrence does not assume that risk-limited costs are Monge. The continuous uncharged expected frontier retains its separate $O(mk)$ guarantee; the certified-catalog theorem uses a general ordered-path algorithm to accommodate actual level charges and risk eligibility. Entropy coding, shared-table compression depending on pairs of levels, repeated-customer state, and arbitrary codebook-dependent costs are not silently reduced to additive ρ .

Why a loss frontier followed by a resource surcharge is insufficient. Use the off-cap model of Proposition 6.3 and catalog $\{1/4, 1/2, 5/8, 3/4\}$ with two symbols and $\delta = 1$. With zero charges, $(1/4, 5/8)$ attains loss $23/1280$. Give only level $5/8$ charge $1/640$. That book now costs $25/1280$, whereas each cap-only pair $(1/4, 1/2)$ and $(1/4, 3/4)$ costs $24/1280$. The optimum changes its selected levels, not merely its capacity. The exact catalog recurrence finds the change; retaining the uncharged loss-minimizing book and adding its storage bill does not.

Risk exposure along the continuous example. The original two-symbol continuous example is retained as an exact illustration beyond a fixed catalog. Its globally optimal top level and loss are

$$z_\delta = 1/2 + \min(\delta, 1/8), \quad D_2^\delta = z_\delta^2/20 - z_\delta/16 + 3/80. \quad (43)$$




Figure 1: The analytically optimized highest level in the heterogeneous three-branch example. The lower level is $1/4$. The displayed function is $z^2/20 - z/16 + 3/80$ on $[1/2, 3/4]$; its interior minimum is $z = 5/8$, not a branch cap.

On the middle branch the high-realization probability is

$$p_{\text{high}}(\delta) = \frac{1/4}{z_{\delta} - 1/4}, \quad \text{maximum overrun} = z_{\delta} - 1/2.$$

At zero tolerance this branch uses its cap deterministically; at tolerance $1/8$, the high probability is $2/3$ and the possible overrun is $1/8$. Thus probability, mean obligation, and maximum overrun are separate reported quantities. Figure 1 shows the continuous loss profile; the complete global-example proof is in the electronic companion.

12 Pricing Writable Capacity and Other Implementation Resources

Writable state, read-only constants, evaluation work, and compilation are different resources. A finite menu of implementations t for institution I , with charges depending only on capacity m , can be compared by

$$\min_{I,m,t} \{D_m^I + \lambda_w \lceil \log_2 m \rceil + \lambda_R R_{I,t}(m) + \lambda_T Q_{I,t}(m) + C_{I,t}(m)/H_0 + \kappa_I\}. \quad (44)$$

Here H_0 is the number of repeated deployments over which compilation is amortized, and κ_I is an exogenous protocol charge. This is accounting followed by finite comparison, not an additional structural optimization theorem. Its validity requires the stipulated menu charges to be independent of the particular book at a given capacity. Theorem 11.1, in contrast, puts actual additive level charges inside book selection.

With only a writable-bit price, capacity 2^w uses the frontier at $\min(2^w, k)$. For the equal-probability three-cap example, the zero-bit loss is $7/16$, the one-bit losses are $1/96$ under expected and $1/8$ under pathwise participation, and two bits attain full information. Thus a bit price strictly between $1/96$ and $1/8$ selects one bit under expected participation and two under pathwise participation when all other protocol charges agree. A different exogenous κ_I can change that comparison. This example interprets the proved frontiers; it does not estimate the monetary value of certification or customers' tolerance for risk.

13 Inherited Verification and Computational Evidence

The computational package separates exact finite checks, practical measurements, and a standard-quantizer comparison. The theorems establish the all-size and continuous claims; finite checks are regression evidence. Prices and branch populations are synthetic inputs, not empirical estimates.

Preceding promise and catalog checks. Independent enumeration of all catalog subsets and all one/two-point branch supports agrees with every reported finite-catalog frontier, without assuming adjacent support in the comparison oracle. Tests include repeated caps, noncap catalog entries, zero intermediate curvature, nonuniform charges, and zero, intermediate, and unrestricted overrun tolerances. Clipped allocations satisfy exact weighted promises and allocation optimality conditions; direct scaling and saturation lifts verify the value inequalities on heterogeneous policies. Fixed-book responses are checked for concavity across their kinks. Invalid floating-point and malformed inputs are rejected.

The executed suite records 768 catalog frontier equalities, 26,604 direct-support book checks, 600 scaling/lift checks, 660 clipped-allocation checks, 2,460 response-concavity checks, and 128 repeated-cap equalities. The charge-switch example and the interval in (24) are checked in exact fractions.

Comparison with an ordinary scalar quantizer. The baseline recursion over ordered cells is the same recursion used for scalar mean-square quantization; the cell oracle is not. For the three equal-probability caps, both optimal two-cell mean-square partitions have loss $1/96$. Their centroids are $(1/4, 5/8)$ or $(3/8, 3/4)$, respectively. Each merged cell includes a branch whose cap is below its centroid, so that centroid reconstruction is infeasible under the pathwise contract, even with nonnegative intermediate service. This exact comparison identifies the changed objective and constraint rather than benchmarking two algorithms that solve different problems. The expected contract instead uses the classical unbiased encoder with a contractual tail.

Inherited continuous-frontier evidence. The complete preceding exact suite, including completion ties, selected all-padded rows, controller replay, independent PADS substitution, and earlier regression suites, is rerun in isolated directories. Historical source and result files are not overwritten. The preceding process-isolated continuous-frontier timings through 16,384 branches and reduced-network HiGHS certificates are retained with their original provenance in the companion. The HiGHS comparison validates the finite reduction, not the cap-anchoring theorem or performance at the larger scale. PADS substitutes the search layer only. Neither comparison is presented as a new experiment on operational data.

Additional catalog landscapes. The executed catalog study uses uniform caps, two separated cap clusters, and a population with concentrated branch weights. Each family is evaluated at $(k, N) = (32, 16)$ and $(128, 32)$, eight-symbol capacity, and tolerances 0, $1/8$, 1, with heterogeneous curvatures and level prices. All 18 cases reconstruct and replay their selected books exactly. Times include catalog construction inside the solver, the dynamic program, reconstruction, and direct replay in one process; they are not directly

compared with the predecessor’s process-isolated timings. Full numerical records and the timing definition accompany the code.

The nonsaturated extension additionally passes 1715 exact primal/dual-event comparisons, 6955 supporting branch maxima, 72 saturated catalog/DP equalities, 48 interior-promise catalog comparisons, and 480 arbitrary-promise transport checks. The supporting-price equality certifies each fixed-book solution; finite regression is not used as proof of the general theorem.

14 Joint-Design Parameter Studies

The new computations ask how the optimal continuous book changes with the promise, realization ceiling, and selected-level charges, and how a finite grid approaches that optimum. They are synthetic studies of the declared model, not an empirical calibration. Every optimum in the first two studies is obtained by the global rational face search of Theorem 9.3; its reconstructed book is then checked by the independent fixed-book multiplier certificate. This is stronger than treating a locally optimized book as a continuous benchmark, but it is not a substitute for the theorem’s proof.

The institutional gap across promises. Take caps $(1/4, 1/2, 3/4)$ with equal probabilities, $r = 2$, $q = 1$, $\gamma_j = 1$, $m = 2$, and no codeword charges. Table 2 reports eight promises and compares the unrestricted expected, pathwise, and deterministic optima. The expected/pathwise gap is zero at each of the five sampled promises through $5/12$, and strictly positive at each sampled promise from $11/24$ to saturation. At $B = 11/24$, the exact gap is $11/192$ even though the old transport certificate gives zero. Thus the sufficient interval $11/24 < B \leq 1/2$ is conservative; it is not the maximal separation interval. These samples alone do not identify the exact transition point. Pathwise and deterministic values coincide along this particular uncharged sample path, not universally in the joint charged model: Proposition 9.4 gives an exact counterexample to that extrapolation.

Risk and actual level prices. For caps $(1/4, 1/2, 3/4)$ and probabilities $(7/20, 3/5, 1/20)$, retain the same quadratic primitives and $m = 2$. We cross four promises $B \in \{1/3, 3/8, 2/5, 17/40\}$, four ceilings $\delta \in \{0, 1/16, 1/8, 1\}$, and two price schedules: zero and $\rho(c) = c/50 + 1/200$. This gives 32 joint continuous designs. Table 3 displays the $B = 2/5$ and saturated slices; the archived rational records include every target, intermediate tier, lottery, and enabled level. Because charges enter before optimization, changes in the selected book are not post hoc accounting adjustments. The algorithm permits off-cap levels, endogenous branch targets, and level collisions throughout these comparisons.

Continuous benchmark versus certified grids. At $B = 2/5$, $\delta = 1/16$, and $\rho(c) = c/50 + 1/200$, the continuous net optimum is $7169/10400$. Table 4 compares it with uniform grids of 9 through 129 points. The realized gap falls from $19/10400$ to $57/166400$ along these grids. The theorem’s worst-case bound is deliberately conservative: here $K_2 = 493/200$, whereas the observed errors are much smaller than $K_2 h$. The unchanged optimum over some successive grids is not a violation of convergence; nested grids give

monotone values but need not improve at every refinement. All displayed grid policies satisfy the original promise and realization ceiling exactly.

Scaling of the new problem. Table 5 reports separate measurements for all-promise charged catalog enumeration, rather than reusing the earlier 16,384-branch saturated benchmark. We use $m = 2$, caps $b_j = 1/4 + j/[2(k + 2)]$, equal probabilities, $\gamma_j = 1 + ((j - 1) \bmod 4)$, $B = 3\bar{b}/4$, $\delta = 1/16$, and the preceding affine charge. These small exact-arithmetic runs expose the dependence on both branch count and catalog size; they do not establish large-scale performance of the exponential continuous face algorithm. The continuous reference runs have three branches, except for the one-branch pathwise example and the two-branch regression cases. Each timing is one measured solve, excluding import and reader compilation, with no averaging or hardware-independent interpretation. The environment and machine-readable timings are retained separately from exact values.

15 Conclusion

A limited execution alphabet is not only a communication budget. When upstream obligations differ and intermediate service is costly, participation determines which terminal lotteries can implement each accepted promise. The saturated quadratic problem has a sharp structural solution: a contractual tail selects an off-cap highest level, and a crossing-difference identity allows classical linear-time matrix search after continuous minimization.

Away from saturation, the branch promises become decisions and the original cap anchors are no longer sufficient. The eligibility-prefix theorem handles those endogenous promises together with a realization ceiling and actual level charges. Closed quadratic cells give an exact continuous global algorithm, including a pathwise comparison, while a downward-rounding and probability-repair theorem connects numerical catalogs to continuous design without relaxing feasibility. The general exact algorithm is exponential; the catalog alternative has a uniform additive guarantee for a fixed symbol budget. Neither result rebrands the saturated complexity as a general bound.

The theory identifies what is purchased by relaxing realization exposure or opening an additional certified execution level. It also identifies where memory is immaterial, when identical branches can be aggregated, and why a strict off-cap advantage survives sufficiently small nonnegative charges. The synthetic computations make those implications auditable. Models with post-draw exit, hidden reports, coupled codebook costs, or additional history channels require their own constraints; they are not tacit conclusions of the stated renewal architecture.

Code and data availability. The package supplies the exact algorithms, deterministic input generators, complete parameter-path results, independent support and symbolic-linear-algebra checks, and reader-build commands. Its entry point is `revisions/or-r42-certified-joint-design-20260924/code/build.py`. New tables come from the current result files; machine-dependent timings and unchanged predecessor evidence are identified separately. No proprietary or personal data are used.

References

- Aardal K, van den Berg P, Gijswijt D, Li S (2014) Approximation algorithms for hard capacitated k -facility location problems. Author manuscript, arXiv:1311.4759v4, September 12.
- Aggarwal A, Klawe MM, Moran S, Shor P, Wilber R (1987) Geometric applications of a matrix-searching algorithm. *Algorithmica* 2:195–208. doi:10.1007/BF01840359.
- Chan TM (2021) (Near-)linear-time randomized algorithms for row minima in Monge partial matrices and related problems. *Proceedings of the 2021 ACM–SIAM Symposium on Discrete Algorithms*, 1465–1482. doi:10.1137/1.9781611976465.88.
- Croci M, Fasi M, Higham NJ, Mary T, Mikaitis M (2022) Stochastic rounding: Implementation, error analysis and applications. *Royal Society Open Science* 9:211631. doi:10.1098/rsos.211631.
- Del Pia A, Dey SS, Molinaro M (2017) Mixed-integer quadratic programming is in NP. *Mathematical Programming* 162:225–240. doi:10.1007/s10107-016-1036-0.
- Eppstein D (2005) PADS: Python Algorithms and Data Structures. *SMAWK.py*, offline matrix-search implementation, originally written 2002 and revised 2005. University of California, Irvine. MIT-licensed distribution; source version and excerpt are identified in the code supplement.
- Fu M (2024) A tutorial on quantized feedback control. *IEEE/CAA Journal of Automatica Sinica* 11(1):5–17. doi:10.1109/JAS.2023.123972.
- Grønlund A, Larsen KG, Mathiasen A, Nielsen JS, Schneider S, Song M (2017) Fast exact k -means, k -medians and Bregman divergence clustering in 1D. Preprint, arXiv:1701.07204, version 4.
- Huangfu Q, Hall JAJ (2018) Parallelizing the dual revised simplex method. *Mathematical Programming Computation* 10(1):119–142. doi:10.1007/s12532-017-0130-5.
- Jourdain B, Pagès G (2021) Optimal dual quantizers of 1D log-concave distributions: Uniqueness and Lloyd like algorithm. *Journal of Approximation Theory* 267:105581. doi:10.1016/j.jat.2021.105581.
- Kaplan H, Mozes S, Nussbaum Y, Sharir M (2017) Submatrix maximum queries in Monge matrices and partial Monge matrices, and their applications. *ACM Transactions on Algorithms* 13(2), Article 26. doi:10.1145/3039873.
- Klawe MM, Kleitman DJ (1990) An almost linear time algorithm for generalized matrix searching. *SIAM Journal on Discrete Mathematics* 3(1):81–97. doi:10.1137/0403009.
- Pagès G, Wilbertz B (2012) Intrinsic stationarity for vector quantization: Foundation of dual quantization. Author manuscript, arXiv:1010.4642v2, March 27. Original preprint 2010.
- Spear SE, Srivastava S (1987) On repeated moral hazard with discounting. *Review of Economic Studies* 54(4):599–617. doi:10.2307/2297484.

- Stouffer K, Pease M, Tang C, Zimmerman T, Pillitteri V, Lightman S, Hahn A, Saravia S, Sherule A, Thompson M (2023) Guide to operational technology (OT) security. NIST Special Publication 800-82 Rev. 3. doi:10.6028/NIST.SP.800-82r3.
- Thomas J, Worrall T (1988) Self-enforcing wage contracts. *Review of Economic Studies* 55(4):541–554. doi:10.2307/2297404.
- Virtanen P, Gommers R, Oliphant TE, et al. (2020) SciPy 1.0: Fundamental algorithms for scientific computing in Python. *Nature Methods* 17:261–272. doi:10.1038/s41592-019-0686-2.
- Wu X (1991) Optimal quantization by matrix searching. *Journal of Algorithms* 12(4):663–673. doi:10.1016/0196-6774(91)90039-2.
- Yao FF (1980) Efficient dynamic programming using quadrangle inequalities. *Proceedings of the Twelfth Annual ACM Symposium on Theory of Computing*, 429–435. doi:10.1145/800141.804691.
- Zhang H (2012) Solving an infinite horizon adverse selection model through finite policy graphs. *Operations Research* 60(4):850–864. doi:10.1287/opre.1120.1056.

Table 1: Theorem-level comparison with established principles.

Result here	Predecessor principle	Contract-specific statement proved here
Deterministic cells	Ordered quantizer DP; Wu	Participation fixes lower-cap reconstruction with heterogeneous intermediate costs.
Adjacent lottery	Stochastic rounding; dual quantization	Expected participation selects that primitive jointly with shortfall service; pathwise participation changes feasibility.
Continuous top level	Optimized quantizer grids; Jourdain–Pagès	Cap-anchored representation combined with a heterogeneous, one-sided contractual tail; exact off-cap optimum.
Linear frontier	Monge DP and classical matrix search	Tail crossing difference, optimized Monge order, and exact oracle/domain verification make the classical bound applicable.
Promise robustness	Concave allocation and Lipschitz comparison	Same-alphabet transport gives all-promise bounds, exact fixed-book certificates, and a strict nonsaturated advantage.
Charged risk catalog	Ordered shortest paths; fixed-opening costs	Exact risk-eligibility edge costs and actual level charges jointly optimize service and implementation.
Joint continuous design	Classical rational quadratic optimization	Eligibility-prefix and closed-cell elimination solve endogenous promises, risk and level charges together.
Certified mesh design	Quantizer approximation and rounding	Downward rounding plus probability repair preserves every target and ceiling, with a uniform additive certificate.

Table 2: Continuous institutional gaps at eight root promises.

B	$V_2^E - V_2^P$	Transport lower bound	$V_2^P - V_2^D$
$\frac{1}{8}$	0	0	0
$\frac{1}{4}$	0	0	0
$\frac{1}{3}$	0	0	0
$\frac{3}{8}$	0	0	0
$\frac{5}{12}$	0	0	0
$\frac{11}{24}$	$\frac{11}{192}$	0	0
$\frac{23}{48}$	$\frac{11}{128}$	$\frac{11}{192}$	0
$\frac{1}{2}$	$\frac{11}{96}$	$\frac{11}{96}$	0

Equal branch probabilities; zero selected-level charges. These are globally optimized continuous books. A zero sampled gap does not establish equality at every intermediate promise.

Table 3: Joint continuous books with endogenous promises, risk ceilings, and charges.

B	δ	Charge	Optimal book	Net payoff
$\frac{2}{5}$	0	0	$(\frac{1}{4}, \frac{25}{52})$	$\frac{297}{416}$
$\frac{2}{5}$	0	$c/50 + 1/200$	$(\frac{1}{4}, \frac{25}{52})$	$\frac{7169}{10400}$
$\frac{2}{5}$	$\frac{1}{16}$	0	$(\frac{1}{4}, \frac{25}{52})$	$\frac{297}{416}$
$\frac{2}{5}$	$\frac{1}{16}$	$c/50 + 1/200$	$(\frac{1}{4}, \frac{25}{52})$	$\frac{7169}{10400}$
$\frac{2}{5}$	$\frac{1}{8}$	0	$(\frac{1}{4}, \frac{25}{52})$	$\frac{297}{416}$
$\frac{2}{5}$	$\frac{1}{8}$	$c/50 + 1/200$	$(\frac{1}{4}, \frac{25}{52})$	$\frac{7169}{10400}$
$\frac{2}{5}$	1	0	$(\frac{1}{4}, \frac{25}{52})$	$\frac{297}{416}$
$\frac{2}{5}$	1	$c/50 + 1/200$	$(\frac{1}{4}, \frac{25}{52})$	$\frac{7169}{10400}$
$\frac{17}{40}$	0	0	$(\frac{1}{4}, \frac{1}{2})$	$\frac{117}{160}$
$\frac{17}{40}$	0	$c/50 + 1/200$	$(\frac{1}{4}, \frac{1}{2})$	$\frac{113}{160}$
$\frac{17}{40}$	$\frac{1}{16}$	0	$(\frac{1}{4}, \frac{9}{16})$	$\frac{3747}{5120}$
$\frac{17}{40}$	$\frac{1}{16}$	$c/50 + 1/200$	$(\frac{1}{4}, \frac{1}{2})$	$\frac{113}{160}$
$\frac{17}{40}$	$\frac{1}{8}$	0	$(\frac{1}{4}, \frac{5}{8})$	$\frac{937}{1280}$
$\frac{17}{40}$	$\frac{1}{8}$	$c/50 + 1/200$	$(\frac{1}{4}, \frac{1}{2})$	$\frac{113}{160}$
$\frac{17}{40}$	1	0	$(\frac{1}{4}, \frac{5}{8})$	$\frac{937}{1280}$
$\frac{17}{40}$	1	$c/50 + 1/200$	$(\frac{1}{4}, \frac{1}{2})$	$\frac{113}{160}$

Caps $(1/4, 1/2, 3/4)$, probabilities $(7/20, 3/5, 1/20)$, and two available symbols. The full 32-case archive also records targets, intermediate tiers, and lotteries.

Table 4: Exact grid values compared with the continuous optimum $7169/10400$.

N	Grid book	Actual gap	Certified gap $K_2 h$	Seconds
9	$(\frac{1}{4}, \frac{1}{2})$	$\frac{19}{10400}$	$\frac{493}{1600}$	0.007
17	$(\frac{1}{4}, \frac{1}{2})$	$\frac{19}{10400}$	$\frac{493}{3200}$	0.023
33	$(\frac{1}{4}, \frac{1}{2})$	$\frac{19}{10400}$	$\frac{493}{6400}$	0.083
65	$(\frac{1}{4}, \frac{31}{64})$	$\frac{57}{166400}$	$\frac{493}{12800}$	0.311
129	$(\frac{1}{4}, \frac{31}{64})$	$\frac{57}{166400}$	$\frac{493}{25600}$	1.213

Interior promise $B = 2/5$, tolerance $\delta = 1/16$, and charge $c/50 + 1/200$. Every policy meets its promise and ceiling without slack. Timings are single reference-implementation solves, not averaged performance estimates.

Table 5: Separate scaling of all-promise charged catalog enumeration.

Branches k	Catalog N	Budget m	Feasible books checked	Seconds
8	17	2	75	0.058
8	33	2	285	0.221
8	65	2	1110	0.880
32	17	2	75	0.221
32	33	2	261	0.790
32	65	2	969	2.966
128	17	2	75	0.877
128	33	2	261	3.131

Each row solves the new nonsaturated problem with a realization ceiling and actual codeword charges. These measurements do not reuse the inherited saturated-frontier timings and do not measure the exponential continuous face search.